

Sep 22, 15 16:14

adc.c

Page 1/7

```

#
/*
 * ADC interface module
 *
 * This module provides the interface from the Reader to Comedi. It
 * handles interaction with the Comedi device and mapping the Comedi
 * data buffer.
 *
 * The routines (apart from adc_new(), which returns a pointer to the
 * semi-opaque adc structure representing this object) in this module
 * return 0 on success and -1 on failure; they leave error information
 * in the adc structure from which it can be retrieved with the
 * adc_error() method.
 */

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <stdint.h>
#include <errno.h>
#include <string.h>
#include <comedi.h>
#include <comedilib.h>

#include <time.h>
#include <sys/time.h>
#include <sys/resource.h>

#include "assert.h"
#include "strbuf.h"
#include "chunk.h"
#include "lut.h"
#include "mman.h"
#include "queue.h"
#include "adc.h"

/*
 * Private information for the ADC module.
 */

#define N_USBDUX_CHANS 16

#define MIN_SAMPLING_FREQUENCY 6e4 /* Minimum sampling frequency per channel [Hz] */
#define MAX_SAMPLING_FREQUENCY 3.75e5 /* Maximum sampling frequency per channel [Hz] */
#define MIN_COMEDI_BUF_SIZE 8 /* Minimum Comedi Buffer size [MiB] */
#define MAX_COMEDI_BUF_SIZE 256 /* Maximum Comedi Buffer size [MiB] */

struct _adc {
    const char *a_path; /* The path to the Comedi device (assumed permanent string) */
    comedi_t *a_device; /* Comedi device handle */
    int a_devflags; /* Comedi device flags */
    int a_fd; /* Device file descriptor */
    int a_req_bufsz_mib; /* Requested buffer size [MiB] */
    int a_bufsz_bytes; /* Size of the buffer in bytes */
    int a_bufsz_samples; /* Size of the buffer in samples */
    sampl_t *a_comedi_ring; /* Ring buffer for the device */
    double a_totfreq; /* Total sampling frequency */
    int a_intersample_ns; /* Time between samples [ns] */
    int a_range; /* Current conversion range */
    int a_raw; /* Don't convert the data, deliver it raw */
    convertfn a_convert; /* Current conversion function */
    comedi_cmd a_command; /* Comedi command descriptor structure */
    unsigned a_chans[N_USBDUX_CHANS]; /* Channel descriptors for hardware channels */
    int a_running; /* True when an ADC data conversion is running */
    int a_live; /* True when we have seen data, and a_start_time is set */
    uint64_t a_start_time; /* Time the current data conversion stream started */
    uint64_t a_head_time; /* Timestamp of latest buffer sample */
    uint64_t a_head; /* Latest sample present in the ring buffer */
    uint64_t a_tail; /* Earliest sample present in the ring buffer */
};

#define USBDUXFAST_COMEDI_500mV 1 /* Bit 3 control output is 0 iff the CR_RANGE is one */
#define USBDUXFAST_COMEDI_750mV 0 /* Bit 3 control output is 1 iff the CR_RANGE is zero */

```

Sep 22, 15 16:14

adc.c

Page 2/7

```

/*
 * ADC is a singleton class, so we can get away with defining a single private structure.
 */

private struct _adc snapshot_adc;

/*
 * Error string function for strbuf module.
 */

private int comedi_error_set_up = 0;

private const char *comedi_error() {
    return comedi_strerror( comedi_errno() );
}

/*
 * Allocate and set up a new ADC descriptor.
 */

public adc adc_new(strbuf e) {
    adc ret = &snapshot_adc;

    if( !comedi_error_set_up++ ) { /* Install the routine to interpolate %C strings */
        int ret = register_error_percent_handler('C', comedi_error);
        assertv(ret==0, "Failed to register handler for Comedi errors (%mC): %m\n");
    }
    ret->a_fd = -1;
    return ret;
}

/*
 * Release ADC resources and free an ADC structure.
 */

public void adc_destroy(adc a) {
    adc_stop_data_transfer(a);

    if(a->a_fd >= 0)
        close(a->a_fd);

    if(a->a_device)
        comedi_close(a->a_device);

    if(a->a_comedi_ring)
        munmap(a->a_comedi_ring, a->a_bufsz_bytes);

    /* Zero the structure -- back to initial state */
    bzero(a, sizeof(struct _adc));
}

/*
 * Set the device path
 */

public int adc_set_device(adc a, const char *device) {
    a->a_path = device;
    return 0;
}

/*
 * Set the total capture sampling frequency from the per-channel frequency.
 */

public int adc_set_chan_frequency(adc a, strbuf e, double *freq) {
    double f = *freq;

    if(f < MIN_SAMPLING_FREQUENCY || f > MAX_SAMPLING_FREQUENCY) {
        strbuf_appendf(e, "Sampling frequency %g not within compiled-in ADC limits [%g,%g] Hz",
            f, MIN_SAMPLING_FREQUENCY, MAX_SAMPLING_FREQUENCY);
        return -1;
    }

    int ns = 1e9 / (f*NCHANNELS); /* Inter-sample period */
    int xtra = ns % 100;

```

Sep 22, 15 16:14

adc.c

Page 3/7

```

/* Adjust period for 30[MHz] USBDUXfast clock rate */
ns = 100 * (ns / 100);
if( xtra > 17 && xtra < 50 )
    ns += 33;
if( xtra >= 50 && xtra < 83 )
    ns += 67;
if( xtra >= 84 )
    ns += 100;
a->a_intersample_ns = ns; /* Need a plausible value at all times for computing snapshot data */
a->a_totfrequency = 1e9 / ns;
*freq = a->a_totfrequency / NCHANNELS;
return 0;
}

/*
 * Set the desired ring buffer size.
 */

public int adc_set_bufsz(adc a, strbuf e, int bufsz) {
    if(bufsz < MIN_COMEDI_BUF_SIZE || bufsz > MAX_COMEDI_BUF_SIZE) {
        strbuf_appendf(e, "Comedi buffer size %d MiB outwith compiled-in range [%d,%d] MiB",
            bufsz, MIN_COMEDI_BUF_SIZE, MAX_COMEDI_BUF_SIZE);
        return -1;
    }
    a->a_req_bufsz_mib = bufsz;
    return 0;
}

/*
 * Set the desired ADC range
 */

public int adc_set_range(adc a, strbuf e, int range) {
    /* Set up the conversion function: 500mV or 750mV FSD */
    switch(range) {
        case 500: /* Narrow FSD range */
            a->a_convert = a->a_raw? convert_raw_raw : convert_raw_500mV;
            a->a_range = USBDUXFAST_COMEDI_500mV;
            break;

        case 750: /* Wide FSD range */
            a->a_convert = a->a_raw? convert_raw_raw : convert_raw_750mV;
            a->a_range = USBDUXFAST_COMEDI_750mV;
            break;

        default:
            strbuf_appendf(e, "Comedi range spec %d unknown", range);
            return -1;
    }
    return 0;
}

/*
 * Set ADC to raw mode, i.e. don't range-map the incoming data.
 */

public void adc_set_raw_mode(adc a, int on) {
    a->a_raw = (on != 0);
    if(a->a_raw)
        a->a_convert = convert_raw_raw;
    else {
        if(a->a_range == USBDUXFAST_COMEDI_500mV)
            a->a_convert = convert_raw_500mV;
        if(a->a_range == USBDUXFAST_COMEDI_750mV)
            a->a_convert = convert_raw_750mV;
    }
}

/*
 * Initialise the ADC structure for data capture.
 */

public int adc_init(adc a, strbuf e) {
    int ret;

```

Sep 22, 15 16:14

adc.c

Page 4/7

```

int i;

if( !a->a_path ) {
    strbuf_appendf(e, "Comedi device path not set");
    return -1;
}

/* Open the Comedi device */
a->a_device = comedi_open(a->a_path);
if(a->a_device == NULL) {
    strbuf_appendf(e, "Comedi device %s failure setting up ADC structure: %C", a->a_path);
    return -1;
}
a->a_fd = comedi_fileno(a->a_device);
a->a_devflags = comedi_get_subdevice_flags(a->a_device, 0);

/* Initialise Comedi streaming buffer */
int request = a->a_req_bufsz_mib * 1024 * 1024;
ret = comedi_get_buffer_size(a->a_device, 0);
if( request > ret ) {
    ret = comedi_get_max_buffer_size(a->a_device, 0);
    if( request > ret ) {
        ret = comedi_set_max_buffer_size(a->a_device, 0, request);
        if( ret < 0 ) {
            strbuf_appendf(e, "Comedi set max buffer to %dMiB failed: %C", a->a_req_bufsz_mib);
            return -1;
        }
    }
    ret = comedi_set_buffer_size(a->a_device, 0, request);
    if( ret < 0 ) {
        strbuf_appendf(e, "Comedi set streaming buffer to %dMiB failed: %C", a->a_req_bufsz_mib);
        return -1;
    }
}

a->a_bufsz_bytes = comedi_get_buffer_size(a->a_device, 0);
a->a_bufsz_samples = a->a_bufsz_bytes / sizeof(sampl_t);
comedi_set_global_or_behavior(COMEDI_OOR_NUMBER);

/* Initialise the command structure */
ret = comedi_get_cmd_generic_timed(a->a_device, 0, &a->a_command, N_USBDUX_CHANS, 0);
if(ret < 0) {
    strbuf_appendf(e, "Comedi generic command setup failed: %C");
    return -1;
}

/* Set the command parameters from the reader parameter values */
for(i=0; i<N_USBDUX_CHANS; i++)
    a->a_chans[i] = CR_PACK_FLAGS(i, a->a_range, AREF_GROUND, 0);
a->a_command.chanlist = &a->a_chans[0];
a->a_command.stop_src = TRIG_NONE;
a->a_command.stop_arg = 0;
a->a_command.convert_arg = a->a_intersample_ns;

/* Ask the driver to check the command structure and complete any omissions */
(void) comedi_command_test(a->a_device, &a->a_command);
ret = comedi_command_test(a->a_device, &a->a_command);
if( ret < 0 ) {
    strbuf_appendf(e, "Comedi second command test fails: %C");
    return -1;
}

/* Check the timing: a difference here means a problem with the driver */
if(a->a_command.convert_arg != a->a_intersample_ns) {
    a->a_intersample_ns = a->a_command.convert_arg;
    a->a_totfrequency = 1e9 / a->a_command.convert_arg;
    /* TODO: consider logging a warning here */
}

/* Map the Comedi buffer into memory, duplicated */
void *map = mmap_and_lock(a->a_fd, 0, a->a_bufsz_bytes, PROT_RDONLY|PREFAULT_RDONLY|MAL_LOCKED|MAL_DOUBLED);
if(map == NULL) {
    strbuf_appendf(e, "Unable to mmap Comedi streaming buffer: %m");
    return -1;
}

```

Sep 22, 15 16:14

adc.c

Page 5/7

```

a->a_comedi_ring = map;

/* Initialise the sample position indices */
a->a_head = 0;
a->a_tail = 0;
a->a_start_time = 0;
a->a_head_time = 0;
a->a_running = 0;
return 0;
}

/*
 * Start the ADC data collection.
 */

public int adc_start_data_transfer(adc a, strbuf e) {
    int ret;

    /* Execute the command to initiate data acquisition */
    ret = comedi_command(a->a_device, &a->a_command);
    if(ret < 0) {
        strbuf_appendf(e, "Comedi command failed: %C");
    }
    else {
        a->a_running = 1;
    }
    return ret;
}

/*
 * Stop the ADC data collection.
 */

public void adc_stop_data_transfer(adc a) {
    if(a->a_running) {
        comedi_cancel(a->a_device, 0);
        a->a_running = 0;
        a->a_live = 0;
    }
}

/*
 * Convert a sample index into an ADC ring pointer. This is used by
 * adc_setup_chunk(). It depends on the fact that the Comedi buffer
 * is double-mapped so the pointer is always the start of a contiguous
 * block of memory that will at some time hold the data for the chunk.
 */

private sampl_t *adc_sample_to_ring_ptr(adc a, uint64_t sample) {
    return &a->a_comedi_ring[sample % a->a_bufsz_samples];
}

/*
 * Set up the ADC-dependent information in a chunk, and determine whether the chunk is recordable.
 * In case of error, set the c_error strbuf and set the c_status code to SNAPSHOT_ERROR.
 */

public void adc_setup_chunk(adc a, chunk_t *c) {
    if(c->c_first < a->a_tail) { /* Too late */
        strbuf_appendf(c->c_error, "Chunk was %d [us] too late", (int)((a->a_tail - c->c_first)/1000));
        c->c_ring = NULL;
        return;
    }
    c->c_ring = adc_sample_to_ring_ptr(a, c->c_first);
    return;
}

/*
 * Convert times to sample indices and vice versa
 */

public uint64_t adc_time_to_sample(adc a, uint64_t time) {
    uint64_t ret;

    ret = (time - a->a_start_time) / a->a_intersample_ns;

```

Sep 22, 15 16:14

adc.c

Page 6/7

```

}
return ret;
}

public uint64_t adc_sample_to_time(adc a, uint64_t sample) {
    uint64_t ret;

    ret = a->a_start_time + sample*a->a_intersample_ns;
    return ret;
}

/*
 * Read-only access to some ADC parameters
 */

public int adc_ns_per_sample(adc a) {
    return a->a_intersample_ns;
}

public double adc_tot_frequency(adc a) {
    return a->a_totfrequency;
}

public uint64_t adc_capture_start_time(adc a) {
    return a->a_start_time;
}

public uint64_t adc_capture_head_time(adc a) {
    return a->a_head_time;
}

public int adc_is_running(adc a) {
    return a && a->a_running;
}

public int adc_is_live(adc a) {
    return a && a->a_live;
}

public uint64_t adc_ring_head(adc a) {
    return a->a_head;
}

public uint64_t adc_ring_tail(adc a) {
    return a->a_tail;
}

/*
 * The buffer strategy implied below is an explicit one of
 * periodically advancing the tail to avoid buffer overrun. The data
 * bounded by the tail and head pointers in the ring buffer is valid,
 * under this explicit strategy.
 */

/*
 * Recognise data in the Comedi buffer: ask Comedi how much new data
 * is available, set the local data structure to match, tell Comedi we
 * have accepted the data. If this is the first data received this
 * time, we compute the start time, i.e. the timestamp for sample
 * index 0, from the current head timestamp and the amount of data
 * obtained.
 */

public int adc_data_collect(adc a) {
    uint64_t now;
    int nb;
    int ns;

    /* Retrieve any new data if possible */
    nb = comedi_get_buffer_contents(a->a_device, 0);
    now = monotonic_ns_clock();
    if(nb) {
        ns = nb / sizeof(sampl_t);
        a->a_head_time = now;
        a->a_head = a->a_tail + ns; /* Assume that nb accumulates if mark read not called */
        if( !a->a_live ) { /* Estimate the timestamp of sample index 0 */

```

Sep 22, 15 16:14

adc.c

Page 7/7

```
        a->a_start_time = a->a_head_time - ns*a->a_intersample_ns;
        a->a_live++;
    }
}
return nb;
}

/*
 * Purge data from the tail of the ring buffer if explicit data
 * lifetime management is used.
 */

public int adc_data_purge(adc a, int ns) {
    int nb = ns*sizeof(sampl_t);
    int ret;

    ret = comedi_mark_buffer_read(a->a_device, 0, nb);
    if(ret != nb)
        return -1;
    if( a->a_head >= a->a_bufsz_samples ) {
        a->a_tail += ns;
    }
    return 0;
}
```

Sep 22, 15 15:11

adc.h

Page 1/1

```
#
#ifndef _ADC_H
#define _ADC_H

#include "general.h"

/*
 * Descriptor structure for Reader ADC interface.
 */

#define NCHANNELS      8      /* Public number of channels offered */

typedef struct _adc *adc;

export adc  adc_new(strbuf);
export void adc_destroy(adc);

export int  adc_set_frequency(adc, strbuf, double *);
export int  adc_set_bufsz(adc, strbuf, int);
export int  adc_set_range(adc, strbuf, int);
export int  adc_set_device(adc, const char *);
export void adc_set_raw_mode(adc, int);

export int  adc_init(adc, strbuf);
export int  adc_start_data_transfer(adc, strbuf);
export void adc_stop_data_transfer(adc);
export void adc_setup_chunk(adc, chunk_t *);

export uint64_t adc_time_to_sample(adc, uint64_t);
export uint64_t adc_sample_to_time(adc, uint64_t);

export int  adc_ns_per_sample(adc);
export double adc_tot_frequency(adc);
export uint64_t adc_capture_start_time(adc);

export uint64_t adc_ring_head(adc);
export uint64_t adc_ring_tail(adc);

export int  adc_data_collect(adc);
export int  adc_data_purge(adc,int);

#endif /* _ADC_H */
```


Sep 12, 15 16:49

argtab-helper.h

Page 1/1

```

#
#ifndef _ARGTAB_HELPER_H
#define _ARGTAB_HELPER_H

#include "general.h"
#include "assert.h"
#include "param.h"

/*
 * Simplify definition of command line parsing tables
 *
 * The argument syntax definition, in the form of calls to argtable2
 * constructors, goes between the BEGIN_CMD_SYNTAX() and APPLY_CMD_DEFAULTS()
 * macro calls, enclosed in { } as an initialiser list (comma-separated). The
 * default assignments are placed between { } as a statement list, following
 * APPLY_CMD_DEFAULTS() and finishing with END_CMD_SYNTAX(). The result is that
 * the default assignments are placed inside the constructor built by the whole
 * macro set.
 *
 * Keeping the { } outside the macros, and following the END_CMD_SYNTAX() call
 * with a semicolon (null statement) allows emacs font-lock to keep up :-))
 */

#define BEGIN_CMD_SYNTAX(name) void **arg_make_ ## name () { void **ret, *argtable[] =
#define APPLY_CMD_DEFAULTS(name) ; \
\
if( arg_nullcheck(argtable) ) { \
    arg_freetable(argtable, sizeof(argtable)/sizeof(void *)); \
    return NULL; \
} \
\
ret = malloc(sizeof(argtable)); \
if( !ret ) { \
    arg_freetable(argtable, sizeof(argtable)/sizeof(void *)); \
    return NULL; \
} \
do \
#define INCLUDE_PARAM_DEFAULTS(ps,nps) \
    int rv = arg_defaults_from_params(argtable, \
        sizeof(argtable)/sizeof(void *), \
        (ps), (nps)); \
    assertv(rv == 0, "Argtable has no end mark\n");
#define END_CMD_SYNTAX(name) \
\
while(0); \
\
int n = sizeof(argtable)/sizeof(void *); \
while( n-- > 0 ) { \
    ret[n] = argtable[n]; \
} \
return ret; \
}

#endif /* _ARGTAB_HELPER_H */

```

Sep 12, 15 16:54

argtab-int16.c

Page 1/2

```

#
#include "general.h"
#include <argtable2.h>
#include <stdlib.h>
#include "argtab-int16.h"

/*
 * Callback functions for this argument class -- based closely on the
 * argtable2 examples.
 */

/* Private error codes for this type */
enum {OK=0,EMINCOUNT,EMAXCOUNT,EBADVALUE};

/* Reset the parent argument count */
private void resetfn(struct arg_int16 *parent)
{
    parent->count=0;
}

/* Read a value from an argument string */
private int scanfn(struct arg_int16 *parent, const char *argval)
{
    long long int val;
    char *left;

    if (parent->count == parent->hdr.maxcount)
    {
        /* maximum number of arguments exceeded */
        return EMAXCOUNT;
    }
    if (!argval)
    {
        /* an argument with no argument value was given. */
        /* This happens when an optional argument value was invoked. */
        /* leave parent argument value unaltered but still count the argument. */
        parent->count++;
        return 0;
    }

    /* Try to convert the argument string */
    val = strtoll(argval, &left, 0);

    if (*left == '\0') {
        /* success; value was scanned ok, and it is within our desired range. */
        parent->data[parent->count++] = val;
        return OK;
    }

    /* failure; command line string was not a valid integer */
    return EBADVALUE;
}

/* Check for presence of required arguments */
private int checkfn(struct arg_int16 *parent)
{
    /* return EMINCOUNT if the minimum argument count has not been satisfied */
    if( parent->count < parent->hdr.mincount )
        return EMINCOUNT;
    else
        return OK;
}

/* Error handler function */
private void errorfn(struct arg_int16 *parent, FILE *fp, int errorcode, const char *argval, const char *progname)
{
    const char *shortopts = parent->hdr.shortopts;
    const char *longopts  = parent->hdr.longopts;
    const char *datatype  = parent->hdr.datatype;

    /* make argval NULL safe */
    argval = argval ? argval : "";

```

Sep 12, 15 16:54

argtab-int16.c

Page 2/2

```

fprintf(fp, "%s: ", progname);
switch(errorcode)
{
    case EMINCOUNT:
        /* We expected more arg_int16 arguments than we received. */
        fputs("missing option \"", fp);
        arg_print_option(fp, shortopts, longopts, datatype, "\\n");
        break;

    case EMAXCOUNT:
        /* We received more arg_int16 arguments than we expected. */
        fputs("excess option \"", fp);
        arg_print_option(fp, shortopts, longopts, argval, "\\n");
        break;

    case EBADVALUE:
        /* An arg_int16 option was given with an invalid value */
        fprintf(fp, "invalid argument \"%s\" to option ", argval);
        arg_print_option(fp, shortopts, longopts, datatype, "\\n");
        break;
}
}

/* Generic constructor for an arg_int16 structure */
struct arg_int16* arg_int16n(const char* shortopts, const char* longopts,
                           const char *datatype,
                           int mincount, int maxcount, const char *glossary) {
    int bytes;
    struct arg_int16 *ret;

    bytes = sizeof(struct arg_int16) + maxcount*sizeof(uint16_t);
    ret = (struct arg_int16 *)calloc(1, bytes);
    if( ret ) {
        ret->hdr.flag      = ARG_HASVALUE;
        ret->hdr.shortopts = shortopts;
        ret->hdr.longopts  = longopts;
        ret->hdr.datatype  = datatype ? datatype : "<[u]int16_t>";
        ret->hdr.glossary  = glossary;
        ret->hdr.mincount  = mincount;
        ret->hdr.maxcount  = maxcount;
        ret->hdr.parent    = ret;
        ret->hdr.resetfn   = (arg_resetfn *)resetfn;
        ret->hdr.scanfn    = (arg_scanfn *)scanfn;
        ret->hdr.checkfn   = (arg_checkfn *)checkfn;
        ret->hdr.errorfn   = (arg_errorfn *)errorfn;
        ret->count = 0;
        ret->data = (int16_t *)&ret[1];
    }
    return ret;
}

/* Special case: 0 or 1 arguments */
struct arg_int16* arg_int160(const char* shortopts, const char* longopts,
                           const char *datatype, const char *glossary) {
    return arg_int16n(shortopts, longopts, datatype, 0, 1, glossary);
}

/* Special case: exactly 1 argument */
struct arg_int16* arg_int161(const char* shortopts, const char* longopts,
                           const char *datatype, const char *glossary) {
    return arg_int16n(shortopts, longopts, datatype, 1, 1, glossary);
}

```

Sep 12, 15 16:49

argtab-int16.h

Page 1/1

```
#
/*
 * Public definitions for a int16_t argument consistent with argtable2.
 */

#ifndef _ARGTAB_INT16_H
#define _ARGTAB_INT16_H

#include <stdint.h>
#include <argtable2.h>

struct arg_16b
{
    struct arg_hdr hdr;      /* The mandatory argtable header struct */
    int count;               /* Number of matching command line arguments found */
    int16_t *data;           /* Array of matching command line argument data */
};

struct arg_16b* arg_16b0(const char* shortopts, const char* longopts, const char *datatype,
                        const char *glossary);

struct arg_16b* arg_16b1(const char* shortopts, const char* longopts, const char *datatype,
                        const char *glossary);

struct arg_16b* arg_16bn(const char* shortopts, const char* longopts, const char *datatype,
                        int mincount, int maxcount, const char *glossary);

#endif /* _ARGTAB_INT16_H */
```

Sep 12, 15 16:55

argtab-int64.c

Page 1/2

```

#
#include "general.h"
#include <argtable2.h>
#include <stdlib.h>
#include "argtab-int64.h"

/*
 * Callback functions for this argument class -- based closely on the
 * argtable2 examples.
 */

/* Private error codes for this type */
enum {OK=0,EMINCOUNT,EMAXCOUNT,EBADVALUE};

/* Reset the parent argument count */
private void resetfn(struct arg_int64 *parent)
{
    parent->count=0;
}

/* Read a value from an argument string */
private int scanfn(struct arg_int64 *parent, const char *argval)
{
    long long int val;
    char *left;

    if (parent->count == parent->hdr.maxcount)
    {
        /* maximum number of arguments exceeded */
        return EMAXCOUNT;
    }
    if (!argval)
    {
        /* an argument with no argument value was given. */
        /* This happens when an optional argument value was invoked. */
        /* leave parent argument value unaltered but still count the argument. */
        parent->count++;
        return 0;
    }

    /* Try to convert the argument string */
    val = strtoll(argval, &left, 0);

    if (*left == '\0') {
        /* success; value was scanned ok, and it is within our desired range. */
        parent->data[parent->count++] = val;
        return OK;
    }

    /* failure; command line string was not a valid integer */
    return EBADVALUE;
}

/* Check for presence of required arguments */
private int checkfn(struct arg_int64 *parent)
{
    /* return EMINCOUNT if the minimum argument count has not been satisfied */
    if( parent->count < parent->hdr.mincount )
        return EMINCOUNT;
    else
        return OK;
}

/* Error handler function */
private void errorfn(struct arg_int64 *parent, FILE *fp, int errorcode, const char *argval, const char *progname)
{
    const char *shortopts = parent->hdr.shortopts;
    const char *longopts  = parent->hdr.longopts;
    const char *datatype  = parent->hdr.datatype;

    /* make argval NULL safe */
    argval = argval ? argval : "";

```

Sep 12, 15 16:55

argtab-int64.c

Page 2/2

```

fprintf(fp, "%s: ", progname);
switch(errorcode)
{
    case EMINCOUNT:
        /* We expected more arg_int64 arguments than we received. */
        fputs("missing option \"", fp);
        arg_print_option(fp, shortopts, longopts, datatype, "\\n");
        break;

    case EMAXCOUNT:
        /* We received more arg_int64 arguments than we expected. */
        fputs("excess option \"", fp);
        arg_print_option(fp, shortopts, longopts, argval, "\\n");
        break;

    case EBADVALUE:
        /* An arg_int64 option was given with an invalid value */
        fprintf(fp, "invalid argument \"%s\" to option ", argval);
        arg_print_option(fp, shortopts, longopts, datatype, "\\n");
        break;
}
}

/* Generic constructor for an arg_int64 structure */
struct arg_int64* arg_int64n(const char* shortopts, const char* longopts,
                           const char *datatype,
                           int mincount, int maxcount, const char *glossary) {
    int bytes;
    struct arg_int64 *ret;

    bytes = sizeof(struct arg_int64) + maxcount*sizeof(uint64_t);
    ret = (struct arg_int64 *)calloc(1, bytes);
    if( ret ) {
        ret->hdr.flag = ARG_HASVALUE;
        ret->hdr.shortopts = shortopts;
        ret->hdr.longopts = longopts;
        ret->hdr.datatype = datatype ? datatype : "<[u]int64_t>";
        ret->hdr.glossary = glossary;
        ret->hdr.mincount = mincount;
        ret->hdr.maxcount = maxcount;
        ret->hdr.parent = ret;
        ret->hdr.resetfn = (arg_resetfn *)resetfn;
        ret->hdr.scanfn = (arg_scanfn *)scanfn;
        ret->hdr.checkfn = (arg_checkfn *)checkfn;
        ret->hdr.errorfn = (arg_errorfn *)errorfn;
        ret->count = 0;
        ret->data = (int64_t *)&ret[1];
    }
    return ret;
}

/* Special case: 0 or 1 arguments */
struct arg_int64* arg_int640(const char* shortopts, const char* longopts,
                           const char *datatype, const char *glossary) {
    return arg_int64n(shortopts, longopts, datatype, 0, 1, glossary);
}

/* Special case: exactly 1 argument */
struct arg_int64* arg_int641(const char* shortopts, const char* longopts,
                           const char *datatype, const char *glossary) {
    return arg_int64n(shortopts, longopts, datatype, 1, 1, glossary);
}

```

Sep 12, 15 16:49

argtab-int64.h

Page 1/1

```
#
/*
 * Public definitions for a int64_t argument consistent with argtable2.
 */

#ifndef _ARGTAB_INT64_H
#define _ARGTAB_INT64_H

#include <stdint.h>
#include <argtable2.h>

struct arg_64b
{
    struct arg_hdr hdr;      /* The mandatory argtable header struct */
    int count;               /* Number of matching command line arguments found */
    int64_t *data;           /* Array of matching command line argument data */
};

struct arg_64b* arg_64b0(const char* shortopts, const char* longopts, const char *datatype,
                        const char *glossary);

struct arg_64b* arg_64b1(const char* shortopts, const char* longopts, const char *datatype,
                        const char *glossary);

struct arg_64b* arg_64bn(const char* shortopts, const char* longopts, const char *datatype,
                        int mincount, int maxcount, const char *glossary);

#endif /* _ARGTAB_INT64_H */
```

Sep 12, 15 16:49

assert.h

Page 1/1

```
#
#ifndef _LOCAL_ASSERT_H
#define _LOCAL_ASSERT_H

/*
 * Local version of assert, bit more informative than system version
 */

#ifdef USE_SYSTEM_ASSERT
#include <assert.h>
#define assertv(cond, ...) assert(cond)
#else
#include <stdio.h>
#include <stdlib.h>

#define assertv(cond,fmt, ...) do { \
    if(!(cond)) { \
        fprintf(stderr, "FAILED ASSERTION -- %s:%d %s %s\n" fmt, __FILE__, __LINE__, __FUNCTION__, "' ' #cond '" , ## __VA_ARGS__ ); \
        abort(); \
    } } while(0)

#endif

#endif /* _LOCAL_ASSERT_H */
```


Sep 22, 15 15:23

chunk.c

Page 1/4

```
#
#include "general.h"
#include <stdlib.h>
#include <stdint.h>
#include <errno.h>
#include <comedilib.h>

#include "assert.h"
#include "queue.h"
#include "mman.h"
#include "strbuf.h"
#include "chunk.h"
#include "writer.h"

struct _frame {
    queue f_Q;
    block f_map;
};

/*
 * Set up the mmap frames for data transfer to snapshot files.
 */

private int    nframes; /* The number of simultaneous mmap frames */
private frame *framelist; /* The list of mmap frame descriptors */
private int    n_frame_Q = 0;

private QUEUE_HEADER(frameQ);

public int init_frame_system(strbuf e, int nfr, int ram, int chunk) {

    framelist = (frame *)calloc(nfr, sizeof(frame));
    if( framelist ) {
        void *map = mmap_locate(ram*1024*1024, 0);
        int    n;

        if(map == NULL) {
            strbuf_appendf(e, "Cannot mmap %d MiB of locked transfer RAM: %m", ram);
            free((void *) framelist );
            return -1;
        }
        for(n=0; n<nfr; n++) { /* Initialise the frame memory pointers, leave sizes as 0 */
            framelist[n].f_map.b_data = map;
            map += chunk;
            init_queue(&framelist[n].f_Q);
            queue_ins_before(&frameQ, &framelist[n].f_Q);
            n_frame_Q++;
        }
    }
    else {
        strbuf_appendf(e, "Cannot allocate frame list memory for %d frames: %m", nfr);
        return -1;
    }
    nframes = nfr;
    return 0;
}

/*
 * Scan the frame list and pull any free frame descriptors into the free queue.
 *
 * A descriptor is free if its byte count is zero, and it is not in
 * the free queue if its queue structure is a singleton.
 */

private void scan_framelist() {
    int    n;
    frame *f;

    for(n=0; f=framelist; n<nframes; n++, f++) {
        if(f->f_map.b_bytes) /* If non-zero, it's in use */
            continue;
        if( !queue_singleton(&f->f_Q) )
            continue;
    }
}
```

Sep 22, 15 15:23

chunk.c

Page 2/4

```

    queue_ins_before(&frameQ, &f->f_Q);
    n_frame_Q++;
}
}

/*
 * Allocate a frame descriptor.
 */

private frame *alloc_frame() {
    frame *f;

    if( !n_frame_Q ) {
        assertv(queue_singleton(&frameQ), "Frame queue count is zero for non-empty queue\n");
        scan_framelist();
    }
    if( !n_frame_Q ) {
        errno = EBUSY;
        return NULL;
    }
    f = (frame *)de_queue(queue_next(&frameQ));
    assertv(f != NULL, "Frame queue count %d but queue is empty\n", n_frame_Q);
    n_frame_Q--;
    f->f_map.b_bytes = 1;      /* In-use; real size is filled in by caller */
    return f;
}

/*
 * Release a frame descriptor. N.B. this is done in the TIDY thread,
 * so must be atomic: the frame is released by setting the bytes value
 * to zero. The munmap() here complements the mmap call in the chunk
 * mapper below.
 */

public void release_frame(frame *f) {
    if(f->f_map.b_data == NULL)
        return;
    munmap(f->f_map.b_data, f->f_map.b_bytes);
    f->f_map.b_bytes = 0;
}

/*
 * Report the index of a frame pointer in the table.
 */

public int frame_nr(frame *f) {
    return f - framelist;
}

/*
 * Functions for dealing with transfer chunk descriptors.
 */

private uint16_t chunk_counter;

#define N_CHUNK_ALLOC    (4096/sizeof(chunk_t))

private QUEUE_HEADER(chunkQ);
private int N_in_chunkQ = 0;

/*
 * Allocate n new chunk descriptors, chained using the WRITER queue descriptor
 */

public chunk_t *alloc_chunk(int nr) {
    queue *ret;

    if( N_in_chunkQ < nr ) {      /* The queue doesn't have enough */
        int n;

        for(n=0; n<N_CHUNK_ALLOC; n++) {
            queue *q = (queue *)calloc(1, sizeof(chunk_t));

            if( !q ) {              /* Allocation failed */
                if( N_in_chunkQ >= nr )

```

Sep 22, 15 15:23

chunk.c

Page 3/4

```

        break;                /* But we have enough now anyway */
        return NULL;
    }
    init_queue(q);
    queue_ins_after(&chunkQ, q);
    N_in_chunkQ++;
}

ret = de_queue(queue_next(&chunkQ));
chunk_t *c = qp2chunk(ret);
init_queue(&c->c_rQ);
c->c_name = ++chunk_counter;

while(--nr > 0) {             /* Collect enough to satisfy request */
    chunk_t *c = qp2chunk(de_queue(queue_next(&chunkQ)));

    init_queue(&c->c_wQ);       /* Redundant... */
    init_queue(&c->c_rQ);
    c->c_name = ++chunk_counter;
    queue_ins_before(ret, chunk2qp(c));
}

return c;
}

/*
 * Finished with chunk descriptors chained using the writer queue descriptor.
 * Assume the reader queue descriptor is detached.
 */

public void release_chunk(chunk_t *c) {
    queue *q = chunk2qp(c);
    queue *p;

    while( (p = de_queue(queue_next(q))) != NULL ) {
        init_queue(p);
        queue_ins_before(&chunkQ, p);
        N_in_chunkQ++;
    }
    init_queue(q);
    queue_ins_before(&chunkQ, q);
    N_in_chunkQ++;
}

/*
 * Find a frame for a chunk and map the chunk into memory. This may
 * take arbitrary time since this is where Linux has to find us new
 * pages. This code runs in the WRITER thread.
 */

public int map_chunk_to_frame(chunk_t *c) {
    frame *fp = alloc_frame();
    void *map;

    if(fp == NULL) {           /* This is not an error: there may be no frames available for transient reasons */
        return -1;
    }

    fp->f_map.b_bytes = c->c_samples*sizeof(sampl_t);
    /* Would really like to do WRONLY here, but I *think* that will break */
    map = mmap_and_lock_fixed(c->c_fd, c->c_offset, fp->f_map.b_bytes, PROT_RDWR|PREFAULT_RDWR|MAL_LOCKED, fp->f_map.b_data);
    if(map != fp->f_map.b_data) { /* A (fatal) mapping error occurred... */
        strbuf_appendf(c->c_error, "Unable to map chunk c:%04hx to frame %d: %m", c->c_name, frame_nr(fp));
        c->c_status = SNAPSHOT_ERROR;
        return -1;
    }
    c->c_frame = fp;           /* Succeeded, chunk now has a mapped frame */
    return 0;
}

/*
 * Copy the data for a chunk from the ring buffer into the frame.
 * Apply the appropriate ADC conversion.
 */

```

Sep 22, 15 15:23

chunk.c

Page 4/4

```

public void copy_chunk_data(chunk_t *c) {
    convertfn fn = c->c_convert;

    (*fn)((sampl_t *)c->c_frame->f_map.b_data, (sampl_t *)c->c_ring, c->c_samples);
    c->c_status = SNAPSHOT_WRITTEN;
}

/*
 * Generate a debugging line for a chunk descriptor. Put it in the buffer buf.
 * Return the actual size, no greater than the space available.
 */

#define qp2cname(p)      (qp2chunk(p)->c_name)
#define rq2cname(p)      (rq2chunk(p)->c_name)

public int debug_chunk(char buf[], int space, chunk_t *c) {
    import const char *snapshot_status(int);
    import uint16_t    snapfile_name(snapfile_t *);
    int used;

    used = snprintf(buf, space,
        "chunk c:%04hx at %p"
        "wQ[c:%04hx,c:%04hx] "
        "rQ[c:%04hx,c:%04hx] "
        "RG %p FR %p PF f:%04hx status %s "
        "S:%08lx F:%016llx L:%016llx\n",
        c->c_name, c,
        qp2cname(queue_prev(&c->c_wQ)), qp2cname(queue_next(&c->c_wQ)),
        rq2cname(queue_prev(&c->c_rQ)), rq2cname(queue_next(&c->c_rQ)),
        c->c_ring, c->c_frame, snapfile_name(c->c_parent), snapshot_status(c->c_status),
        c->c_samples, c->c_first, c->c_last
    );
    if(used >= space)
        used = space;
    return used;
}

```

Sep 21, 15 16:11

chunk.h

Page 1/1

```

#
#ifndef _CHUNK_H
#define _CHUNK_H

#include "general.h"
#include <comedilib.h>
#include "lut.h"

/* Structure for a memory block */

typedef struct {
    void *b_data;
    int   b_bytes;
}
    block;

typedef struct _frame frame;
typedef struct _sfile snapfile_t;

#include "queue.h"

typedef struct {
    queue      c_Q[2];      /* Q header for READER capture queue and WRITER file chunk list*/
#define c_wQ c_Q[0]        /* Chunk Q linkage associated with the file */
#define c_rQ c_Q[1]        /* Chunk Q linkage associated with the data flow */
    frame      *c_frame;    /* Mmap'd file buffer for this chunk */
    strbuf     c_error;     /* Error buffer, for error messages (copy from snapshot_t origin) */
    snapfile_t *c_parent;   /* Chunk belongs to this file */
    uint64_t    c_first;    /* First sample of this chunk */
    uint64_t    c_last;     /* First sample beyond this chunk */
    int16_t     c_ring;     /* Ring buffer start for this chunk */
    convertfn   c_convert;  /* Function to copy samples into frame with conversion */
    uint32_t    c_samples;  /* Number of samples to copy */
    uint32_t    c_offset;   /* File offset for this chunk */
    int         c_status;   /* Status of this capture chunk */
    int         c_fd;       /* File descriptor for this chunk */
    uint16_t    c_name;     /* Unique name for this chunk */
}
    chunk_t;

#define qp2chunk(q)      ((chunk_t *) (q))
#define chunk2qp(c)      (&(c)->c_Q[0])

#define chunk2rq(c)      (&(c)->c_rQ)
#define rq2chunk(q)      ((chunk_t *) &((q)[-1]))

export chunk_t *alloc_chunk(int);
export void release_chunk(chunk_t *);
export int map_chunk_to_frame(chunk_t *);

export int debug_chunk(char [], int, chunk_t *);

export void release_frame(frame *);

#endif /* _CHUNK_H */

```

Sep 12, 15 17:20

general.h

Page 1/1

```
#
#ifndef _GENERAL_H
/*
 * Macro definitions to make static and extern more explicit.
 */
#define public
#define import  extern
#define export  extern
#define private static
#define persist static
#endif /* _GENERAL_H */
```

Sep 05, 15 14:07

grab.c

Page 1/5

```

#
/*
 * Program to grab data from USB DUXfast via Comedi.
 *
 * Arguments:
 * --verbose/-v      Increase reporting level
 * --freq/-f         Sampling frequency in [Hz], default 2.5 [MHz]
 * --range/-r        ADC range 'hi' (750 mVpk) or 'lo' (500 mVpk)
 * --raw            ADC output as raw data
 * --device/-d       Comedi device to use, default /dev/comedi0
 * --bufsz/-B        Comedi buffer size to request [MiB], default 40 [MiB]
 * --help/-h         Print usage message
 * --version         Print program version
 */

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <stdint.h>
#include "assert.h"
#include <argtable2.h>
#include <regex.h>
#include <comedi.h>
#include <comedilib.h>
#include <fcntl.h>
#include <errno.h>
#include <string.h>

#include "argtab.h"
#include "mman.h"
#include "lut.h"

#define N_CHANS      16
#define BUFSZ        4096
#define BUFSPSZ      (BUFSZ/sizeof(sampl_t))

char read_buf[BUFSZ];

#define COMEDI_DEVICE    "/dev/comedi0"

#define COMEDIBUFFERSIZE (*40)
#define COMEDIBUFFERSPLS (COMEDIBUFFERSIZE/sizeof(sampl_t))

#define PROGRAM_VERSION    "2.0"
#define VERSION_VERBOSE_BANNER "MCLURS ADC toolset...\n"

/* Standard arguments + flags */
int  verbose = 0;
char *program = NULL;

/* Command line syntax options */

struct arg_lit *h1, *vn1, *v1;
struct arg_end *e1;

BEGIN_CMD_SYNTAX(help) {
    v1 = arg_litn("v", "verbose", 0, 2, "Increase verbosity"),
    h1 = arg_lit0("h", "help", "Print usage help message"),
    vn1 = arg_lit0(NULL, "version", "Print program version string"),
    e1 = arg_end(20)
} APPLY_CMD_DEFAULTS(help) {
    /* No defaults to apply here */
} END_CMD_SYNTAX(help)

struct arg_lit *v2, *rw2;
struct arg_int *b2;
struct arg_dbl *f2;
struct arg_str *d2;
struct arg_rex *rn2;
struct arg_end *e2;

BEGIN_CMD_SYNTAX(main) {
    v2 = arg_litn("v", "verbose", 0, 2, "Increase verbosity"),
    f2 = arg_dbl0("f", "freq", "<real>", "Sampling frequency [Hz], default 2.5 [MHz]"),

```

Sep 05, 15 14:07

grab.c

Page 2/5

```

b2 = arg_int0("B", "bufsz", "<int>", "Comedi buffer size [MiB], default 40 [MiB]"),
d2 = arg_str0("d", "device", "<path>", "Comedi device to open, default /dev/comedi0"),
rn2 = arg_rex0("r", "range", "hi|lo", NULL, REG_EXTENDED, "Specify range in {hi, lo}, default hi"),
rw2 = arg_lit0(NULL, "raw", "Emit raw ADC sample values"),
e2 = arg_end(20)
} APPLY_CMD_DEFAULTS(main) {
    *f2->dval = 2.5e6; /* Default frequency 2.5 [MHz] */
    *b2->ival = 40; /* Default buffer size 40 [MiB] */
    *d2->sval = COMEDI_DEVICE; /* Default device for Comedi */
    *rn2->sval = "hi"; /* Default ADC range (hi) */
} END_CMD_SYNTAX(main);

/* Standard help routines: display the version banner */
void print_version(FILE *fp, int verbosity) {
    fprintf(fp, "%sVn.%s\n", program, PROGRAM_VERSION);
    if(verbosity > 0) { /* Verbose requested... */
        fprintf(fp, VERSION_VERBOSE_BANNER);
    }
}

/* Standard help routines: display the usage summary for a syntax */
void print_usage(FILE *fp, void **argtable, int verbosity, char *program) {
    if( !verbosity ) {
        fprintf(fp, "Usage: %s ", program);
        arg_print_syntax(fp, argtable, "\n");
        return;
    }
    if( verbosity ) {
        char *suffix = verbosity>1? "\n\n" : "\n";
        fprintf(fp, "Usage: %s ", program);
        arg_print_syntaxv(fp, argtable, suffix);
        if( verbosity > 1 )
            arg_print_glossary(fp, argtable, "%-25s %s\n");
    }
}

/*
 * The main() entry point.
 */

int main(int argc, char *argv[]) {
    float      sr_total;
    int        bufsz;
    char       *device;
    int        range = 0; /* Default range is +/- 750mV */

    int        buf_samples;
    unsigned int convert_arg;
    comedi_t   *dev;
    int        errs, ret, i;
    unsigned int chanlist[N_CHANS];
    void       *map;
    sampl_t    *start;
    uint64_t    head, tail;
    int        data_coming;
    void        (*convert)(sampl_t *, sampl_t *, int);

    program = argv[0];

    /* Create and parse the command lines */
    void **cmd_help = arg_make_help();
    void **cmd_main = arg_make_main();

    /* Try first syntax */
    int err_help = arg_parse(argc, argv, cmd_help);
    if( !err_help ) { /* Assume this was the desired command syntax */
        if(vn1->count)
            print_version(stdout, v1->count);
        if(hl->count || !vn1->count) {
            print_usage(stdout, cmd_help, v1->count>0, program);
            print_usage(stdout, cmd_main, v1->count, program);
        }
        exit(0);
    }
}

```


Sep 05, 15 14:07

grab.c

Page 3/5

```

/* Try second syntax */
int err_main = arg_parse(argc, argv, cmd_main);
if( err_main ) {
    /* This is the default desired syntax; give full usage */
    arg_print_errors(stderr, e2, program);
    print_usage(stderr, cmd_help, v2->count>0, program);
    print_usage(stderr, cmd_main, v2->count, program);
    exit(1);
}

/* The second syntax was correctly parsed, so retrieve the important values from the table */
errs = 0;

/* Deal with the sampling frequency */
sr_total = f2->dval[0];
if(sr_total < 5e4 || sr_total > 3e6) {
    fprintf(stderr, "%s: Error -- total sample rate %g [Hz] out of sensible range (50 [kHz] to 3 [MHz])\n", program, sr_total);
    errs++;
}
convert_arg = (unsigned int) 1e9 / sr_total;

/* Deal with the requested buffer size */
bufsz = b2->ival[0];
if(bufsz < 8 || bufsz > 256) {
    fprintf(stderr, "%s: Error -- requested buffer size %d [MiB] out of sensible range (8 to 256 [MiB])\n", program, bufsz);
    errs++;
}
bufsz *= 1048576;
buf_samples = bufsz / sizeof(sampl_t);

/* Deal with the Comedi device */
device = (char *) d2->sval[0];
if( !(dev = comedi_open(device)) ) {
    fprintf(stderr, "%s: Error -- cannot open %s:\n", program, device, comedi_strerror(comedi_errno()));
    errs++;
}

/* Deal with specification of range and raw */
range = !strcmp(rn2->sval[0], "hi") ? 0 : 1;
if(rw2->count) {
    /* Requested raw */
    convert = convert_raw_raw;
} else {
    /* Convert from specified range */
    convert = range? convert_raw_500mV : convert_raw_750mV;
}

/* Record desired verbosity */
verbose = v2->count;

/* All finished with the argument syntax tables */
arg_free(cmd_main);
arg_free(cmd_help);

/* Exit 2 if argument errors */
if(errs) {
    exit(2);
}

fprintf(stderr, "%s %s\n", program, PROGRAM_VERSION);
fprintf(stderr, "Total sample rate requested = %g [Hz]\n", sr_total);
fprintf(stderr, "Using ADC range +/- %s [mV] full-scale\n", range? "500" : "750");

comedi_cmd *cmd = (comedi_cmd *) calloc(1, sizeof(comedi_cmd));
if( !cmd ) {
    fprintf(stderr, "%s: Error -- failed to get memory for Comedi command structure: %s\n", program, strerror(errno));
    exit(3);
}

ret = comedi_get_max_buffer_size(dev, 0);

if(ret < bufsz) {
    if(comedi_set_max_buffer_size(dev, 0, bufsz) < 0 ) {
        fprintf(stderr, "%s: Error -- failed to set %s max buffer size to %d bytes: %s\n", program, device, bufsz, comedi_strerror(comedi_errno()));
        exit(3);
    }
}
else {

```

Sep 05, 15 14:07

grab.c

Page 4/5

```

    if(verbose)
        fprintf(stderr, "%s: Comedi maximum buffer size requested %d, actual %d\n", program, bufsz, ret);
}

if(comedi_set_buffer_size(dev, 0, bufsz) < 0) {
    fprintf(stderr, "%s: Error -- failed to set %s buffer size to %d bytes: %s\n", program, device, bufsz, comedi_strerror(comedi_errno()));
    exit(3);
}

for(i=0; i<N_CHANS; i++)
    chanlist[i] = CR_PACK(i, range, AREF_GROUND);

/* Print numbers for clipped inputs */
comedi_set_global_oor_behavior(COMEDI_OOR_NUMBER);

/* get the correct command structure to run usbdxfast */
if((ret = comedi_get_cmd_generic_timed(dev, 0, cmd, N_CHANS, 0)) < 0) {
    fprintf(stderr, "%s: Error -- comedi_get_cmd_generic_timed failed for %s: %s\n", program, device, comedi_strerror(comedi_errno()));
    exit(3);
}

populate_conversion_luts();

/* adjust some cmd parameters */
cmd->chanlist = chanlist;
cmd->stop_src = TRIG_NONE;
cmd->stop_arg = 0;
convert_arg = (unsigned int) 1e9 / sr_total;
cmd->convert_arg = convert_arg;

/* call test twice because different things are tested?
 * if tests are successful run sampling command */
if( (ret = comedi_command_test(dev, cmd)) != 0 && verbose > 1 ) {
    fprintf(stderr, "First test, err: %s:", comedi_strerror(comedi_errno()));
    fprintf(stderr, " cmd->convert_arg = %d\n", cmd->convert_arg);
}
if( (ret = comedi_command_test(dev, cmd)) != 0 && verbose > 1 ) {
    fprintf(stderr, "Second test, err: %s:", comedi_strerror(comedi_errno()));
    fprintf(stderr, " cmd->convert_arg = %d\n", cmd->convert_arg);
}

map = mmap_and_lock(comedi_fileno(dev), 0, bufsz, PROT_RDONLY|PREFAULT_RDONLY|MAL_DOUBLED|MAL_LOCKED);
if(map == NULL) {
    fprintf(stderr, "%s: Error -- failed to map Comedi buffer to RAM: %s\n", program, strerror(errno));
    exit(3);
}

if(verbose)
    fprintf(stderr, "%s: Comedi buffer (size %u bytes) mapped at 0x%p\n", program, bufsz, start);

if( (ret = comedi_command(dev, cmd)) < 0 ) {
    fprintf(stderr, "%s: Error -- Comedi command returns %d: %s\n", program, ret, comedi_strerror(comedi_errno()));
    exit(4);
}

if(verbose > 1) {
    fprintf(stderr, "comedi_command returns %d\n", ret);
    fprintf(stderr, "stop src = %d\n", cmd->stop_src);
    fprintf(stderr, "stop arg = %d\n", cmd->stop_arg);
    fprintf(stderr, "convert arg = %d\n", cmd->convert_arg);
}

if(verbose)
    fprintf(stderr, "%s: Total sample rate allocated = %g Hz\n", program, 1e9 / cmd->convert_arg);

head=tail=0;
data_coming = 1000; /* Is data arriving? After this many pauses with no data, exit... */
while( 1 ) {
    int nb = comedi_get_buffer_contents(dev, 0); /* Find out how many new bytes there are */
    sampl_t *back = &start[ tail % buf_samples ];
    ret = 0;

    if(nb <= 0) {
        usleep(10000);
        if( --data_coming == 0 ) break;
    }
}

```

Sep 05, 15 14:07

grab.c

Page 5/5

```
    }
    continue;
}

data_coming = 100;      /* Some data has come, use a smaller value henceforth */
head += nb/sizeof(sampl_t); /* This many new samples have arrived */
nb = head - tail;      /* And this is how many remain to process */

while( nb >= BUFSPSZ ) {
    /* Convert and dump a buffer-full to stdout, repeat while possible */
    (*convert)((sampl_t *)read_buf, back, BUFSPSZ);
    ret = comedi_mark_buffer_read(dev, 0, BUFSZ);
    if(ret < 0) {
        fprintf(stderr, "%s: Error -- comedi_mark_buffer_read during loop: %s\n", program, comedi_strerror(comedi_errno()));
        break;
    }
    fwrite(read_buf, sizeof(sampl_t), BUFSPSZ, stdout);
    back += BUFSPSZ;
    tail += BUFSPSZ;
    nb -= BUFSPSZ;
}
}

fprintf(stderr, "%s: Error? -- Comedi data flow interrupted for more than 1 second\n", program);

comedi_cancel(dev, 0);
exit(0);
}
```

Sep 12, 15 16:58

lut.c

Page 1/2

```

#
#include "general.h"

#include "assert.h"
#include <string.h>
#include <comedi.h>

#include "lut.h"

/*
 * Construct a look up table to map the USBDUXfast ADC outputs into 1V
 * pk s16 representation.
 *
 * The raw data goes from 000 to FFF, or-ed with 0x1000 if an overflow occurs.
 * The table is indexed with the raw data value to generate the s16 value.
 *
 * In the case of overflow, the value converted is 1 more than the
 * maximum raw value returnable by the ADC.
 *
 * There is one table for when the USBduxFAST is in 0.5V scale mode
 * and one for 0.75V scale.
 */

#define CONVERT_BY_LUT          /* Use a lookup table to do the complete conversion from raw to normalised */

#define ADC_BITS                12

#define USBDUXFAST_OOR   (1<<ADC_BITS)
#define USBDUXFAST_SIGN (1 << (ADC_BITS-1))
#define USBDUXRAW_MIN    0
#define USBDUXRAW_MAX    (~(0)<<ADC_BITS))

#define RAW_500mV_TO_OUT_500mV(raw)      (((short)((raw)<<4) ^ 0x8000)) >> 1) /* Shift up and correct sign bit, then arithmetic shift back 1 */
#define OUT_500mV_TO_OUT_750mV(raw)      ((raw)+(short)((raw) >> 1)) /* Add 0.5 times value you first thought of... */

#define USBDUXFAST_OOR_POS_500mV      (RAW_500mV_TO_OUT_500mV(USBDEXRAW_MAX)+1)
#define USBDUXFAST_OOR_NEG_500mV      (RAW_500mV_TO_OUT_500mV(USBDEXRAW_MIN)-1)
#define USBDUXFAST_OOR_POS_750mV      (OUT_500mV_TO_OUT_750mV(RAW_500mV_TO_OUT_500mV(USBDEXRAW_MAX))+1)
#define USBDUXFAST_OOR_NEG_750mV      (OUT_500mV_TO_OUT_750mV(RAW_500mV_TO_OUT_500mV(USBDEXRAW_MIN))-1)

/* Define the look-up tables for the conversion */
/* Using LUT only doubles the table size (but probably saves some time) */
#ifndef CONVERT_BY_LUT
#define TABLE_SIZE      (2*(1<<ADC_BITS))
#else
#define TABLE_SIZE      (1<<ADC_BITS)
#endif

private sampl_t lut_raw_to_1Vpk_500mV[TABLE_SIZE];
private sampl_t lut_raw_to_1Vpk_750mV[TABLE_SIZE];

private int lut_not_ready = 1;

public void populate_conversion_luts() {
    short raw;

    assertv(sizeof(sampl_t) == 2, "sizeof(sampl_t) is %d not 2\n", sizeof(sampl_t)); /* Check type definitions on this architecture */
    assertv(RAW_500mV_TO_OUT_500mV(USBDEXRAW_MAX) > 0, "ADC mapped max not positive\n"); /* Should work if sampl_t is signed short */
    assertv(RAW_500mV_TO_OUT_500mV(USBDEXRAW_MIN) < 0, "ADC mapped min not negative\n");

    if( !lut_not_ready ) /* i.e. the tables are already ready */
        return;

    for(raw=0; raw<=0xFFF; raw++) {
        short conv = RAW_500mV_TO_OUT_500mV(raw);

        lut_raw_to_1Vpk_500mV[raw] = conv; /* Raw value maps to itself x 8 with sign corrected */
        lut_raw_to_1Vpk_750mV[raw] = OUT_500mV_TO_OUT_750mV(conv); /* Values in 0.75pk range are scaled by 1.5 */
    }
    #ifndef CONVERT_BY_LUT
        lut_raw_to_1Vpk_500mV[raw+0x1000] = (raw&0x800)? USBDUXFAST_OOR_POS_500mV : USBDUXFAST_OOR_NEG_500mV;
        lut_raw_to_1Vpk_750mV[raw+0x1000] = (raw&0x800)? USBDUXFAST_OOR_POS_750mV : USBDUXFAST_OOR_NEG_750mV;
    #endif
}
lut_not_ready = 0; /* The tables are ready now... */

```

Sep 12, 15 16:58

lut.c

Page 2/2

```
}

public void convert_raw_500mV(sampl_t *dst, sampl_t *src, int nsamples) {
    if(lut_not_ready)
        populate_conversion_luts();

    while(nsamples-- > 0) {
#ifdef CONVERT_BY_LUT
        *dst++ = lut_raw_to_1Vpk_500mV[*src++ & (USBDEXFAST_OOR | USBDEXRAW_MAX)];
#else
        sampl_t s = *src++ & (USBDEXFAST_OOR | USBDEXRAW_MAX);
        if(s&USBDEXFAST_OOR) {
            *dst++ = (s&0x800)? USBDEXFAST_OOR_POS_500mV : USBDEXFAST_OOR_NEG_500mV;
            continue;
        }
        else
            *dst++ = lut_raw_to_1Vpk_500mV[*src++];
#endif
    }
}

public void convert_raw_750mV(sampl_t *dst, sampl_t *src, int nsamples) {
    if(lut_not_ready)
        populate_conversion_luts();

    while(nsamples-- > 0) {
#ifdef CONVERT_BY_LUT
        *dst++ = lut_raw_to_1Vpk_750mV[*src++ & (USBDEXFAST_OOR | USBDEXRAW_MAX)];
#else
        sampl_t s = *src++ & (USBDEXFAST_OOR | USBDEXRAW_MAX);
        if(s&USBDEXFAST_OOR) {
            *dst++ = (s&0x800)? USBDEXFAST_OOR_POS_750mV : USBDEXFAST_OOR_NEG_750mV;
            continue;
        }
        else
            *dst++ = lut_raw_to_1Vpk_500mV[*src++];
#endif
    }
}

public void convert_raw_raw(sampl_t *dst, sampl_t *src, int nsamples) {
    if(dst == src)
        return;
    memcpy(dst, src, nsamples*sizeof(sampl_t));
}
```

Sep 16, 15 15:23

lut.h

Page 1/1

```
#
#ifndef _LUT_H
#define _LUT_H

#include "general.h"

export void populate_conversion_luts();

export void convert_raw_500mV(sampl_t *, sampl_t *, int);
export void convert_raw_750mV(sampl_t *, sampl_t *, int);
export void convert_raw_raw(sampl_t *, sampl_t *, int);

typedef void (*convertfn)(sampl_t *, sampl_t *, int);

#endif /* _LUT_H */
```

Sep 12, 15 17:20

mman.c

Page 1/2

```

#
#include "general.h"

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <assert.h>
#include <sys/mman.h>
#include "mman.h"

/*
 * Useful utility function to ensure pages are pre-faulted.
 */

public void prefault_pages(void *p, int n, int w) {
    int ret = 0;

    while( n-- > 0 ) {
        if( (w&PREFFAULT_RDONLY) )           /* Read page */
            ret = *(int *)p;
        if( (w&PREFFAULT_WRONLY) )           /* Write page */
            *(int *)p = ret;
        p += sysconf(_SC_PAGESIZE);
    }
}

/*
 * Locate a region of memory where one could map a file of size size.
 */

public void *mmap_locate(size_t length, int flags) {
    void *map;

    if( flags & MAL_DOUBLED ) length *= 2;

    map = mmap(NULL, length, PROT_NONE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0);
    if(map == NULL || map == (void *)-1)
        return NULL;

    return map;
}

/*
 * Map and lock a region of a file into memory at given fixed address.
 */

public void *mmap_and_lock_fixed(int fd, off_t offset, size_t length, int flags, void *fixed) {
    void *map;
    int mflags = 0;

    if( flags&PROT_RDONLY )
        mflags |= PROT_READ;
    if( flags&PROT_WRONLY )
        mflags |= PROT_WRITE;

    if( !mflags )
        mflags = PROT_NONE;

    map = mmap(fixed, length, mflags, MAP_SHARED, fd, offset);
    if(map == NULL || map == (void *)-1)
        return NULL;

    if( (flags&MAL_LOCKED) && mlock(map, length) < 0 ) {
        munmap(map, length);
        return NULL;
    }

    if( flags & PREFFAULT_RDWR )
        prefault_pages(map, length / sysconf(_SC_PAGESIZE), (flags & PREFFAULT_RDWR));

    return map;
}

/*

```

Sep 12, 15 17:20

mman.c

Page 2/2

```
* Map and lock a region of a file into memory, don't care where...
*/

public void *mmap_and_lock(int fd, off_t offset, size_t length, int flags) {
    void *map;
    void *ms;

    map = mmap_locate(length, flags);
    if( !map )
        return NULL;

    if( mmap_and_lock_fixed(fd, offset, length, flags, map) == NULL )
        return NULL;

    if( flags & MAL_DOUBLED ) {
        if( mmap_and_lock_fixed(fd, offset, length, flags, map+length) == NULL ) {
            munmap(map, length);
            return NULL;
        }
        length *= 2;
    }

    if( flags & PREFFAULT_RDWR )
        prefault_pages(map, length / sysconf(_SC_PAGESIZE), (flags & PREFFAULT_RDWR));

    return map;
}
```


Sep 12, 15 16:50

mman.h

Page 1/1

```
#
#ifndef _MMAN_H
#define _MMAN_H

#include "general.h"

/* Memory mapping and locking utilities */

export void prefault_pages(void *, int, int);
export void *mmap_locate(size_t, int);
export void *mmap_and_lock_fixed(int, off_t, size_t, int, void *);
export void *mmap_and_lock(int, off_t, size_t, int);

#define PROT_RDONLY 1
#define PROT_WRONLY 2
#define PROT_RDWR (PROT_RDONLY|PROT_WRONLY)

#define PREFALT_RDONLY 4
#define PREFALT_WRONLY 8
#define PREFALT_RDWR (PREFALT_RDONLY|PREFALT_WRONLY)

#define MAL_LOCKED 16
#define MAL_DOUBLED 32

#endif /* _MMAN_H */
```

Sep 12, 15 17:00

param.c

Page 1/8

```

#
#include "general.h"

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>
#include "assert.h"
#include <ctype.h>
#include "argtab.h"
#include "param.h"

/*
 * Parameter types, represented by constant strings...
 */

PARAM_TYPE_DECL(bool,    int,      "%d",    "%d");
PARAM_TYPE_DECL(int16,   uint16_t, "%hi",    "%hu");
PARAM_TYPE_DECL(int32,   uint32_t, "%li",    "%u");
PARAM_TYPE_DECL(int64,   uint64_t, "%Li",    "%llu");
PARAM_TYPE_DECL(double,  double,   "%lg",    "%g");
PARAM_TYPE_DECL(string,  char *,    NULL,     "%s");

/*
 * Reset the str and val pointers in a param_t structure. Free strings
 * as needed, but assume that a string val that is dynamic is dealt with
 * by the caller.
 */

public void reset_param(param_t *p) {
    if(p->p_type == PARAM_TYPE(string)) { /* Special case of dynamic string in two places */
        if( p->p_val && *(char **)p->p_val == p->p_str) { /* String copy is in both places */
            p->p_dyn = 0; /* Ignore dynamic: caller is responsible */
        }
    }
    if(p->p_dyn)
        free( (void *)p->p_str );
    p->p_str = NULL;
    p->p_dyn = 0;
    if(p->p_val)
        p->p_val = NULL;
}

/*
 * Set the val pointer for a param.
 */

public void setval_param(param_t *p, void **val) {
    p->p_val = val;
}

/*
 * Push an extra value for the given parameter onto its value stack,
 * if there is room. If p_dyn is set, the last (top) value is an
 * allocated copy; free it and overwrite the slot. Otherwise the
 * value is a permanent buffer, so just push.
 */

public int set_param_value(param_t *p, char *v) {
    /* fprintf(stderr, "Pushing value %s\n", v); */
    if( !p->p_source ) {
        errno = EPERM;
        return -1;
    }
    if( p->p_dyn && p->p_str ) {
        if(p->p_val && *(const char **)p->p_val == p->p_str)
            *(const char **)p->p_val = NULL;
        free((void *)p->p_str);
        p->p_str = NULL;
    }
    p->p_str = v;
    return 0;
}

```

Sep 12, 15 17:00

param.c

Page 2/8

```

/*
 * Locate the parameter descriptor in the ps array for the named
 * parameter, if it exists.
 */

public param_t *find_param_by_name(const char *name, int sz, param_t ps[], int nps) {
    int i;

    /* fprintf(stderr, "Looking for name "); fwrite(name, sz, 1, stderr); fprintf(stderr, "\n"); */
    errno = 0;
    for(i=0; i<nps; i++)
        if( !strcmp(name, ps[i].p_name, sz) )
            return &ps[i];
    errno = EBADSLT;
    return NULL;
}

/*
 * Scan the parameters in the ps array and check whether environment
 * variables provide values for any of them. The environment variable
 * name must match the parameter using case insensitive matching. The
 * environment variable's value replaces the parameter's value, if any.
 */

public int set_param_from_env(char *env[], param_t ps[], int nps) {
    int i;

    if( !env )
        return 0;

    for(i=0; i<nps; i++) {
        char **e, *p;
        int sz;

        if( !(ps[i].p_source & PARAM_SRC_ENV) ) /* Only look for params with environment source */
            continue;
        for(e=env; p=*e; e++) {
            for(sz=0; *p && *p != '='; p++, sz++); /* Find the = */
            if( !strncasecmp(ps[i].p_name, *e, sz) ) { /* Unless true, it's not this one */
                if( ps[i].p_name[sz] ) /* If true, there is more name left over: not this one */
                    continue;
                if( set_param_value(&ps[i], (*p ? p+1 : p)) < 0 )
                    return -1;
            }
        }
    }
    return 0;
}

/*
 * Read the presented string and look for Name=Value where Name is the
 * name of a parameter. If found, push a pointer to the value. The
 * cmd string is assumed NUL-terminated after the value.
 */

public int set_param_from_cmd(char *cmd, param_t ps[], int nps) {
    char *s;
    param_t *p;

    /* fprintf(stderr, "Working on cmd %s\n", cmd); */
    if( !cmd )
        return 0;
    for(s=cmd; *s && *s != '='; s++);
    p = find_param_by_name(cmd, s-cmd, ps, nps);
    if( !p )
        return -1;
    if( !(p->p_source & PARAM_SRC_CMD) ) {
        errno = EPERM;
        return -1;
    }
    if( !*s++ ) { /* If *s non-zero, step over the = */
        errno = EINVAL; /* Name=Value string has no '=Value' part */
        return -1;
    }
    return set_param_value(p, s);
}

```

Sep 12, 15 17:00

param.c

Page 3/8

```

}

/*
 * Given a string comprising a set of space/comma/semicolon separated
 * Name=Value pairs, instantiate parameters from them. Use strtok_r
 * to parse the string, which alters the input string by replacing
 * separators with NUL characters. Each string returned by strtok_r
 * is a single NUL-terminated Name=Value element. On error, return
 * the negative of the position in the string of the current token
 * start.
 */

private int do_set_params_from_string(char *str, int opt, param_t ps[], int nps) {
    char *save;
    char *cur;
    int done;
    int ret;

    /* Initialise the strtok_r scan: skip to space */
    cur = strtok_r(str, "\\t;", &save);
    if( cur == NULL ) {
        errno = EBADMSG;
        return opt? 0 : -1; /* If parameters are optional, may succeed here for empty */
    }

    /* First parameter Name=Value should come next */
    done = 0;
    while( (cur = strtok_r(NULL, "\\t;", &save)) != NULL ) {
        if( !isalpha(*cur) ) {
            errno = EBADMSG;
            return str-cur;
        }
        ret = set_param_from_cmd(cur, ps, nps);
        if( ret < 0 )
            return str-cur;
        done++;
    }
    return (done || opt)? 0 : -1;
}

/* Parameters are compulsory */

public int set_params_from_string(char *str, param_t ps[], int nps) {
    return do_set_params_from_string(str, 0, ps, nps);
}

/* String may be empty of parameters */

public int set_opt_params_from_string(char *str, param_t ps[], int nps) {
    return do_set_params_from_string(str, 1, ps, nps);
}

/*
 * Assign a parameter value, i.e. parse its string value and write the result to
 * the location pointed to by the val pointer, which must be of the correct kind.
 */

public int assign_param(param_t *p) {
    if(p == NULL) {
        errno = EINVAL;
        return -1;
    }
    if( !p->p_val ) { /* Nowhere to put value */
        errno = EFAULT;
        return -1;
    }

    param_type *pt = p->p_type;
    if(pt == PARAM_TYPE(bool)) { /* Special cases for booleans */
        const char *s = p->p_str; /* May be NULL, for a boolean (== false) */

        if( !*s || !strncasecmp(s, "false", 6) || !strncasecmp(s, "no", 3) || !strncasecmp(s, "off", 4) ) {
            *(int *)p->p_val = 0;
            return 0;
        }
    }
}

```

Sep 12, 15 17:00

param.c

Page 4/8

```

    if( !strncasecmp(s, "true", 5) || !strncasecmp(s, "yes", 4) || !strncasecmp(s, "on", 3) ) {
        *(int *)p->p_val = 1;
        return 0;
    }
}

if( !p->p_str )          /* No value to put anywhere */
    return 0;

if(pt == PARAM_TYPE(string)) { /* Special case for strings -- no conversion needed */
    *(const char **)p->p_val = p->p_str;
    return 0;
}

// fprintf(stderr, "Scan param %s with str %s to %p using %s\n",
//      p->p_name, p->p_str, p->p_val, pt->t_scan);
return sscanf(p->p_str, pt->t_scan, p->p_val) == 1? 0 : -1;
}

/*
 * Scan the parameter table and copy out the values present, converting strings to
 * appropriate types and installing them in the external addresses where provided.
 */

public int assign_all_params(param_t ps[], int nps) {
    int n;

    for(n=0; n<nps; n++) {
        param_t *p = &ps[n];

        if(assign_param(p) < 0)
            return -1-n;
    }
    return 0;
}

/*
 * Same as above but only for parameters sourced from commands.
 */

int assign_cmd_params(param_t ps[], int nps) {
    int n;

    for(n=0; n<nps; n++) {
        param_t *p = &ps[n];

        if( p->p_source & PARAM_SRC_CMD ) {
            if(assign_param(p) < 0)
                return -n-1;
        }
    }
    return 0;
}

/*
 * Retrieve the string value of the parameter and store it in the
 * buffer pointed to by vp, which must be suitable to receive it.
 */

public int get_param_str(param_t *p, const char **vp) {
    const char *v = NULL;
    if( !p->p_str ) {
        if( p->p_type == PARAM_TYPE(bool) ) {
            *vp = "false";
            return 0;
        }
        errno = EINVAL;
        return -1;
    }
    v = p->p_str;          /* The string value for the parameter */
    if( !v ){
        errno = EINVAL;
        return -1;
    }
    *vp = v;
}

```

Sep 12, 15 17:00

param.c

Page 5/8

```

    return 0;
}

/*
 * Find parameters that match an argxxx structure.  An arg_xxx structure matches a
 * parameter if (one of) its long name(s) matches a parameter name for which an ARG
 * source has been set.  The long option names are tried in order; only one may match!
 */

private param_t *arg_param_match(const char *a, param_t ps[], int nps) {
    param_t *p, *ret;
    const char *ap;

    if(a == NULL)
        /* There are no long option names */
        return NULL;
    ret = NULL;
    ap = a;
    /* First option name starts here */
    for(ap=a; *a; a=ap) {

        while(*ap && *ap != ',') ap++; /* Skip to end of (first) option name */

        p = find_param_by_name(a, ap-a, ps, nps);
        if(p == NULL)
            /* No match for that name */
            continue;

        if(ret && ret != p) {
            /* Multiple matches! */
            errno = EBADSLT;
            return NULL;
        }

        ret = p;
        if(*ap == ',') ap++;
        /* At least one match found */
        /* Skip a comma, if more to come */
    }
    return ret;
}

/* ASSUME that the count and 'data' values in every argxxx follow the hdr directly */
/* IF THAT IS TRUE, then we can use the ->count and ->data members of ANY arg_xxx */
#define ARG_COUNT(a) (((struct arg_int *) (a))->count)
#define ARG_DATA(a) ((void *) ((struct arg_int *) (a))->ival)

/*
 * Install defaults into an argtable from matching parameter structures.  This assumes
 * internal knowledge of the arg_hdr structures to determine the relevant parameter
 * structure etc. to use.  The parameter's string value is converted using the arg_hdr
 * structure's scan function and is pre-installed in the arg_xxx structure, which is then
 * reset.
 */

public int arg_defaults_from_params(void **argtable, int nargs, param_t ps[], int nps) {
    struct arg_hdr **ate = (struct arg_hdr **) &argtable[nargs-1]; /* The arg_end structure slot */
    param_t *endp = &ps[nps];

    if( !((*ate)->flag & ARG_TERMINATOR) ) {
        errno = EINVAL;
        return -1;
    }

    struct arg_hdr **atp;
    for(atp=(struct arg_hdr **) argtable; atp<ate; atp++) {
        struct arg_hdr *a = *atp;

        param_t *p = arg_param_match(a->longopts, ps, nps);

        if(p == NULL)
            /* No parameter matched this argument */
            continue;
        /* Must be an ARG parameter, or coding problem */
        assertv( (p->p_source & PARAM_SRC_ARG), "Param %s not ARG sourced\n", p->p_name );
        if( !p->p_str )
            /* No string value, no default */
            continue;

        // fprintf(stderr, "Found parameter %s with addr %p, str %s\n", p->p_name, p->p_val, p->p_str);

        /* Copy the parameter's value to the arg structure -- fake an argument parse */
    }
}

```

Sep 12, 15 17:00

param.c

Page 6/8

```

    (*a->resetfn)(a->parent); /* Reset the counter; init the structure */
    int ret = (*a->scanfn)(a->parent, p->p_str);
    /* Else value compatibility error: abort */
    assertv(ret == 0, "Param %s str %s does not pass arg scanfn\n", p->p_name, p->p_str);
    ret = (*a->checkfn)(a->parent);
    /* Else value compatibility error: abort */
    assertv(ret == 0, "Param %s str %s fails arg checkfn\n", p->p_name, p->p_str);
    ARG_COUNT(a) = 0; /* This was a default */
}

return 0;
}

/*
 * Copy the results from a parsed argtable back to the locations pointed to by their
 * matching param structures; the matching structures are determined as for the
 * arg_defaults_from_params() routine above. Unfortunately, there is no way to know what
 * kind of value the argxxx structure describes -- we have to assume that the parameter
 * knows the type (and therefore the size to copy).
 *
 * We also copy the value back into the parameter string form -- this might entail some loss
 * of precision for real number values.
 */

public int arg_results_to_params(void **argtable, param_t ps[], int nps) {
    struct arg_hdr **ate = (struct arg_hdr **)argtable;
    param_t *endp = &ps[nps];

    while( (*ate) && !((*ate)->flag & ARG_TERMINATOR) ) ate++; /* Find the end */
    if( !(*ate) || !((*ate)->flag & ARG_TERMINATOR) ) {
        errno = EINVAL;
        return -1;
    }

    struct arg_hdr **atp;
    for(atp=(struct arg_hdr **)argtable; atp<ate; atp++) {
        struct arg_hdr *a = *atp;
        void *av;

        if( ARG_COUNT(a) == 0 ) /* There is no command-line argument value */
            continue;

        param_t *p = arg_param_match(a->longopts, ps, nps);

        if(p == NULL) /* No parameter matched this argument */
            continue;
        if( !p->p_val ) /* Nowhere to put the value */
            continue;

        av = ARG_DATA(a) + (ARG_COUNT(a)-1)*p->p_type->t_size;
        memcpy(p->p_val, av, p->p_type->t_size);

        /*
         * This one copy back is tricky... If *p->p_val is not already the
         * same as p->p_str, it must have come from a static string from
         * argument or environment, since only the assign code changes the
         * val content and it copies the str. Therefore we copy back the
         * val content and turn off the free-it flag.
         *
         * On the other hand, if p->p_str is in fact *p->p_val, there is
         * nothing further to do.
         */
        if(p->p_type == PARAM_TYPE(string)) {
            const char *v = *(char **)p->p_val;
            if(v != p->p_str) {
                if(p->p_dyn)
                    free((void *)p->p_str);
                p->p_dyn = 0;
                p->p_str = v;
            }
            continue; /* We are done, in this case */
        }

        if(p->p_dyn) /* Free the old str value if necessary */
            free((void *)p->p_str);
    }
}

```

Sep 12, 15 17:00

param.c

Page 7/8

```

    int ret = param_value_to_string(p, &p->p_str);
    p->p_dyn = 1; /* The new value is a dynamic string */
    assertv(ret >= 0, "Update of parameter %s str from val for arg %d failed\n",
            p->p_name, ate-atp+1);
}
return 0;
}

#undef ARG_COUNT
#undef ARG_DATA

/*
 * Generate part of a usage message based on the parameter structures
 */

public void param_option_usage(FILE *f, int spc, param_t ps[], int nps) {
    int i;
    char *buf = malloc(spc+1);

    for(i=0; i<spc; i++) buf[i] = ' ';
    buf[spc] = '\0';
    for(i=0; i<nps; i++) {
        param_t *p = &ps[i];
        if( !(p->p_source & PARAM_SRC_ARG) )
            continue;
        fprintf(f, "%s--%s<%s>:%s\n", buf, p->p_name, p->p_type, p->p_gloss);
    }
}

public void param_brief_usage(char *buf, int sz, param_t ps[], int nps) {
    int i,
        used = 0,
        rest = sz-1;

    for(i=0; i<nps && rest > 0; i++) {
        param_t *p = &ps[i];
        char *type = "NULL";
        int n;
        if( !(p->p_source & PARAM_SRC_ARG) )
            continue;
        n = snprintf(&buf[used], rest, "[--%s<%s>]", p->p_name, p->p_type);
        used += n;
        rest -= n;
    }
    buf[used] = '\0';
}

/*
 * Convert a parameter value and store in a dynamically allocated string
 */

#define LOCALBUF_SIZE 64

public int param_value_to_string(param_t *p, const char **s) {
    param_type *pt = p->p_type;
    char buf[LOCALBUF_SIZE];
    int used = 0;

    if( !p->p_val ) /* Nowhere to get the value from */
        return 0;

    /* These cases are systematically treatable */
    if(pt == PARAM_TYPE(string)) {
        used = snprintf(&buf[0], LOCALBUF_SIZE-1, pt->t_show, *(char **)p->p_val);
    }
    if(pt == PARAM_TYPE(bool)) {
        used = snprintf(&buf[0], LOCALBUF_SIZE-1, pt->t_show, *(int *)p->p_val);
    }
    if(pt == PARAM_TYPE(int16)) {
        used = snprintf(&buf[0], LOCALBUF_SIZE-1, pt->t_show, *(uint16_t *)p->p_val);
    }
    if(pt == PARAM_TYPE(int32)) {
        used = snprintf(&buf[0], LOCALBUF_SIZE-1, pt->t_show, *(uint32_t *)p->p_val);
    }
    if(pt == PARAM_TYPE(int64)) {

```


Sep 12, 15 17:00

param.c

Page 8/8

```

    used = snprintf(&buf[0], LOCALBUF_SIZE-1, pt->t_show, *(uint64_t *)p->p_val);
}
if(pt == PARAM_TYPE(double)) {
    used = snprintf(&buf[0], LOCALBUF_SIZE-1, pt->t_show, *(double *)p->p_val);
}
if( !(*s = strndup(&buf[0], used)) )
    return -1;

return used;
}

public void debug_params(FILE *fp, param_t ps[], int nps) {
    int i;

    for(i=0; i<nps; i++) {
        param_t *p = &ps[i];
        param_type *pt = p->p_type;

        fprintf(fp, "Parameter '%s': type %s addr %p str %p='%s'",
            p->p_name, pt->t_name, p->p_val, p->p_str, p->p_str);
        if(p->p_val) {
            fprintf(fp, " val'");
            if(pt == PARAM_TYPE(bool)) {
                fprintf(fp, pt->t_show, *(int *)p->p_val);
            }
            if(pt == PARAM_TYPE(int16)) {
                fprintf(fp, pt->t_show, *(uint16_t *)p->p_val);
            }
            if(pt == PARAM_TYPE(int32)) {
                fprintf(fp, pt->t_show, *(uint32_t *)p->p_val);
            }
            if(pt == PARAM_TYPE(int64)) {
                fprintf(fp, pt->t_show, *(uint64_t *)p->p_val);
            }
            if(pt == PARAM_TYPE(string)) {
                fprintf(fp, pt->t_show, *(char **)p->p_val);
            }
            if(pt == PARAM_TYPE(double)) {
                fprintf(fp, pt->t_show, *(double *)p->p_val);
            }
        }
        fprintf(fp, "\n");
    }
}

```

Sep 12, 15 16:50

param.h

Page 1/1

```

#
#ifndef _PARAM_H
#define _PARAM_H

#include <stdio.h>
#include <stdint.h>

#include "general.h"

typedef const struct {
    const char *t_name;
    int t_size;
    const char *t_scan;
    const char *t_show;
}
    param_type;

#define PARAM_TYPE(name) param_type_ ## name
#define PARAM_TYPE_DECL(name,size,scan,show) param_type PARAM_TYPE(name)[] = { "<" #name ">" , sizeof(size), scan, show, }
#define PARAM_TYPE_EXPORT(name) export param_type PARAM_TYPE(name)[];

PARAM_TYPE_EXPORT(bool);
PARAM_TYPE_EXPORT(int16);
PARAM_TYPE_EXPORT(int32);
PARAM_TYPE_EXPORT(int64);
PARAM_TYPE_EXPORT(double);
PARAM_TYPE_EXPORT(string);

typedef struct
{
    const char *p_name;          /* Name of this parameter */
    const char *p_str;           /* String value for this parameter */
    void *p_val;                 /* Location where value is to be stored */
    param_type *p_type;          /* Type of the parameter, for value conversion */
    int p_source;                /* Possible sources of the values */
    const char *p_gloss;         /* Explanation of this parameter */
    int p_dyn;                   /* If true, free and replace str on push */
}
    param_t;

#define PARAM_SRC_ENV    0x1
#define PARAM_SRC_ARG    0x2
#define PARAM_SRC_CMD    0x4

export int set_param_value(param_t *, char *);
export param_t *find_param_by_name(const char *, int, param_t [], int);
export int set_param_from_env(char *[], param_t [], int);
export int set_params_from_string(char *, param_t [], int);
export int set_opt_params_from_string(char *, param_t [], int);
export int get_param_str(param_t *, const char **);
// export void param_brief_usage(char *, int, param_t [], int);
// export void param_option_usage(FILE *, int, param_t [], int);
// export const char *pop_param_value(param_t *);
export void reset_param(param_t *);
export void setval_param(param_t *, void **);
export int assign_param(param_t *);
export int assign_all_params(param_t *, int);
export int assign_cmd_params(param_t *, int);
export int param_value_to_string(param_t *, const char **);
export int arg_defaults_from_params(void **, int, param_t [], int);
export int arg_results_to_params(void **, param_t [], int);
export void debug_params(FILE *, param_t [], int);

#endif /* _PARAM_H */

```

Sep 12, 15 17:01

queue.c

Page 1/2

```

#
#include "general.h"

#include <stdlib.h>
#include "assert.h"
#include "queue.h"

/*
 * Implements a doubly-linked queue in ring form.
 *
 * Invariant: every q structure is doubly-linked; new structures are singletons.
 */

public queue *init_queue(queue *p) {
    if( p == NULL ) {
        p = (queue *)calloc(1, sizeof(queue));
        assertv(p != NULL, "Queue alocation failure\n");
    }
    p->q_next = p->q_prev = p;
    return p;
}

/*
 * Remove p from its queue and make it a singleton. You cannot detach
 * a singleton from its queue.
 */

public queue *de_queue(queue *p) {
    if( p->q_next == p )
        return NULL;
    p->q_prev->q_next = p->q_next;
    p->q_next->q_prev = p->q_prev;
    p->q_next = p->q_prev = p;
    return p;
}

/*
 * Splice q and p together so that p immediately follows q and the
 * next and prev chains continue in the correct senses
 */

public queue *splice_queue(queue *q, queue *p) {
    queue *qn, *pp;

    qn = q->q_next;
    q->q_next = p;
    pp = p->q_prev;
    p->q_prev = q;
    qn->q_prev = pp;
    pp->q_next = qn;
    return q;
}

/*
 * Unsplice a queue: cut the ring at start and end and relink. Also
 * join start and end.
 */

public queue *unsplice_queue(queue *start, queue *end) {
}

/*
 * Apply a function to each queue member in [start,end). The function
 * is called with arg as its first argument and the queue structure
 * pointer as its second. The first function, map_queue_nxt,
 * traverses the segment "forward" while the second goes "backward".
 *
 * If start == end or end is not in the list (e.g. end is NULL) the
 * functions traverse the whole list visiting each node exactly once.
 */

public void map_queue_nxt(queue *start, queue *end, void (*fn)(void *, queue *), void *arg) {
    for_nxt_in_Q(queue *p, start, end)

```

Sep 12, 15 17:01

queue.c

Page 2/2

```
    (*fn)(arg, p);  
    end_for_nxt;  
}  
  
public void map_queue_prv(queue *start, queue *end, void (*fn)(void *, queue *), void *arg) {  
    for_prv_in_Q(queue *p, start, end)  
        (*fn)(arg, p);  
    end_for_prv;  
}
```

Sep 12, 15 16:50

queue.h

Page 1/2

```

#
#ifndef _QUEUE_H
#define _QUEUE_H

#include "general.h"

typedef struct q
{
    struct q *q_next;
    struct q *q_prev;
}
queue;

export queue *de_queue(queue *);
export queue *init_queue(queue *);
export queue *splice_queue(queue *, queue *);
export queue *unsplice_queue(queue *, queue *);
export void map_queue_nxt(queue *, queue *, void (*)(void *, queue *, void *));
export void map_queue_prv(queue *, queue *, void (*)(void *, queue *, void *));

#define queue_next(q) ((q)->q_next)
#define queue_prev(q) ((q)->q_prev)

#define queue_ins_after(q,i) splice_queue((q), (i))
#define queue_ins_before(q,i) splice_queue((i), (q))

#define queue_singleton(q) ((q)->q_next == (q) && (q)->q_prev == (q))

#define QUEUE_HEADER(name) queue name = { &name, &name }

/*
 * These macro definitions do essentially the same as the
 * map_queue_nxt and map_queue_prv but they don't leave the current
 * local scope -- so for instance one can break the loop early in this
 * form whereas one cannot in the (default) map function.
 *
 * The var argument is a variable that will hold the current node
 * pointer as the loop proceeds. It can be declared locally to the
 * for_nxt by including its declaration in the macro call:
 *
 * for_nxt_in_Q(queue *ptr,start,end) ...
 *
 * or it can be a variable declared outside the scope of the for_nxt
 * in which case just its name is given as argument and it will
 * persist after the map-loop ends.
 *
 * The macros evaluate start and end exactly once and execute the User
 * Code once for each list element in the range [start,end) with var
 * set to that element. If start==end or end is not actually in the
 * list, the loop traverses the whole list exactly once visiting each
 * node exactly once.
 *
 * Note that it is also possible to remove node __p during the USER
 * CODE because it is neither the node we are about to work on nor the
 * end point node. It may be the start node, however: the user should
 * deal with that case!
 */

#define for_nxt_in_Q(var,start,end) \
do { queue *__s = (start), *__e = (end); \
    queue *__p = __s; \
    int __done = 0; \
    while(!__done) { queue *__n = queue_next(__p); \
        __done = (__n == __s || __n == __e); \
        var = __p; __p = __n; \
        /* USER CODE GOES HERE */ \
    } \
} while(0)

#define end_for_nxt \
} } while(0)

#define for_prv_in_Q(var,start,end) \
do { queue *__s = (start), *__e = (end); \
    queue *__p = __s; \
    int __done = 0; \

```

Sep 12, 15 16:50

queue.h

Page 2/2

```
while(!__done) { queue *__n = queue_prev(__p);    \
  __done = (__n == __s || __n == __e);          \
  var = __p; __p = __n;                          \
  /* USER CODE GOES HERE */

#define end_for_prv \
} } while(0)

#endif /* _QUEUE_H */
```

Sep 22, 15 15:13

reader.c

Page 1/10

```

#
#include "general.h"

#include <stdio.h>
#include <stdlib.h>
#include "assert.h"
#include <time.h>
#include <sys/time.h>
#include <sys/resource.h>
#include <sys/capability.h>

#include <zmq.h>
#include <pthread.h>

#include <comedi.h>
#include <comedilib.h>
#include <sys/mman.h>
#include <fcntl.h>
#include <unistd.h>
#include <errno.h>
#include <string.h>

#include "util.h"
#include "param.h"
#include "queue.h"
#include "strbuf.h"
#include "mman.h"
#include "chunk.h"
#include "adc.h"
#include "snapshot.h"
#include "tidy.h"
#include "reader.h"
#include "writer.h"

/*
 * READER global data structures
 */

public rparams reader_parameters; /* The externally-visible parameters for the reader thread */
public adc reader_adc; /* The ADC object for the READER */

/*
 * READER state machine definitions.
 *
 * The READER state is kept in the rp_state variable, private to the
 * READER thread.
 *
 * ERROR state: this occurs when a serious error happens, normally due
 * to bad parameters. One can leave ERROR state using the Param
 * command.
 *
 * PARAM state: results from initialisation by the main thread routine
 * and after the receipt of a Param command, because of the activity
 * of the verify function. Failure of parameters to verify sends us
 * to ERROR state. Successful verification also results in the
 * creation and parameterisation of an ADC object.
 *
 * RESTING state: a successful execution of the Init command leaves us
 * in RESTING state. In this state, an initialised ADC object is
 * available. Errors in parameter verification or instantiation of
 * the initialised ADC object put us into ERROR state.
 *
 * ARMED state: executing the Go command initiates a data transfer and
 * moves the READER to this state. We stay in ARMED state until the
 * first data has been seen (i.e. the ADC object has changed from
 * running to running and live). Failure of data to arrive within a
 * reasonable time causes an automatic transition to the ERROR
 * state, with the same cleanup as done by the Halt command, which
 * may be issued in this or the RUN state.
 *
 * RUN state: automatic transition from ARMED on receipt of the first
 * data. In ARMED and RUN state the Halt command will terminate
 * data acquisition and return the READER to ERROR state (as a
 * special case; the parameters are valid, but after Halt there is

```

Sep 22, 15 15:13

reader.c

Page 2/10

```

*   no ADC object).
*
* The Quit command issued in any state will cause the READER to shut
* down cleanly.
*
* The WRITER will reject Snap commands unless the READER is in ARMED
* or RUN state (in fact, unless the ADC object exists and reports
* itself as running).
*/

private int rp_state;

#define READER_ERROR    0      /* An error occurred, base start state */
#define READER_PARAM    1      /* There are parameters that need to be verified */
#define READER_RESTING  2      /* READER is ready, Comedi and mmap setup has been done */
#define READER_ARMED    3      /* The ADC has been started */
#define READER_RUN      4      /* Data from the ADC has been seen in the buffers */

/*
 * READER forward definitions
 */

private void drain_reader_chunk_queue();

/*
 * READER thread comms initialisation.
 * Called after the context is created.
 */

private void *writer;
private void *tidy;
private void *log;
private void *command;

private void create_reader_comms() {
    import void *snapshot_zmq_ctx;
    /* Create necessary sockets */
    command = zh_bind_new_socket(snapshot_zmq_ctx, ZMQ_REP, READER_CMD_ADDR); /* Receive commands */
    assertv(command != NULL, "Failed to instantiate reader command socket\n");
    log      = zh_connect_new_socket(snapshot_zmq_ctx, ZMQ_PUSH, LOG_SOCKET); /* Socket for log messages */
    assertv(log != NULL, "Failed to instantiate reader log socket\n");
    writer   = zh_bind_new_socket(snapshot_zmq_ctx, ZMQ_PAIR, READER_QUEUE_ADDR);
    assertv(writer != NULL, "Failed to instantiate reader queue socket\n");
    tidy     = zh_connect_new_socket(snapshot_zmq_ctx, ZMQ_PAIR, TIDY_SOCKET); /* Socket to TIDY thread */
    assertv(tidy != NULL, "Failed to instantiate reader->tidy socket\n");
}

/* Close everything created above. */

private void close_reader_comms() {
    zmq_close(command);
    zmq_close(log);
    zmq_close(writer);
    zmq_close(tidy);
}

/*
 * Copy the necessary capabilities from permitted to effective set (failure is fatal).
 *
 * The READER needs:
 *
 * CAP_IPC_LOCK -- ability to mmap and mlock pages.
 * CAP_SYS_NICE -- ability to set RT scheduling priorities
 *
 * These capabilities should be in the CAP_PERMITTED set, but not in CAP_EFFECTIVE which was cleared
 * when the main thread dropped privileges by changing to the desired non-root uid/gid.
 */

private int set_up_reader_capability() {
    cap_t c = cap_get_proc();
    const cap_value_t vs[] = { CAP_IPC_LOCK, CAP_SYS_NICE, };

    cap_set_flag(c, CAP_EFFECTIVE, sizeof(vs)/sizeof(cap_value_t), &vs[0], CAP_SET);
    return cap_set_proc(c);
}

```


Sep 22, 15 15:13

reader.c

Page 3/10

```

/*
 * Get a value from the monotonic kernel clock and express in nanoseconds.
 */

public uint64_t monotonic_ns_clock() {
    uint64_t ret;
    struct timespec now;

    clock_gettime(CLOCK_MONOTONIC, &now);    /* Timestamp for debugging */
    ret = now.tv_sec;
    ret = ret*1000000000 + now.tv_nsec;
    return ret;
}

/*
 * Process a READER command from MAIN thread. Generate replies as necessary.
 * Returns true if processing messages should continue..
 */

private int process_reader_command(void *s) {
    rparams *rp = &reader_parameters;
    int used;
    int ret;
    strbuf cmd;
    char *cmd_buf;
    strbuf err;

    used = zh_get_msg(s, 0, sizeof(strbuf), &cmd);
    if( !used ) { /* It was a quit message */
        if(rp_state == READER_ARMED || rp_state == READER_RUN || rp_state == READER_RESTING)
            adc_stop_data_transfer(reader_adc);
        return false;
    }

    cmd_buf = strbuf_string(cmd);
    err = strbuf_next(cmd);

    if(verbose > 1)
        zh_put_multi(log, 3, "READER cmd:'", &cmd_buf[0], "'");

    ret = 0;
    switch(cmd_buf[0]) {
    case 'p':
    case 'P':
        if( rp_state != READER_PARAM && rp_state != READER_RESTING && rp_state != READER_ERROR ) {
            strbuf_printf(err, "NO: Param issued but not in PARAM, RESTING or ERROR state");
            ret = -1;
            break;
        }
        ret = set_params_from_string(&cmd_buf[0], globals, n_global_params);
        if( ret < 0 ) {
            strbuf_printf(err, "NO: Param -- parse error at position %d", -ret);
            break;
        }
        ret = assign_cmd_params(globals, n_global_params);
        if( ret < 0 ) {
            strbuf_printf(err, "NO: Param -- assign error on param %d: %m", -ret);
            break;
        }

        /* Otherwise, succeeded in updating parameters */
        strbuf_printf(err, "NO: Param -- verify error: ");
        ret = verify_reader_params(&reader_parameters, err);
        if( ret < 0 ) {
            break;
        }
        strbuf_printf(err, "OK Param");
        rp_state = READER_PARAM;
        break;

    case 'i':
    case 'I':
        if( rp_state != READER_PARAM ) {
            strbuf_printf(err, "NO: Init issued but not in PARAM state");

```

Sep 22, 15 15:13

reader.c

Page 4/10

```

        ret = -1;
        break;
    }
    strbuf_printf(err, "NO: Init --- param verify error: ");
    ret = verify_reader_params(&reader_parameters, err);
    if( ret < 0 ) {
        rp_state = READER_ERROR;
        break;
    }
    ret = adc_init(reader_adc, err);
    if( ret < 0 ) {
        rp_state = READER_ERROR;
        break;
    }
    if(verbose > 0) {
        /* Borrow the err buffer */
        strbuf_printf(err, "READER Init with dev %s, freq %g [Hz], isp %d [ns] and buf %d [MiB]",
            rp->r_device, rp->r_frequency, adc_ns_per_sample(reader_adc), rp->r_bufsz);
        zh_put_multi(log, 1, strbuf_string(err));
    }
    strbuf_printf(err, "OK Init --- nchan %d isp %d [ns]", NCHANNELS, adc_ns_per_sample(reader_adc));
    rp_state = READER_RESTING;
    break;

case 'g':
case 'G':
    if( rp_state != READER_RESTING ) {
        strbuf_printf(err, "NO: Go issued but not in RESTING state");
        ret = -1;
        break;
    }
    ret = adc_start_data_transfer(reader_adc, err);
    if( ret < 0 ) {
        rp_state = READER_ERROR;
        break;
    }
    strbuf_printf(err, "OK Go");
    rp_state = READER_ARMED;
    break;

case 'h':
case 'H':
    if( rp_state != READER_ARMED && rp_state != READER_RUN ) {
        strbuf_printf(err, "NO: Halt issued but not in ARMED or RUN state");
        ret = -1;
        break;
    }
    adc_stop_data_transfer(reader_adc); /* Terminate any transfer in progress */
    drain_reader_chunk_queue();        /* Empty the chunk queue */
    strbuf_printf(err, "OK Halt");
    adc_destroy(reader_adc);
    reader_adc = NULL;
    rp_state = READER_ERROR;
    break;

default:
    strbuf_printf(err, "NO: READER --- Unexpected reader command");
    ret = -1;
    break;
}
if( ret < 0 ) {
    strbuf_revert(cmd);
    zh_put_multi(log, 4, strbuf_string(err), "\n>", &cmd_buf[0], ""); /* Error occurred, log it */
}
strbuf_clear(cmd);
zh_put_msg(s, 0, sizeof(strbuf), (void *)&err); /* return message */
return true;
}

/*
 * Set the READER thread to real-time priority, if RTPRIO is set...
 */

public int set_reader_rt_scheduling() {

    if( reader_parameters.r_schedprio > 0 ) { /* Then there is RT priority scheduling to set up */

```

Sep 22, 15 15:13

reader.c

Page 5/10

```

    if( set_rt_scheduling(reader_parameters.r_schedprio) < 0 )
        return -1;

    /* Successfully applied RT scheduling */
    return 1;
}

/* RT scheduling not applicable: no RTPRIO set */
return 0;
}

/*
 * Handle a message from the WRITER. The message will be a chunk
 * which is ready to add to the READER's pending-work queue. Chunks
 * arrive here with a state of SNAPSHOT_WAITING or SNAPSHOT_ERROR (if
 * they were in transit when an error occurred). The latter are sent
 * straight back to the WRITER, which is counting down pending chunks
 * to file completion, after their frame has been released.
 */

private QUEUE_HEADER(ReaderChunkQ);
private chunk_t *rq_head = NULL;

private int process_queue_message(void *s) {
    rparams *rp = &reader_parameters;
    chunk_t *c;
    int ret;

    ret = zh_get_msg(s, 0, sizeof(chunk_t *), (void *)&c);
    assertv(ret==sizeof(chunk_t *), "Received message from WRITER with wrong size %d (not %d)\n", ret, sizeof(chunk_t *));

    if(rp_state != READER_ARMED && rp_state != READER_RUN) {
        strbuf_appendf(c->c_error, "READER thread ADC is not running");
        c->c_status = SNAPSHOT_ERROR;
    }
    else { /* Check the chunk is still current -- set SNAPSHOT_ERROR state on failure */
        adc_setup_chunk(reader_adc, c);
        if( !c->c_ring )
            c->c_status = SNAPSHOT_ERROR;
    }

    if(c->c_status==SNAPSHOT_ERROR) {
        /* we send it straight back */
        ret = zh_put_msg(writer, 0, sizeof(chunk_t *), (void *)&c);
        assertv(ret==sizeof(chunk_t *), "Message returned to WRITER with wrong size %d (not %d)\n", ret, sizeof(chunk_t *));
        ret = zh_put_msg(tidy, 0, sizeof(frame *), &c->c_frame);
        assertv(ret==sizeof(frame *), "Frame message to TIDY with wrong size %d (not %d)\n", ret, sizeof(frame *));
        c->c_frame = NULL;
        return true;
    }

    assertv(c->c_status==SNAPSHOT_WAITING, "Received chunk c:%04hx has unexpected state %s\n", c->c_name, snapshot_status(c->c_status));

    /* Add the chunk to the READER chunk queue in order of increasing *last* sample */
    queue *pos = &ReaderChunkQ;
    if( !queue_singleton(&ReaderChunkQ) ) {
        for_nxt_in_Q(queue *p, queue_next(&ReaderChunkQ), &ReaderChunkQ);
        chunk_t *h = rq2chunk(p);
        if(h->c_last > c->c_last) {
            pos = p;
            break;
        }
        end_for_nxt;
    }
    queue_ins_before(pos, chunk2rq(c));
    if(pos == &ReaderChunkQ) {
        rq_head = c; /* Points to the chunk at the head of the READER queue, when not NULL */
    }
    return true;
}

/*
 * Abort the chunk which is at the head of the ReaderChunkQ, i.e. it is
 * queue_next(&ReaderChunkQ). This means we must scan for its
 * siblings in the queue, remove them and set their status to

```

Sep 22, 15 15:13

reader.c

Page 6/10

```

* SNAPSHOT_ERROR, and return them to the WRITER. We assume that the
* caller has set the c_error strbuf.
*/

private void abort_queue_head_chunk() {
    snapfile_t *parent = rq2chunk(queue_next(&ReaderChunkQ))->c_parent;
    int ret;

    for_nxt_in_Q(queue *p, queue_next(&ReaderChunkQ), &ReaderChunkQ);
    chunk_t *c = rq2chunk(p);
    if(c->c_parent == parent) {
        de_queue(p);
        c->c_status = SNAPSHOT_ERROR;
        ret = zh_put_msg(tidy, 0, sizeof(frame *), &c->c_frame);
        assertv(ret==sizeof(frame *), "Frame message to TIDY with wrong size %d (not %d)\n", ret, sizeof(frame *));
        c->c_frame = NULL;
        ret = zh_put_msg(writer, 0, sizeof(chunk_t *), (void *)&c);
        assertv(ret==sizeof(chunk_t *), "Abort to WRITER with wrong size %d (not %d)\n", ret, sizeof(chunk_t *));
    }
    end_for_nxt;
    rq_head = queue_singleton(&ReaderChunkQ) ? NULL : rq2chunk(queue_next(&ReaderChunkQ));
}

/*
* Complete the chunk at the head of the ReaderChunkQ: remove the queue
* head and compute new head chunk; copy the data for the old head;
* send the frame to TIDY for release and the chunk pointer back to
* WRITER for book-keeping. Before doing this, check we still have
* the data for the head chunk and if not then abort it.
*/

private void complete_queue_head_chunk() {
    int ret;
    chunk_t *c = rq_head;

    if(c->c_first < adc_ring_tail(reader_adc)) { /* Oops, we are too late */
        abort_queue_head_chunk();
        return;
    }

    de_queue(chunk2rq(rq_head));
    rq_head = queue_singleton(&ReaderChunkQ) ? NULL : rq2chunk(queue_next(&ReaderChunkQ));

    copy_chunk_data(c);

    ret = zh_put_msg(tidy, 0, sizeof(frame *), &c->c_frame);
    assertv(ret==sizeof(frame *), "Frame message to TIDY with wrong size %d (not %d)\n", ret, sizeof(frame *));
    c->c_frame = NULL;

    ret = zh_put_msg(writer, 0, sizeof(chunk_t *), (void *)&c);
    assertv(ret==sizeof(chunk_t *), "Abort to WRITER with wrong size %d (not %d)\n", ret, sizeof(chunk_t *));
}

/*
* Drain the READER chunk queue when turning off the data capture.
* Any snapshots in progress are aborted.
*/

private void drain_reader_chunk_queue() {
    while( !queue_singleton(&ReaderChunkQ) ) {
        chunk_t *c = rq2chunk(queue_next(&ReaderChunkQ));
        strbuf_appendf(c->c_error, "aborted because of READER ADC shutdown");
        abort_queue_head_chunk();
    }
}

/*
* READER thread message loop
*
* The two variables buf_hwm_samples and buf_window_samples are
* determined by the program parameters window and bufhwm and set the
* policy for moving the ring buffer tail pointer. Their values are
* computed in the parameter verify routine for the READER (see below).
*/

```

Sep 22, 15 15:13

reader.c

Page 7/10

```

* Operation is as follows. The routine waits for incoming messages
* up to a certain maximum delay; then on each pass through the loop,
* at least once per delay interval assuming we got some new data, we
* do two things:
*
* - first, try to advance the adc_ring_head position which records
* data placed in the ADC ring buffer by Comedi's ADC driver. If
* the head advances past the last sample index of any chunk we can
* write that chunk out, recomputing the next threshold for head.
*
* - second, check if the head has passed the ring buffer high-water
* mark threshold, which is computed by adding buf_hwm_samples to
* the adc_ring_tail value. If it has, the ring buffer is too full
* and we must move the adc_ring_tail using adc_data_purge(). We
* advance the tail to (at most) buf_window_samples before the
* current head position -- this ensures that we have at least the
* specified 'window' duration in the ring buffer at all times.
*
* In the first step, if the first sample index of the chunk is
* earlier than the current tail, we have been forced to purge data
* (to avoid buffer overrun in Comedi) before we got the complete
* chunk. This can only happen if the chunks are very large compared
* to the buffer, which should be disallowed by parameter checking.
*
* Furthermore, if the main loop is executed for too long without any
* data being captured, we shut down the ADC and enter error state.
*/

#define ADC_DRY_PERIOD_MAX 1000 /* Initial default is 10 [s] */

private int buf_hwm_samples = 0;
private int buf_window_samples = 0;
private int adc_dry_period_max = ADC_DRY_PERIOD_MAX;
private int reader_poll_delay = 100; /* Poll wait time [ms] */

private void reader_thread_msg_loop() { /* Read and process messages */
    uint64_t high_water_mark;
    int adc_dry_period;
    int ret;
    int running;

    /* Main loop: read messages and process messages */
    zmq_pollitem_t poll_list[] =
    { { writer, 0, ZMQ_POLLIN, 0 },
      { command, 0, ZMQ_POLLIN, 0 },
    };

#define N_POLL_ITEMS (sizeof(poll_list)/sizeof(zmq_pollitem_t))
    int (*poll_responders[N_POLL_ITEMS])(void *) =
    { process_queue_message,
      process_reader_command,
    };

    zh_put_multi(log, 1, "READER thread is initialised");
    rp_state = READER_PARAM;

    high_water_mark = adc_ring_tail(reader_adc) + buf_hwm_samples;
    adc_dry_period = adc_dry_period_max;

    reader_parameters.r_running = true;
    running = true;

    while( running && !die_die_die_now ) {
        int ret;
        int nb;
        int delay;
        int n;

        if(reader_adc && adc_is_running(reader_adc)) {
            adc_dry_period--;
            nb = adc_data_collect(reader_adc);
            if( nb ) { /* There was some new data, adc_ring_head has advanced */
                adc_dry_period = adc_dry_period_max;
                /* Once the ADC head pointer has advanced past the READER queue head's end, a chunk is ready */
                while( rq_head && rq_head->c_last <= adc_ring_head(reader_adc) ) {
                    complete_queue_head_chunk();
                }
            }
        }
    }
}

```

Sep 22, 15 15:13

reader.c

Page 8/10

```

    }

    /* Check buffer fullness; if necessary, call adc_data_purge to move adc_ring_tail */
    uint64_t head = adc_ring_head(reader_adc);
    if(head > high_water_mark) {
        uint64_t lwm = head - buf_window_samples;
        uint64_t tail = adc_ring_tail(reader_adc);
        if(lwm > tail) {
            ret = adc_data_purge(reader_adc, (int)(lwm-tail));
            assertv(ret==0, "Comedi mark read failed for %d bytes: %C", (int)(lwm-tail));
            high_water_mark = lwm + buf_hwm_samples;
        }
    }
}

if(adc_dry_period <= 0) { /* Data capture interrupted or failed to start... */
}

}

ret = zmq_poll(&poll_list[0], N_POLL_ITEMS, reader_poll_delay); /* Look for commands here */
if( ret < 0 && errno == EINTR ) { /* Interrupted */
    zh_put_multi(log, 1, "READER loop interrupted");
    break;
}
if(ret < 0)
    break;

for(n=0; n<N_POLL_ITEMS; n++) {
    if( poll_list[n].revents & ZMQ_POLLIN ) {
        running = running & (*poll_responders[n])(poll_list[n].socket); /* N.B. not && */
    }
}
}

}

/*
 * READER thread main routine
 *
 * This loop either waits for a command on the command socket, or
 * loops reading from Comedi. It aborts if it cannot get the sockets
 * it needs.
 */

public void *reader_main(void *arg) {
    int ret;
    char *thread_msg = "normal exit";

    create_reader_comms();

    if( set_up_reader_capability() < 0 ) {
        zh_put_multi(log, 1, "READER thread capabilities are deficient");
    }

    ret = set_reader_rt_scheduling();
    switch(ret) {
    case 1:
        zh_put_multi(log, 1, "READER RT scheduling succeeded");
        break;
    case 0:
        zh_put_multi(log, 1, "READER using normal scheduling: RTPRIO unset");
        break;
    default:
        zh_put_multi(log, 2, "READER RT scheduling setup failed: ", strerror(errno));
        break;
    }

    struct timespec test_stamp;
    ret = clock_gettime(CLOCK_MONOTONIC, &test_stamp);
    assertv(ret == 0, "Test failed to get monotonic clock time\n");

    reader_thread_msg_loop();
    if(rp_state == READER_ARMED || rp_state == READER_RUN || rp_state == READER_RESTING) {
        adc_stop_data_transfer(reader_adc);
        adc_destroy(reader_adc);
    }
}

```

Sep 22, 15 15:13

reader.c

Page 9/10

```

zh_put_msg(tidy, 0, 0, NULL); /* Tell TIDY thread to finish */

zh_put_multi(log, 1, "READER thread terminates by return");

/* Clean up our ZeroMQ sockets */
close_reader_comms();
reader_parameters.r_running = false;
return (void *) thread_msg;
}

/*
 * Verify reader parameters and generate reader state description.
 */

public int verify_reader_params(rparams *rp, strbuf e) {
    import int writer_chunksize_samples();

    if( rp->r_schedprio != 0 ) { /* Check for illegal value */
        int max, min;

        min = sched_get_priority_min(SCHED_FIFO);
        max = sched_get_priority_max(SCHED_FIFO);
        if(rp->r_schedprio < min || rp->r_schedprio > max) {
            strbuf_appendf(e, "RT scheduling priority %d not in kernel's acceptable range [%d,%d]",
                           rp->r_schedprio, min, max);
            return -1;
        }
    }

    if(reader_adc) {
        adc_destroy(reader_adc);
        reader_adc = NULL;
    }
    reader_adc = adc_new(e);

    if( adc_set_chan_frequency(reader_adc, e, &rp->r_frequency) < 0 )
        return -1;

    if(rp->r_window < 1 || rp->r_window > 30) {
        strbuf_appendf(e, "Min. capture window %d seconds outwith compiled-in range [%d,%d] seconds",
                       rp->r_window, 1, 30);
        return -1;
    }

    int pagesize = sysconf(_SC_PAGESIZE)/sizeof(sampl_t);

    /* Compute the size of the desired capture window in samples, rounded up to a full page */
    int rbw_samples = rp->r_window * rp->r_frequency * NCHANNELS;
    rbw_samples = (rbw_samples*sizeof(sampl_t) + pagesize - 1) / pagesize;
    rbw_samples *= pagesize / sizeof(sampl_t);

    if(rp->r_buf_hwm_fraction < 0.5 || rp->r_buf_hwm_fraction > 0.95) {
        strbuf_appendf(e, "Ring buffer high-water mark fraction %g outwith compiled-in range [%g,%g] seconds",
                       rp->r_buf_hwm_fraction, 0.5, 0.95);
        return -1;
    }

    /* Compute ring buffer high-water mark in samples, rounded up to a full page */
    int bhwm_samples = rp->r_buf_hwm_fraction * rp->r_bufsz * 1024 * 1024;
    bhwm_samples = (bhwm_samples + pagesize - 1) / pagesize;
    bhwm_samples = pagesize * bhwm_samples / sizeof(sampl_t);

    if(rbw_samples > bhwm_samples) {
        strbuf_appendf(e, "Capture window of %d [kiB] is bigger than ring buffer high-water mark at %d [kiB]",
                       rbw_samples*sizeof(sampl_t)/1024, bhwm_samples*sizeof(sampl_t)/1024);
        return -1;
    }

    /* Check the window and high-water mark against the chunk size */
    int chunksize = writer_chunksize_samples();
    if(chunksize) {
        if(rbw_samples < chunksize) {
            strbuf_appendf(e, "Capture window of %d [kiB] is smaller than chunk size %d [kiB]",
                           rbw_samples*sizeof(sampl_t)/1024, chunksize*sizeof(sampl_t)/1024);

```

Sep 22, 15 15:13

reader.c

Page 10/10

```
    }
    return -1;
}
if(bhwm_samples+2*chunksize > rp->r_bufsz*1024*1024/sizeof(sampl_t)) {
    strbuf_appendf(e, "Ring overflow region %d [kiB] is smaller than twice the chunk size %d[kiB]",
        (rp->r_bufsz*1024*1024-bhwm_samples*sizeof(sampl_t))/1024, chunksize*sizeof(sampl_t)/1024);
    return -1;
}
}

if( adc_set_bufsz(reader_adc, e, rp->r_bufsz) < 0 )
    return -1;

if( adc_set_range(reader_adc, e, rp->r_range) < 0 )
    return -1;

adc_set_device(reader_adc, rp->r_device); /* Record the path, don't open the device */

/* Determine the READER main loop poll delay from the chunk duration */
double d = 1e-6 * chunksize * adc_ns_per_sample(reader_adc); /* Length of a chunk in [ms] */
reader_poll_delay = ((d+2.5)/5 > 100)? 100.0 : (d+2.5)/5; /* One fifth of a chunk or 100[ms] */

/* Set the tail policy variables */
buf_hwm_samples = bhwm_samples;
buf_window_samples = rbw_samples;

rp_state = READER_PARAM;
return 0;
}
```


Sep 21, 15 18:02

reader.h

Page 1/1

```
#
#include "general.h"

/*
 * The ZMQ addresses for the reader thread
 */

#define READER_CMD_ADDR "inproc://Reader-CMD"
#define READER_QUEUE_ADDR "inproc://Reader-Q"

/*
 * Reader parameter structure.
 */

typedef struct {
    double    r_frequency;      /* Per-channel sampling frequency [Hz] */
    int       r_schedprio;      /* Reader real-time priority */
    int       r_bufsz;          /* Reader buffer size [MiB] */
    int       r_range;          /* ADC full-scale range [mV] */
    double    r_window;         /* Snapshot window [s] (must fit in buffer) */
    double    r_buf_hwm_fraction; /* Ring buffer high-water mark as fraction of size */
    const char *r_device;       /* Comedi device to use */
    int       r_running;        /* Thread is running and ready */
}
rparams;

export rparams    reader_parameters;

export int        verify_reader_params(rparams *, strbuf);
export void       *reader_main(void *);
export uint64_t    monotonic_ns_clock();
```

Sep 12, 15 17:05

rtprio.c

Page 1/2

```

#
#include "general.h"
#define __GNU_SOURCE
#include <syscall.h>
#include <sys/capability.h>
#include <assert.h>
#include <pthread.h>
#include <errno.h>

#include "rtprio.h"

#ifndef __GNU_SOURCE
#define gettid() (syscall(SYS_gettid)) /* No glibc interface, Linux-only call */
/*
 * Routine(s) for establishing threads in RT FIFO scheduling mode using Linux tricks
 */
public int set_rt_scheduling(int p) {
    pid_t me = gettid();
    struct sched_param pri;
    int mode;

    /* Attempt the operation */
    pri.sched_priority = p;
    if( sched_setscheduler(me, SCHED_FIFO, &pri) < 0 ) {
        return -1; /* Failed for some reason */
    }

    /* Verify the operation */
    mode = sched_getscheduler(me);
    if( mode != SCHED_FIFO ) { /* Didn't work, despite no errors... */
        errno = ENOSYS;
        return -1;
    }

    pri.sched_priority = -1; /* Check correct priority was set... */
    if( sched_getparam(me, &pri) < 0
        || pri.sched_priority != p ) {
        errno = ENOSYS;
        return -1;
    }

    /* Successfully applied RT scheduling */
    return 0;
}

#else
/*
 * Routine(s) for establishing threads in RT FIFO scheduling mode using POSIX calls
 */
public int set_rt_scheduling(int p) {
    pthread_t me = pthread_self();
    struct sched_param pri;
    int mode;

    /* Attempt the operation */
    pri.sched_priority = p;
    if( pthread_setschedparam(me, SCHED_FIFO, &pri) < 0 ) {
        return -1; /* Failed for some reason */
    }

    /* Verify the operation */
    pri.sched_priority = -1;
    if( pthread_getschedparam(me, &mode, &pri) < 0
        || mode != SCHED_FIFO
        || pri.sched_priority != p ) { /* Didn't work, despite no errors... */
        errno = ENOSYS;
        return -1;
    }
}

```

Sep 12, 15 17:05

rtprio.c

Page 2/2

```

}

/* Successfully applied RT scheduling */
return 0;
}

#endif

/*
 * Routine to check we have the permitted capabilities needed for program operations
 *
 * The various threads need the following capabilities:
 *
 * CAP_IPC_LOCK (READER and WRITER) -- ability to mmap and mlock pages.
 * CAP_SYS_NICE (READER and WRITER) -- ability to set RT scheduling priorities
 * CAP_SYS_ADMIN (READER) -- ability to set (increase) the Comedi buffer maximum size
 * CAP_SYS_ADMIN (WRITER) -- ability to set RT IO scheduling priorities (unused at present)
 * CAP_SYS_ADMIN (TIDY) -- ability to set RT IO scheduling priorities (unused at present)
 *
 * Otherwise the MAIN thread and the TIDY thread need no special powers. The ZMQ IO thread
 * is also unprivileged, and is currently spawned during context creation from TIDY.
 */

public int check_permitted_capabilities_ok() {
    cap_t c = cap_get_proc();
    cap_flag_value_t v = CAP_CLEAR;

    if( !c )                /* No memory? */
        return -1;

    if( cap_get_flag(c, CAP_IPC_LOCK, CAP_PERMITTED, &v) < 0 || v == CAP_CLEAR ||
        cap_get_flag(c, CAP_SYS_NICE, CAP_PERMITTED, &v) < 0 || v == CAP_CLEAR ||
        cap_get_flag(c, CAP_SYS_ADMIN, CAP_PERMITTED, &v) < 0 || v == CAP_CLEAR
        ) {
        cap_free(c);
        errno = EPERM;
        return -1;
    }

    return 0;
}

```

Sep 12, 15 16:51

rtprio.h

Page 1/1

```
#
#ifndef _RTPRIO_H
#define _RTPRIO_H

#include "general.h"
#include <sys/capability.h>

/*
 * Routine(s) for establishing thread real-time scheduling
 */

export int set_rt_scheduling(int);
export int check_permitted_capabilities_ok();

#endif /* _RTPRIO_H */
```

Sep 06, 15 10:37

snapchat.c

Page 1/4

```

#
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <errno.h>
#include <string.h>
#include "assert.h"

#include <argtable2.h>
#include <zmq.h>

#include <getopt.h>

#include "util.h"
#include "param.h"
#include "argtab.h"

/*
 * Snapshot version
 */

#define PROGRAM_VERSION "1.0"
#define VERSION_VERBOSE_BANNER "MCLURS ADC toolset...\n"

/*
 * Global parameters for the snapshot program
 */

extern char *snapshot_addr;

param_t globals[] = {
    { "snapshot", "ipc://snapshot-CMD", &snapshot_addr, param_type_string, PARAM_SRC_ENV|PARAM_SRC_ARG,
      "address of snapshot command socket"
    },
};

const int n_global_params = (sizeof(globals)/sizeof(param_t));

/*
 * Debugging print out control
 */

int verbose = 0;
char *program = NULL;

/* Command line syntax options */

struct arg_lit *hl, *vn1, *v1, *q1;
struct arg_end *el;

BEGIN_CMD_SYNTAX(help) {
    v1 = arg_litn("v", "verbose", 0, 2, "Increase verbosity"),
    q1 = arg_lit0("q", "quiet", "Decrease verbosity"),
    hl = arg_lit0("h", "help", "Print usage help message"),
    vn1 = arg_lit0(NULL, "version", "Print program version string"),
    el = arg_end(20)
} APPLY_CMD_DEFAULTS(help) {
    /* No defaults to apply here */
} END_CMD_SYNTAX(help)

struct arg_lit *v2, *q2;
struct arg_end *e2;
struct arg_str *u2;
struct arg_str *n2;
struct arg_str *m2;

BEGIN_CMD_SYNTAX(main) {
    v2 = arg_litn("v", "verbose", 0, 3, "Increase verbosity"),
    q2 = arg_lit0("q", "quiet", "Decrease verbosity"),
    u2 = arg_str0("s", "snapshot", "<url>", "URL of snapshotter command socket"),
    m2 = arg_str0("m", "multi", "<prefix>", "Send multiple messages if replies begin with <prefix>"),
    n2 = arg_strn(NULL, NULL, "<args>", 1, 30, "Message content"),
    e2 = arg_end(20)
} APPLY_CMD_DEFAULTS(main) {

```

Sep 06, 15 10:37

snapchat.c

Page 2/4

```

m2->hdr.flag |= ARG_HASOPTVALUE;
m2->sval[0] = "";
INCLUDE_PARAM_DEFAULTS(globals, n_global_params); /* Use defaults from parameter table */
} END_CMD_SYNTAX(main);

/* Standard help routines: display the version banner */
void print_version(FILE *fp, int verbosity) {
    fprintf(fp, "%s Vn.%s\n", program, PROGRAM_VERSION);
    if(verbosity > 0) { /* Verbose requested... */
        fprintf(fp, VERSION_VERBOSE_BANNER);
    }
}

/* Standard help routines: display the usage summary for a syntax */
void print_usage(FILE *fp, void **argtable, int verbosity, char *program) {
    if( !verbosity ) {
        fprintf(fp, "Usage: %s ", program);
        arg_print_syntax(fp, argtable, "\n");
        return;
    }
    if( verbosity ) {
        char *suffix = verbosity>1? "\n\n" : "\n";
        fprintf(fp, "Usage: %s ", program);
        arg_print_syntaxv(fp, argtable, suffix);
        if( verbosity > 1 )
            arg_print_glossary(fp, argtable, "%-25s %s\n");
    }
}

/*
 * Snapchat globals...
 */

void      *zmq_main_ctx;      /* ZMQ context for messaging */
char      *snapshot_addr;    /* The URL of the snapshotter */

/*
 * Print a reply message to stdout
 */

int print_message(char *msg, int size) {
    if( msg[size-1] != '\n' ) {
        msg[size] = '\n';
        fwrite(msg, size+1, 1, stdout);
    }
    else {
        fwrite(msg, size, 1, stdout);
    }
    fflush(stdout);
}

/*
 * Return true if the string p is an initial prefix of str
 */

int checked_prefix(const char *p, const char *str) {
    while(*p && *str && *p == *str) {
        if( *p != *str ) /* Mismatch with prefix */
            return 0;
        p++, str++;
    }
    return *p? 0 : 1; /* True iff prefix has run out */
}

/*
 * Main entry point
 */

#define LOGBUF_SIZE      1024

int main(int argc, char *argv[], char *envp[]) {
    const char *prefix = NULL;
    char      buf[LOGBUF_SIZE];
    void      *snapshot;
    param_t   *p;

```

Sep 06, 15 10:37

snapchat.c

Page 3/4

```

int         ret, n;

program = argv[0];

/* Set up the standard parameters */
/* 1. Process parameters: internal default, then environment. */
set_param_from_env(envp, globals, n_global_params);

/* 2. Process parameters: push values out to program globals */
ret = assign_all_params(globals, n_global_params);
assertv(ret == 0, "Push parameters failed on param %d out of %d\n", -ret, n_global_params);

/* 3. Create and parse the command lines -- installs defaults from parameter table */
void **cmd_help = arg_make_help();
void **cmd_main = arg_make_main();

/* Try first syntax */
int err_help = arg_parse(argc, argv, cmd_help);
if( !err_help ) { /* Assume this was the desired command syntax */
    if(vn1->count)
        print_version(stdout, vn1->count);
    if(hl->count || !vn1->count) {
        int verbose = vn1->count - ql->count;
        print_usage(stdout, cmd_help, verbose>0, program);
        print_usage(stdout, cmd_main, verbose, program);
    }
    exit(0);
}

/* Try second syntax */
int err_main = arg_parse(argc, argv, cmd_main);
verbose = v2->count - q2->count;
if( err_main ) { /* This is the default desired syntax; give full usage */
    arg_print_errors(stderr, e2, program);
    print_usage(stderr, cmd_help, verbose>0, program);
    print_usage(stderr, cmd_main, verbose, program);
    exit(1);
}

/* 4. Process parameters: copy argument values back through the parameter table */
ret = arg_results_to_params(cmd_main, globals, n_global_params);

/* 5. Process parameters: deal with non-parameter table arguments where necessary */
if(m2->count) { /* Repeat-mode with prefix */
    prefix = m2->sval[0];
}

if(verbose > 2) /* Dump global parameters for debugging purposes */
    debug_params(stderr, globals, n_global_params);

/* Create the ZMQ contexts */
zmq_main_ctx = zmq_ctx_new();
if( !zmq_main_ctx ) {
    fprintf(stderr, "%s: Error -- ZeroMQ context creation failed: %s\n", program, strerror(errno));
    exit(2);
}

/* Create the socket to talk to the snapshot program */
snapshot = zh_connect_new_socket(zmq_main_ctx, ZMQ_REQ, snapshot_addr);
if( snapshot == NULL ) {
    fprintf(stderr, "%s: Error -- unable to create socket to snapshot at %s: %s\n",
        program, snapshot_addr, strerror(errno));
    zmq_ctx_term(zmq_main_ctx);
    exit(2);
}

const char **msg = n2->sval;
int         parts = n2->count;

if(prefix && verbose > 0)
    fprintf(stderr, "Sending %d parts in multi-message mode with reply prefix '%s'\n", parts, prefix);

do {
    int used, left;

```

Sep 06, 15 10:37

snapchat.c

Page 4/4

```

/* Send the message, wait for the reply; data is in arg_str *n2 */
if(verbose > 0)
    fprintf(stderr, "Sending message to %s...\n", snapshot_addr);

if(verbose > 1)
    fprintf(stderr, "Build:");

if( !prefix ) {
    used = 0;
    left = LOGBUF_SIZE-1;
    for(n=0; n<parts; n++) {
        int len;

        len = snprintf(&buf[used], left, "%s ", msg[n]);
        if(len >= left) {
            len=left;
            fprintf(stderr, "%s: ran out of space composing message '%s'\n", program, &buf[0]);
            exit(2);
        }
        if(verbose > 1)
            fprintf(stderr, " [%s]", buf);
        used += len;
        left -= len;
    }
    if(used)
        used--;
}
else {
    used = snprintf(&buf[0], LOGBUF_SIZE-1, "%s", *msg++);
    parts--;
    if(verbose > 1)
        fprintf(stderr, " [%s]", buf);
}

if(verbose > 1)
    fprintf(stderr, "\n");

/* Send the message, omit the final null */
ret = zh_put_msg(snapshot, 0, used, buf);
if( ret < 0 ) {
    fprintf(stderr, "\n%s: Error -- sending message failed: %s\n", program, strerror(errno));
    zmq_close(snapshot);
    zmq_ctx_term(zmq_main_ctx);
    exit(3);
}

/* Wait for reply */
if(verbose > 0)
    fprintf(stderr, "Awaiting reply from %s...\n", snapshot_addr);
used = zh_collect_multi(snapshot, &buf[0], LOGBUF_SIZE-1, "");
buf[LOGBUF_SIZE-1] = '\0';
if(verbose >= 0)
    print_message(&buf[0], used);
} while( prefix && parts > 0 && checked_prefix(prefix, &buf[0]) );

/* Clean up ZeroMQ sockets and context */
zmq_close(snapshot);
zmq_ctx_term(zmq_main_ctx);
exit(0);
}

```


Sep 22, 15 16:12

snapshot.c

Page 1/12

```

#
#include "general.h"

#define _GNU_SOURCE    /* Linux-specific code below (O_PATH) */

#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/capability.h>
#include <sys/prctl.h>
#include <fcntl.h>
#include <pwd.h>
#include <grp.h>
#include <signal.h>
#include <argtable2.h>
#include "argtab.h"

#include <zmq.h>
#include <pthread.h>

#include <unistd.h>
#include <errno.h>
#include <string.h>
#include "assert.h"
#include <sched.h>

#include <comedi.h>
#include <comedilib.h>

#include "util.h"
#include "param.h"
#include "queue.h"
#include "strbuf.h"
#include "chunk.h"
#include "rtprio.h"
#include "snapshot.h"
#include "reader.h"
#include "writer.h"
#include "tidy.h"

/*
 * Snapshot version
 */

#define PROGRAM_VERSION "1.1"
#define VERSION_VERBOSE_BANNER "MCLURS ADC toolset...\n"

/*
 * Global parameters for the snapshot program
 */

public int die_die_die_now = 0;

import rparams    reader_parameters;
import wparams    writer_parameters;
import const char *tmpdir_path;
private const char *snapshot_addr;
private const char *snapshot_user;
private const char *snapshot_group;
private int        schedprio;

public param_t globals[] = {
    { "tmpdir",    "tmp",
      &tmpdir_path,
      PARAM_TYPE(string), PARAM_SRC_ENV|PARAM_SRC_ARG,
      "directory for creation of temporary files"
    },
    { "freq",      "312.5e3",
      &reader_parameters.r_frequency,
      PARAM_TYPE(double), PARAM_SRC_ENV|PARAM_SRC_ARG|PARAM_SRC_CMD,
      "sampling frequency (divided by 8) of the ADC [Hz]"
    },
    { "snapshot",  "ipc://snapshot-CMD",

```

Sep 22, 15 16:12

snapshot.c

Page 2/12

```

&snapshot_addr,
PARAM_TYPE(string), PARAM_SRC_ENV|PARAM_SRC_ARG,
"address of snapshot command socket"
},
{
"snapdir", "snap",
&writer_parameters.w_snapdir,
PARAM_TYPE(string), PARAM_SRC_ENV|PARAM_SRC_ARG,
"directory where samples are written"
},
{
"dev", "/dev/comedi0",
&reader_parameters.r_device,
PARAM_TYPE(string), PARAM_SRC_ENV|PARAM_SRC_ARG,
"the Comedi device to open"
},
{
"range", "750",
&reader_parameters.r_range,
PARAM_TYPE(int32), PARAM_SRC_ENV|PARAM_SRC_ARG|PARAM_SRC_CMD,
"the ADC converter full-scale range [mV]"
},
{
"bufsz", "32",
&reader_parameters.r_bufsz,
PARAM_TYPE(int32), PARAM_SRC_ENV|PARAM_SRC_ARG|PARAM_SRC_CMD,
"size of the Comedi buffer [MiB]"
},
{
"window", "10",
&reader_parameters.r_window,
PARAM_TYPE(double), PARAM_SRC_ENV|PARAM_SRC_ARG|PARAM_SRC_CMD,
"guaranteed window in the ring buffer [s]"
},
{
"bufhwm", "0.9",
&reader_parameters.r_buf_hwm_fraction,
PARAM_TYPE(double), PARAM_SRC_ENV|PARAM_SRC_ARG|PARAM_SRC_CMD,
"ring buffer high-water mark fraction"
},
{
"rtprio", NULL,
&schedprio,
PARAM_TYPE(int32), PARAM_SRC_ENV|PARAM_SRC_ARG,
"priority of real-time threads [0-99]"
},
{
"rdprio", NULL,
&reader_parameters.r_schedprio,
PARAM_TYPE(int32), PARAM_SRC_ENV|PARAM_SRC_ARG,
"priority of real-time reader thread [0-99]"
},
{
"wrprio", NULL,
&writer_parameters.w_schedprio,
PARAM_TYPE(int32), PARAM_SRC_ENV|PARAM_SRC_ARG,
"priority of real-time writer thread [0-99]"
},
{
"user", NULL,
&snapshot_user,
PARAM_TYPE(string), PARAM_SRC_ENV|PARAM_SRC_ARG,
"user/UID for file system access and creation"
},
{
"group", NULL,
&snapshot_group,
PARAM_TYPE(string), PARAM_SRC_ENV|PARAM_SRC_ARG,
"group/GID for file system access and creation"
},
{
"ram", "64",
&writer_parameters.w_lockedram,
PARAM_TYPE(int32), PARAM_SRC_ENV|PARAM_SRC_ARG,
"amount of data RAM to lock [MiB]"
},
{
"wof", "0.5",
&writer_parameters.w_writeahead,
PARAM_TYPE(double), PARAM_SRC_ENV|PARAM_SRC_ARG,
"write overbooking fraction"
},
{
"chunk", "1024",
&writer_parameters.w_chunksize,
PARAM_TYPE(int32), PARAM_SRC_ENV|PARAM_SRC_ARG,
"size of a transfer chunk [KiB]"
},
},
};

```

Sep 22, 15 16:12

snapshot.c

Page 3/12

```

public const int n_global_params =      (sizeof(globals)/sizeof(param_t));

/*
 * Debugging print out control
 */

public int   verbose;
public char *program   = NULL;

/* Command line syntax options -- there are no mandatory arguments on the main command line! */

private struct arg_lit *h1, *vn1, *v1, *q1;
private struct arg_end *e1;

BEGIN_CMD_SYNTAX(help) {
    v1 = arg_litn("v", "verbose", 0, 3, "Increase verbosity"),
    q1 = arg_lit0("q", "quiet", "Decrease verbosity"),
    h1 = arg_lit0("h", "help", "Print usage help message"),
    vn1 = arg_lit0(NULL, "version", "Print program version string"),
    e1 = arg_end(20)
} APPLY_CMD_DEFAULTS(help) {
    /* No defaults to apply here */
} END_CMD_SYNTAX(help)

private struct arg_lit *v2, *q2;
private struct arg_end *e2;

BEGIN_CMD_SYNTAX(main) {
    v2 = arg_litn("v", "verbose", 0, 3, "Increase verbosity"),
    q2 = arg_lit0("q", "quiet", "Decrease verbosity"),
    arg_str0("s", "snapshot", "<url>", "URL of snapshotter command socket"),
    arg_str0(NULL, "tmpdir", "<path>", "Path to temporary directory"),
    arg_str0("S", "snapdir", "<path>", "Path to samples directory"),
    arg_dbl0("f", "freq", "<real>", "Per-channel sampling frequency [Hz]"),
    arg_dbl0("w", "window", "<real>", "Min. capture window length [s]"),
    arg_dbl0("B", "bufhwm", "<real>", "Ring buffer High-water mark fraction"),
    arg_str0("d", "dev", "<path>", "Comedi device to use"),
    arg_int0("P", "rtprio", "<1-99>", "Common thread RT priority"),
    arg_int0("R", "rdprio", "<1-99>", "Reader thread RT priority"),
    arg_int0("W", "wrprio", "<1-99>", "Writer thread RT priority"),
    arg_str0("u", "user", "<uid/name>", "User to run as"),
    arg_str0("g", "group", "<gid/name>", "Group to run as"),
    arg_int0("b", "bufsz", "<int>", "Comedi ring buffer Size [MiB]"),
    arg_int0("m", "ram", "<int>", "Data Transfer RAM size [MiB]"),
    arg_int0("r", "range", "<int>", "ADC full-scale range [mV]"),
    arg_int0("c", "chunk", "<int>", "File transfer chunk size [kiB]"),
    arg_dbl0("W", "wof", "<real>", "Write Overbooking Fraction"),
    e2 = arg_end(20)
} APPLY_CMD_DEFAULTS(main) {
    INCLUDE_PARAM_DEFAULTS(globals, n_global_params);
} END_CMD_SYNTAX(main)

/* Standard help routines: display the version banner */
private void print_version(FILE *fp, int verbosity) {
    fprintf(fp, "%s: Vn.%s\n", program, PROGRAM_VERSION);
    if(verbosity > 0) {
        /* Verbose requested... */
        fprintf(fp, VERSION_VERBOSE_BANNER);
    }
}

/* Standard help routines: display the usage summary for a syntax */
private void print_usage(FILE *fp, void **argtable, int verbosity, char *program) {
    if( !verbosity ) {
        fprintf(fp, "Usage: %s ", program);
        arg_print_syntax(fp, argtable, "\n");
        return;
    }
    if( verbosity ) {
        char *suffix = verbosity>1? "\n\n" : "\n";
        fprintf(fp, "Usage: %s ", program);
        arg_print_syntaxv(fp, argtable, suffix);
        if( verbosity > 1 )
            arg_print_glossary(fp, argtable, "%-25s %s\n");
    }
}

```

Sep 22, 15 16:12

snapshot.c

Page 4/12

```

}

/*
 * Snapshot globals for this file.
 */

private const char *snapshot_addr = NULL; /* The address of the main command socket */
private const char *snapshot_user = NULL; /* The user we should run as, after startup */
private const char *snapshot_group = NULL; /* The group to run as, after startup */
private int schedprio; /* Real-time priority for reader and writer */

/*
 * Snapshot globals shared between threads
 */

public void snapshot_zmq_ctx; /* ZMQ context for messaging -- created by the TIDY thread */

public int tmpdir_dirfd; /* The file descriptor obtained for the TMPDIR directory */
public const char *tmpdir_path; /* The path for the file descriptor above */

/*
 * Thread handles for reader and writer
 */

private pthread_t reader_thread,
writer_thread,
tidy_thread;

private pthread_attr_t reader_thread_attr,
writer_thread_attr,
tidy_thread_attr;

/*
 * Establish main comms: this routine runs last, so it mostly does connect() calls.
 * It must run when the other three threads are already active.
 */

private void *log_socket; /* N.B. This socket is opened by the TIDY thread, but not used there */
private void *reader;
private void *writer;
private void *command;

private int create_main_comms() {
    int ret;

    /* Create and initialise the sockets: reader and writer command sockets */
    reader = zh_connect_new_socket(snapshot_zmq_ctx, ZMQ_REQ, READER_CMD_ADDR);
    if( reader == NULL ) {
        fprintf(stderr, "%s: Error -- unable to connect internal socket to reader: %s\n", program, strerror(errno));
        return -1;
    }
    writer = zh_connect_new_socket(snapshot_zmq_ctx, ZMQ_REQ, WRITER_CMD_ADDR);
    if( writer == NULL ) {
        fprintf(stderr, "%s: Error -- unable to connect internal socket to writer: %s\n", program, strerror(errno));
        return -1;
    }

    /* Create and initialise the external command socket */
    command = zh_bind_new_socket(snapshot_zmq_ctx, ZMQ_REP, snapshot_addr);
    if( command == NULL ) {
        fprintf(stderr, "%s: Error -- unable to bind external command socket %s: %s\n",
            program, snapshot_addr, strerror(errno));
        return -1;
    }

    return 0;
}

/* Close everything created above */

private void close_main_comms() {
    zmq_close(reader);
    zmq_close(writer);
    zmq_close(command);
}

```

Sep 22, 15 16:12

snapshot.c

Page 5/12

```

/*
 * Sort out the capabilities required by the process. (If not running
 * as root, check that we have the capabilities we require.) Release
 * any capabilities not needed and lock against dropping privilege.
 *
 * The threads need the following capabilities:
 *
 * CAP_IPC_LOCK (Reader and Writer) -- ability to mmap and mlock pages.
 * CAP_SYS_NICE (Reader and Writer) -- ability to set RT scheduling priorities
 * CAP_SYS_ADMIN (Reader) -- ability to set (increase) the Comedi buffer maximum size
 * CAP_SYS_ADMIN (Writer) -- ability to set RT IO scheduling priorities (unused at present)
 * CAP_SYS_ADMIN (Tidy) -- ability to set RT IO scheduling priorities (unused at present)
 *
 * CAP_SETUID (Main)
 * CAP_SETGID (Main) -- ability to change user ID
 *
 * Otherwise the main thread and the tidy thread need no special powers. The ZMQ IO thread
 * is also unprivileged, and is currently spawned during context creation from tidy.
 */

private int snap_adjust_capabilities() {
    cap_t c = cap_get_proc();
    uid_t u = geteuid();
    int ret = 0;

    if( !c ) /* No memory? */
        return -1;

    if( check_permitted_capabilities_ok() < 0 ) {
        fprintf(stderr, "%s: Error --- I do not have the necessary capabilities to operate\n", program);
        return -1;
    }

    if( !u ) {
        const cap_value_t vs[] = { CAP_IPC_LOCK, CAP_SYS_NICE, CAP_SYS_ADMIN, CAP_SETUID, CAP_SETGID, };

        /* So we are root and have the capabilities we need. Prepare to drop the others... */
        /* Keep the EFFECTIVE capabilities as long as we stay root */
        cap_clear(c);
        cap_set_flag(c, CAP_PERMITTED, sizeof(vs)/sizeof(cap_value_t), &vs[0], CAP_SET);
        cap_set_flag(c, CAP_EFFECTIVE, sizeof(vs)/sizeof(cap_value_t), &vs[0], CAP_SET);
        if( prctl(PR_SET_KEEPCAPS, 1L) < 0 ) {
            cap_free(c);
            fprintf(stderr, "%s: Error --- unable to keep required capabilities on user change\n", program);
            return -1;
        }

        ret = cap_set_proc(c);
    }

    cap_free(c);
    return ret;
}

/*
 * Drop privileges and capabilities when appropriate.
 */

private int main_adjust_capabilities(uid_t uid, gid_t gid) {
    cap_t c = cap_get_proc();
    const cap_value_t vs[] = { CAP_SETUID, CAP_SETGID, };

    /* Drop all capabilities except CAP_SETUID/GID from effective set */

    if(c) {
        cap_clear_flag(c, CAP_EFFECTIVE);
        cap_set_flag(c, CAP_EFFECTIVE, sizeof(vs)/sizeof(cap_value_t), &vs[0], CAP_SET);
        if( cap_set_proc(c) < 0 ) {
            cap_free(c);
            fprintf(stderr, "%s: Error --- MAIN thread fails to clear capabilities: %s\n", program, strerror(errno));
            return -1;
        }
        cap_free(c);
    }
}

```

Sep 22, 15 16:12

snapshot.c

Page 6/12

```

/* Drop all user and group privileges: set all uids to uid and all gids to gid */
/* Complain if that fails -- we were not root and uid/gid were not in our set */
if( setresgid(gid, gid, gid) < 0 ) {
    fprintf(stderr, "%s: Error -- MAIN thread unable to change to gid %d: %s\n", program, gid, strerror(errno));
    return -1;
}
if( setresuid(uid, uid, uid) < 0 ) {
    fprintf(stderr, "%s: Error -- MAIN thread unable to change to uid %d: %s\n", program, uid, strerror(errno));
    return -1;
}

c = cap_get_proc();
if(c) {
    cap_set_flag(c, CAP_PERMITTED, sizeof(vs)/sizeof(cap_value_t), &vs[0], CAP_CLEAR);
    if( cap_set_proc(c) < 0 ) {
        cap_free(c);
        fprintf(stderr, "%s: Error -- MAIN thread keeps setuid/gid capabilities: %s\n", program, strerror(errno));
        return -1;
    }
    cap_free(c);
}

/* Now check we still have the required permitted capabilities */
if( check_permitted_capabilities_ok() < 0 ) {
    fprintf(stderr, "%s: Error -- MAIN thread lost capabilities on changing user!\n", program);
    return -1;
}

return 0;
}

/*
 * Deal nicely with the interrupt signal.
 * Basically, the signal sets the die_die_die_now flag which the various threads notice.
 * CURRENTLY NOT WORKING PROPERLY SO DISABLED
 */

private void intr_handler(int i) {
    die_die_die_now++;
}

private int set_intr_sig_handler() {
    struct sigaction a;

    bzero(&a, sizeof(a));
    a.sa_handler = intr_handler;
    if( sigaction(SIGINT, &a, NULL) < 0 ) {
        fprintf(stderr, "%s: Error -- unable to install INT signal handler: %s\n", program, strerror(errno));
        return -1;
    }
    return 0;
}

/*
 * Process a (possibly multipart) log message.
 * Collect the various pieces and write to stderr
 */

#define LOGBUF_SIZE    MSGBUFSIZE
private char pfx[] = "Log: ";

private int process_log_message(void *s) {
    char log_buffer[MSGBUFSIZE];
    int used;

    memcpy(&log_buffer[0], &pfx[0], sizeof(pfx));
    used = sizeof(pfx)-1;
    used += zh_collect_multi(s, &log_buffer[used], LOGBUF_SIZE-1, "");
    if( log_buffer[used-1] != '\n' ) {
        log_buffer[used] = '\n';
        fwrite(log_buffer, used+1, 1, stderr);
    }
    else {
        fwrite(log_buffer, used, 1, stderr);
    }
}

```

Sep 22, 15 16:12

snapshot.c

Page 7/12

```

}
fflush(stderr);
return 0;
}

/*
 * Handle replies from READER and WRITER threads. The reply message
 * is a pointer to a set of error strbufs. We collect and join all
 * the strings in the reply buffer. The collector maintains as
 * invariant that "used==0 || reply_buffer[used-1] is not NUL" and that
 * "b == &reply_buffer[used]".
 */

#define REPLY_BUFSIZE MSGBUFSIZE
private char reply_buffer[REPLY_BUFSIZE];

private int process_reply(void *s) {
    int size;
    strbuf err;
    char *b = &reply_buffer[0];
    int used;

    size = zh_get_msg(s, 0, sizeof(strbuf), (void *)&err);
    assertv(size==sizeof(err), "Reply message of wrong size %d\n", size);

    /* Establish invariants */
    *b = '\0'; used = 0;

    /* Traverse the strbuf chain once collecting data, then release */
    for_nxt_in_Q(queue *q, strbuf2qp(err), (queue *)NULL)
        strbuf s = qp2strbuf(q);
        int n = strbuf_used(s);
        if(n) {
            strbuf_revert(s); /* Empty strbuf, nothing to do */
            if(n > REPLY_BUFSIZE-used) { /* Remove any internal NUL characters */
                n = REPLY_BUFSIZE-used-1; /* There is too much data */
                /* We can manage this much of it */
            }
            /* fprintf(stderr, "strbuf %p, used %d, ptr %p, string '%s'\n",
             * s, n, b, strbuf_string(s));
             * memcpy(b, strbuf_string(s), n); /* Copy the data */
             * b += n; used += n; /* Now we have used this much space */
             * while( b[-1] == '\0' ) b--,used--; /* Skip back over any NULs */
             * fprintf(stderr, "strbuf %p, ptr now %p, total used now %d\n", s, b, used);
            */
        }
    end_for_nxt;

    release_strbuf(err); /* Free the entire link of strbufs */

    if( b[-1] != '\n' ) /* Ensure final newline */
        *b = '\n';

    /* Send the complete reply */
    used = b - &reply_buffer[0];
    zh_put_msg(command, 0, used, &reply_buffer[0]);
    return 0;
}

/*
 * Handle commands sent to the snapshotter. These are forwarded
 * either to the reader thread or the writer thread, and their replies
 * are returned to the originator. Using the REP socket ensures only
 * one outstanding message is in process, so simplifies the reply routing.
 */

private int process_snapshot_command() {
    strbuf c,e; /* Command and Error buffers */
    char *buf;
    int size, ret;
    int fwd;

    c = alloc_strbuf(2);
    e = strbuf_next(c);

    buf = strbuf_string(c);
    size = zh_get_msg(command, 0, strbuf_space(c), buf);

```

Sep 22, 15 16:12

snapshot.c

Page 8/12

```

if( !size ) {
    ret = zh_put_msg(command, 0, 0, NULL); /* If empty message received, send empty reply at once */
    release_strbuf(c);
    assertv(ret == 0, "Reply to command failed,%d\n", ret);
    return 0;
}
strbuf_setpos(c, size);
buf[size] = '\0';
// fprintf(stderr, "Msg '%c' (%d)\n", buf[0], buf[0]);
fwd = 0;
switch(buf[0]) {
case 'q':
case 'Q':
    /* Deal specially with Quit command, to close down nicely... */
    ret = zh_put_msg(reader, 0, 0, NULL); /* Forward zero length message to the READER thread */
    assertv(ret == 0, "Quit to READER failed,%d\n", ret);
    ret = zh_put_msg(writer, 0, 0, NULL); /* Forward zero length message to the WRITER thread */
    assertv(ret == 0, "Quit to WRITER failed,%d\n", ret);
    ret = zh_put_msg(command, 0, 7, "OK Quit"); /* Reply to Quit here */
    assertv(ret == 7, "Quit reply failed,%d\n", ret);
    break;

case 'g':
case 'G':
case 'h':
case 'H':
case 'i':
case 'I':
case 'p':
case 'P':
    /* Forward these commands to the READER thread */
    ret = zh_put_msg(reader, 0, sizeof(strbuf), (void *)&c);
    assertv(ret == sizeof(c), "Forward to READER failed,%d\n", ret);
    fwd++;
    break;

case 'd':
case 'D':
case 's':
case 'S':
case 'z':
case 'Z':
    /* Forward snapshot and dir commands to WRITER */
    ret = zh_put_msg(writer, 0, sizeof(strbuf), (void *)&c);
    assertv(ret == sizeof(c), "Forward to WRITER failed,%d\n", ret);
    fwd++;
    break;

case '?':
    buf[0] = '!';
    ret = zh_put_msg(command, 0, size, buf); /* Reply to 'ping' message */
    assertv(ret > 0, "Reply to ping failed,%d\n", ret);
    break;

default:
    strbuf_printf(e, "NO: Unknown command: '%s'\n", buf);
    fprintf(stderr, "%s: %s", program, strbuf_string(e));
    ret = zh_put_msg(command, 0, strbuf_used(e), strbuf_string(e));
    assertv(ret == strbuf_used(e), "Reject unknown reply failed,%d\n", ret);
    break;
}
if( !fwd )
    release_strbuf(c);
return 0;
}

/*
 * MAIN thread message loop
 */

#define MAIN_LOOP_POLL_INTERVAL 20

private void main_thread_msg_loop() { /* Read and process messages */
    int poll_delay;
    int running;
    zmq_pollitem_t poll_list[] =

```


Sep 22, 15 16:12

snapshot.c

Page 9/12

```

    { { log_socket, 0, ZMQ_POLLIN, 0 },
      { command, 0, ZMQ_POLLIN, 0 },
      { reader, 0, ZMQ_POLLIN, 0 },
      { writer, 0, ZMQ_POLLIN, 0 },
    };
#define N_POLL_ITEMS (sizeof(poll_list)/sizeof(zmq_pollitem_t))
int (*poll_responders[N_POLL_ITEMS])(void *) =
{ process_log_message,
  process_snapshot_command,
  process_reply,
  process_reply,
};

fprintf(stderr, "Log: starting MAIN thread polling loop with %d items\n", N_POLL_ITEMS);
running = true;
poll_delay = MAIN_LOOP_POLL_INTERVAL;
while(running && !die_die_die_now) {
    int n;
    int ret = zmq_poll(&poll_list[0], N_POLL_ITEMS, poll_delay);

    if( ret < 0 && errno == EINTR ) { /* Interrupted */
        fprintf(stderr, "%s: MAIN thread loop interrupted\n", program);
        break;
    }
    if(ret < 0)
        break;
    running = reader_parameters.r_running || writer_parameters.w_running;
    if( !running ) /* Flush out last messages */
        poll_delay = 1000;
    for(n=0; n<N_POLL_ITEMS; n++) {
        if( poll_list[n].revents & ZMQ_POLLIN ) {
            ret = (*poll_responders[n])(poll_list[n].socket);
            assertv(ret >= 0, "Error in message processing in MAIN poll loop, ret %d\n", ret);
            running = true;
        }
    }
}
}
}

/*
 * Snapshot main routine.
 */

public int main(int argc, char *argv[], char *envp[]) {
    char *thread_return = NULL;
    int ret, running, poll_delay;
    char *cmd_addr;
    param_t *p;

    program = argv[0];

    /* Set up the standard parameters */
    /* 1. Process parameters: internal default, environment, then command-line argument. */
    set_param_from_env(envp, globals, n_global_params);

    /* 2. Process parameters: push values out to program globals */
    ret = assign_all_params(globals, n_global_params);
    assertv(ret == 0, "Push parameters failed on param %d out of %d\n", -ret, n_global_params);

    if(verbose > 2) {
        fprintf(stderr, "Params before cmdline...\n");
        debug_params(stderr, globals, n_global_params);
    }

    /* 3. Create and parse the command lines -- installs defaults from parameter table */
    void **cmd_help = arg_make_help();
    void **cmd_main = arg_make_main();

    /* Try first syntax -- reject empty command lines */
    int err_help = arg_parse(argc, argv, cmd_help);
    if( !err_help && (vnl->count || hl->count) ) { /* Assume this was the desired command syntax */
        if(vnl->count)
            print_version(stdout, vl->count);
        if(hl->count || !vnl->count) {
            print_usage(stdout, cmd_help, vl->count>0, program);
        }
    }
}

```

Sep 22, 15 16:12

snapshot.c

Page 10/12

```

    }
    print_usage(stdout, cmd_main, v1->count, program);
}
exit(0);
}

/* Try second syntax -- may be empty, means use default or environment variable parameters */
int err_main = arg_parse(argc, argv, cmd_main);
if( err_main ) {
    /* This is the default desired syntax; give full usage */
    arg_print_errors(stderr, e2, program);
    print_usage(stderr, cmd_help, v2->count>0, program);
    print_usage(stderr, cmd_main, v2->count, program);
    exit(1);
}

verbose = v2->count - q2->count;
if(verbose > 2) {
    fprintf(stderr, "Params before reverse pass...\n");
    debug_params(stderr, globals, n_global_params);
}

/* 4. Process parameters: copy argument values back through the parameter table */
ret = arg_results_to_params(cmd_main, globals, n_global_params);

/* 5. Process parameters: deal with non-parameter table arguments where necessary */

if(verbose > 1) {
    fprintf(stderr, "Params before checking...\n");
    debug_params(stderr, globals, n_global_params);
}

/* 5a. Verify parameters required by the main program/thread */
tmpdir_dirfd = open(tmpdir_path, O_PATH|O_DIRECTORY); /* Verify the TMPDIR path */
if( tmpdir_dirfd < 0 ) {
    fprintf(stderr, "%s: Error -- cannot access given TMPDIR '%s': %s\n", program, tmpdir_path, strerror(errno));
    exit(2);
}

/* Compute the UID and GID for unprivileged operation.
 *
 * If the GID parameter is set, use that for the group; if not, but
 * the UID parameter is set, get the group from that user and set
 * the uid from there too. If neither is set, use the real uid/gid
 * of the thread.
 */

gid_t gid = -1;
if(snapshot_group) {
    struct group *grp = getgrnam(snapshot_group);

    if(grp == NULL) {
        /* The group name was invalid */
        fprintf(stderr, "%s: Error -- given group %s is not recognised\n", program, snapshot_group);
        exit(2);
    }
    gid = grp->gr_gid;
}

uid_t uid = -1;
if(snapshot_user) { /* Got a UID value */
    struct passwd *pwd = getpwnam(snapshot_user);

    if(pwd == NULL) {
        /* The user name was invalid */
        fprintf(stderr, "%s: Error -- given user %s is not recognised\n", program, snapshot_user);
        exit(2);
    }

    uid = pwd->pw_uid; /* Use this user's UID */
    if(gid == -1)
        gid = pwd->pw_gid; /* Use this user's principal GID */
}
else {
    uid = getuid(); /* Use the real UID of this thread */
    gid = getgid(); /* Use the real GID of this thread */
}

/* 5b. Check capabilities and drop privileges */

```

Sep 22, 15 16:12

snapshot.c

Page 11/12

```

if( snap_adjust_capabilities() < 0 ) {
    exit(2);
}
if( main_adjust_capabilities(uid, gid) < 0 ) {
    exit(2);
}

/* Check the supplied parameters; WRITER must come first as READER needs chunk size */
strbuf e = alloc_strbuf(1); /* Catch parameter error diagnostics */

/* 5c. Verify and initialise parameters for the WRITER thread */
if( !writer_parameters.w_schedprio )
    writer_parameters.w_schedprio = schedprio;
strbuf_printf(e, "%s: Error -- WRITER Params:", program);
ret = verify_writer_params(&writer_parameters, e);
if( ret < 0 ) {
    fprintf(stderr, "%s\n", strbuf_string(e));
    exit(3);
}

/* 5d. Verify and initialise parameters for the READER thread */
if( !reader_parameters.r_schedprio )
    reader_parameters.r_schedprio = schedprio;
strbuf_printf(e, "%s: Error -- READER Params:", program);
ret = verify_reader_params(&reader_parameters, e);
if( ret < 0 ) {
    fprintf(stderr, "%s\n", strbuf_string(e));
    exit(3);
}

release_strbuf(e);

#if 0
/* Exit nicely on SIGINT: this is done by setting the die_die_die_now flag. */
if( set_intr_sig_handler() < 0 ) {
    exit(3);
}
#endif

/* Create the TIDY thread */
pthread_attr_init(&tidy_thread_attr);
if( pthread_create(&tidy_thread, &tidy_thread_attr, tidy_main, &log_socket) < 0 ) {
    fprintf(stderr, "%s: Error -- TIDY thread creation failed: %s\n", program, strerror(errno));
    exit(4);
}

/* Wait here for log_socket */
while( !die_die_die_now && !log_socket ) {
    usleep(10000);
}

if( !die_die_die_now ) {
    pthread_attr_init(&reader_thread_attr); /* Create the READER thread */
    if( pthread_create(&reader_thread, &reader_thread_attr, reader_main, NULL) < 0 ) {
        fprintf(stderr, "%s: Error -- READER thread creation failed: %s\n", program, strerror(errno));
        exit(4);
    }

    pthread_attr_init(&writer_thread_attr); /* Create the WRITER thread */
    if( pthread_create(&writer_thread, &writer_thread_attr, writer_main, NULL) < 0 ) {
        fprintf(stderr, "%s: Error -- WRITER thread creation failed: %s\n", program, strerror(errno));
        exit(4);
    }
}

/* Wait for the threads to establish comms etc. DON'T WAIT TOO LONG */
while( !die_die_die_now ) {
    usleep(10000); /* Wait for 10ms */
    if( reader_parameters.r_running && writer_parameters.w_running )
        break; /* Now ready to start main loop */
}

/* Run the MAIN thread sevice loop here */
if( create_main_comms() < 0 ) {
    die_die_die_now++;
}

```

Sep 22, 15 16:12

snapshot.c

Page 12/12

```
}
main_thread_msg_loop();

/* Clean up the various threads */
if(reader_thread) {
    if( pthread_join(reader_thread, (void *)&thread_return) < 0 ) {
        fprintf(stderr, "%s: Error -- READER thread join error: %s\n", program, strerror(errno));
        thread_return = NULL;
    }
    else {
        if( thread_return ) {
            fprintf(stderr, "Log: READER thread rejoined -- %s\n", thread_return);
            thread_return = NULL;
        }
    }
}
if(writer_thread) {
    if( pthread_join(writer_thread, (void *)&thread_return) < 0 ) {
        fprintf(stderr, "%s: Error -- WRITER thread join error: %s\n", program, strerror(errno));
        thread_return = NULL;
    }
    else {
        if( thread_return ) {
            fprintf(stderr, "Log: WRITER thread rejoined -- %s\n", thread_return);
            thread_return = NULL;
        }
    }
}

if( pthread_join(tidy_thread, (void *)&thread_return) < 0 ) {
    fprintf(stderr, "%s: Error -- TIDY thread join error: %s\n", program, strerror(errno));
    thread_return = NULL;
}
else {
    if( thread_return ) {
        fprintf(stderr, "Log: TIDY thread rejoined -- %s\n", thread_return);
        thread_return = NULL;
    }
}

/* Clean up our ZeroMQ sockets */
close_main_comms();

/* These were created by the TIDY thread */
zmq_close(log_socket);
zmq_ctx_term(snapshot_zmq_ctx);
exit(0);
}
```

Sep 17, 15 17:46

snapshot.h

Page 1/1

```
#
#include "general.h"
/* Shared globals */
export void      *zmq_main_ctx;
export param_t    globals[];
export const int  n_global_params;
export int        verbose;
export int        die_die_die_now;
export int        tmpdir_dirfd;
export const char *tmpdir_path;
/* Common definitions */
#define LOG_SOCKET      "inproc://Main-LOG"
#define MSGBUFSIZE      8192
```

Sep 17, 15 17:34

strbuf.c

Page 1/4

```
#
#include "general.h"
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include "strbuf.h"
#include "queue.h"

#define N_STRBUF_ALLOC 8          /* Allocate this many buffers at one go */
#define MAX_STRBUF_SIZE (512-sizeof(queue)) /* Strbuf will hold 496 characters maximum */

struct _strbuf {
    queue s_Q;          /* Queue header to avoid malloc() calls */
    int s_used;         /* Pointer to next free space in buffer */
    char s_buffer[MAX_STRBUF_SIZE]; /* Buffer space when in use */
};

/*
 * Return the usable string space in a strbuf.
 */

public int strbuf_space(strbuf s) {
    return MAX_STRBUF_SIZE;
}

/*
 * Allocate and free strbufs, using a queue to avoid excessive malloc()
 */

private QUEUE_HEADER(sbufQ);
private int N_in_Q = 0;

public strbuf alloc_strbuf(int nr) {
    queue *ret;

    if( N_in_Q < nr ) { /* The queue doesn't have enough */
        int n;

        for(n=0; n<N_STRBUF_ALLOC; n++) {
            queue *q = (queue *)calloc(1, sizeof(struct _strbuf));

            if( !q ) { /* Allocation failed */
                if( N_in_Q >= nr )
                    break; /* But we have enough now anyway */
                return NULL;
            }
            init_queue(q);
            queue_ins_after(&sbufQ, q);
            N_in_Q++;
        }
        ret = de_queue(queue_next(&sbufQ));
        while(--nr > 0) { /* Collect enough to satisfy request */
            queue *p = de_queue(queue_next(&sbufQ));

            init_queue(p);
            ((strbuf)p->s_used = 0;
            ((strbuf)p->s_buffer[0] = '\0';
            queue_ins_before(ret, p);
        }
        return (strbuf)ret;
    }
}

private void free_strbuf(strbuf s) {
    free( (void *)s );
}

public void release_strbuf(strbuf s) {
    queue *p;

    while( (p = de_queue(queue_next(&s->s_Q))) != NULL ) {
        init_queue(p);
        queue_ins_before(&sbufQ, p);
    }
}
```

Sep 17, 15 17:34

strbuf.c

Page 2/4

```

    N_in_Q++;
}
queue_ins_before(&sbufQ, &s->s_Q);
N_in_Q++;
}

/*
 * Get the string pointer from an strbuf (since the latter is opaque, we need a function for this).
 */

public char *strbuf_string(strbuf s) {
    return &s->s_buffer[0];
}

/*
 * Return the number of characters printed into a strbuf so far
 */

public int strbuf_used(strbuf s) {
    return s->s_used;
}

/*
 * Mark the current used position.
 */

public int strbuf_setpos(strbuf s, int pos) {
    if( !s ) {
        errno = EINVAL;
        return -1;
    }
    if(pos < 0 || pos > MAX_STRBUF_SIZE) {
        errno = ERANGE;
        return -1;
    }
    int used = s->s_used;
    s->s_used = pos;
    return used;
}

/*
 * Do a formatted print into an strbuf, starting at pos.
 */

private int strbuf_vprintf(strbuf s, int pos, const char *fmt, va_list ap) {
    int rest;
    int used;
    char *buf;

    if(pos < 0)
        pos = s->s_used ? s->s_used : 0;
    buf = &s->s_buffer[pos];
    rest = MAX_STRBUF_SIZE - pos; /* There should be this much space remaining */
    if(rest < 0) {
        errno = EINVAL;
        return -1;
    }
    used = vsnprintf(buf, rest, fmt, ap);
    s->s_used = used >= rest ? MAX_STRBUF_SIZE : s->s_used + used;
    return used;
}

/*
 * Do fixup of the format: deal with all standard printf % options,
 * plus any extras. When we see a pc_flag in one of the structures in
 * the list below, then we call the associated pc_func and interpolate
 * the string it returns.
 */

typedef struct _percent *percent;
struct _percent {
    percent pc_link;
    char pc_flag; /* If we see this flag ... */
    const char *(*pc_func)(); /* ... then this function gives us the string */
};

```

Sep 17, 15 17:34

strbuf.c

Page 3/4

```

private percent percent_list = NULL;

private percent find_in_list(char c) {
    percent p = percent_list;

    for( ; p; p=p->pc_link) {
        if(p->pc_flag == c)
            return p;
    }
    return NULL;
}

private void do_extra_percents(char *buf, int size, const char *fmt) {
    /* Copy the format into the buffer, checking each % modifier against the list */
    for( ; size > 1 && *fmt; size-- ) {
        if( (*buf++ = *fmt++) != '%' )
            continue;

        size--; /* Count the previous % that was copied */
        percent p = find_in_list(*buf++ = *fmt++); /* Look up the next character */

        if(p == NULL)
            continue; /* Not an extension, keep copying */

        int used = snprintf(buf-2, size, "%s", (*p->pc_func)()); /* Interpolate at most 'size' chars */
        if(used >= size) { /* Output was truncated */
            buf += size-1;
            break;
        }
        buf += used-2; /* We copied 'used' characters, overwriting 2 */
        size -= used-3; /* The for loop will decrement one more... */
    }
    *buf = '\0';
}

/* Start printing into the buffer at position pos */
public int strbuf_printf_pos(strbuf s, int pos, const char *fmt, ...) {
    va_list ap;
    int used;
    char fmt_buf[MAX_STRBUF_SIZE];

    if( strbuf_setpos(s, pos) < 0 )
        return -1;
    do_extra_percents(&fmt_buf[0], MAX_STRBUF_SIZE, fmt);
    va_start(ap, fmt);
    used = strbuf_vprintf(s, pos, fmt_buf, ap);
    va_end(ap);
    return used;
}

/* Start printing into the buffer at position 0 */
public int strbuf_printf(strbuf s, const char *fmt, ...) {
    va_list ap;
    int used;
    char fmt_buf[MAX_STRBUF_SIZE];

    if( strbuf_setpos(s, 0) < 0 )
        return -1;
    do_extra_percents(&fmt_buf[0], MAX_STRBUF_SIZE, fmt);
    va_start(ap, fmt);
    used = strbuf_vprintf(s, 0, fmt_buf, ap);
    va_end(ap);
    return used;
}

/* Start printing into the buffer at the current position */
public int strbuf_appendf(strbuf s, const char *fmt, ...) {
    va_list ap;
    int used;
    char fmt_buf[MAX_STRBUF_SIZE];

```


Sep 17, 15 17:34

strbuf.c

Page 4/4

```

do_extra_percents(&fmt_buf[0], MAX_STRBUF_SIZE, fmt);
va_start(ap, fmt);
used = strbuf_vprintf(s, -1, fmt_buf, ap);
va_end(ap);
return used;
}

/*
 * Register new percent interpreters.
 */

public int register_error_percent_handler(char c, const char *(*fn)()) {
    percent p = calloc(1, sizeof(struct _percent));

    if(p == NULL) {
        return -1;
    }

    p->pc_link = percent_list;
    p->pc_flag = c;
    p->pc_func = fn;
    percent_list = p;
    return 0;
}

/*
 * Revert a strbuf -- remove extra NUL characters inserted by tokenising
 */

public void strbuf_revert(strbuf s) {
    char *p = &s->s_buffer[0];
    int n;

    for(n=0; n<s->s_used; n++,p++)
        if( !*p ) *p = ' ';
    n = (n == MAX_STRBUF_SIZE)? n-1 : n;
    s->s_buffer[n] = '\0';
}

/*
 * Debug a strbuf
 */

public void debug_strbuf(FILE *fp, strbuf s) {
    char *str = strbuf_string(s);
    char buf[MAX_STRBUF_SIZE+64];
    char *b;
    int used;
    int n;

    used = snprintf(&buf[0], 64, "s=%p,n=%p,q=%p:data[0..%d]='",
        s, s->s_Q.q_next, s->s_Q.q_prev, s->s_used);
    b = &buf[used];
    n = MAX_STRBUF_SIZE+64-used-1;
    if( n > s->s_used )
        n = s->s_used;
    while(n-- > 0) if( (*b++ = *str++) == '\0' ) b[-1] = ' ';
    *b++ = '\0';
    *b = '\0';
    fwrite(&buf[0], 1, b-&buf[0], fp);
}

```

Sep 17, 15 17:35

strbuf.h

Page 1/1

```
#
#ifndef _STRBUF_H
#define _STRBUF_H

#include "general.h"

/*
 * Error buffer structure.
 */

typedef struct _strbuf *strbuf; /* Opaque object */

export strbuf alloc_strbuf();
export void   release_strbuf(strbuf);
export char  *strbuf_string(strbuf);
export int   strbuf_space(strbuf);
export int   strbuf_used(strbuf);
export int   strbuf_setpos(strbuf,int);

#include <stdio.h>
#include <stdarg.h>

export int  strbuf_printf(strbuf, const char *, ...);
export int  strbuf_appendf(strbuf, const char *, ...);
export int  strbuf_printf_pos(strbuf, int, const char *, ...);
export int  register_error_percent_handler(char, const char *(*)());
export void strbuf_revert(strbuf);
export void debug_strbuf(FILE *, strbuf);

#define strbuf_clear(s) ((void) strbuf_setpos(s, 0))

#define strbuf_next(s)  ((strbuf)queue_next((queue *) (s)))
#define strbuf_prev(s)  ((strbuf)queue_prev((queue *) (s)))

#define strbuf2qp(s)    ((queue *) (s))
#define qp2strbuf(q)    ((strbuf) (q))

#endif /* _STRBUF_H */
```

Sep 15, 15 19:58

tidy.c

Page 1/2

```

#
#include "general.h"

/*
 * Low-priority thread that unlocks pages after they've been filled.
 */

#include <stdio.h>
#include <stdlib.h>
#include "assert.h"
#include <errno.h>
#include <sys/mman.h>
#include <pthread.h>
#include <zmq.h>

#include "util.h"
#include "strbuf.h"
#include "chunk.h"
#include "param.h"
#include "tidy.h"
#include "snapshot.h"

import void *snapshot_zmq_ctx;

private void *tidy;
private void *log;

/*
 * Establish tidy comms: this routine gets called first of all threads, so it
 * creates the context.
 */

private char *create_tidy_comms(void **s) {
    if( !snapshot_zmq_ctx )
        snapshot_zmq_ctx = zmq_ctx_new();
    if( !snapshot_zmq_ctx ) {
        return "failed to create ZMQ context";
    }

    /* Create and initialise the sockets: */

    /* MAIN thread's log socket */
    *s = zh_bind_new_socket(snapshot_zmq_ctx, ZMQ_PULL, LOG_SOCKET);
    if( *s == NULL ) {
        return "unable to create MAIN thread log socket";
    }

    /* TIDY's socket for work messages */
    tidy = zh_bind_new_socket(snapshot_zmq_ctx, ZMQ_PAIR, TIDY_SOCKET);
    if(tidy == NULL)
        return "unable to create TIDY thread listener";

    /* TIDY's socket for log messages */
    log = zh_connect_new_socket(snapshot_zmq_ctx, ZMQ_PUSH, LOG_SOCKET);
    if(log == NULL)
        return "unable to create TIDY thread log socket";

    return NULL;
}

/* Close the TIDY thread's comms channels */

private void close_tidy_comms() {
    zmq_close(tidy);
    zmq_close(log);
}

/*
 * Unmap data blocks after writing. Runs as a thread which continues
 * until a zero-length message is received signalling the end of the
 * unmap requests. The argument passed is the address for the MAIN thread's
 * log receiver socket, which is created here along with the context.
 */

```

Sep 15, 15 19:58

tidy.c

Page 2/2

```
public void *tidy_main(void *arg) {
    char *err;
    int ret;
    frame *f;

    err = create_tidy_comms((void **)arg);
    if(err) {
        die_die_die_now++;
        return (void *) err;
    }

    zh_put_multi(log, 1, "TIDY thread initialised");

    while( ret = zh_get_msg(tidy, 0, sizeof(frame *), &f) && !die_die_die_now ) {
        assertv(ret==sizeof(frame *), "TIDY read message error, ret=%d\n", ret);
        release_frame(f);
    }

    zh_put_multi(log, 1, "TIDY thread terminates by return");

    /* Clean up our ZeroMQ sockets */
    close_tidy_comms();
    return (void *) "normal exit";
}
```

Sep 12, 15 16:52

tidy.h

Page 1/1

```
#
#ifndef _TIDY_H
#define _TIDY_H

#include "general.h"

export void *tidy_main(void *);

#define TIDY_SOCKET "inproc://snapshot-TIDY"

#endif /* _TIDY_H */
```

Sep 10, 15 14:15

trig.c

Page 1/8

```

#
/*
 * Program to generate triggered snapshots manually. This communicates with the
 * snapshotter via ZMQ.
 *
 * Arguments:
 * --verbose/-v      Increase reporting level
 * --quiet/-q        Decrease reporting level
 * --snapshot/-s     The snapshotter socket address
 * --pre            Pre-trigger interval
 * --post           Post-trigger interval
 * --trigger         Timepoint of trigger
 * --wait-for-it/-w  Wait for a key-press to generate trigger
 * --repeat/-r       Generate multiple triggers instead of just one
 * --auto/-a         Generate the snapshot name automatically
 * --help/-h         Print usage message
 * --version         Print program version
 */

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <stdint.h>
#include "assert.h"

#include <zmq.h>
#include <argtable2.h>
#include <regex.h>
#include <errno.h>
#include <string.h>
#include <getopt.h>
#include <time.h>

#include "argtab.h"
#include "util.h"
#include "param.h"
// #include "snapshot.h"

/*
 * Program source version
 */

#define PROGRAM_VERSION "1.0"
#define VERSION_VERBOSE_BANNER "MCLURS ADC toolset...\n"

/*
 * Auto-name format options
 */

#define AUTO_NAME_FORMAT_DEFAULT "iso"
#define AUTO_NAME_FORMAT_REX     "hex|iso|utc|tai|seq"

/*
 * Code for automatic snapshot name generation -- interprets format options
 */

typedef enum {
    SNAPNAME = 0,
    HEXADECIMAL,
    TAI64N,
    ISOUTC,
    ISODATE,
    SEQUENTIAL,
    SPECIAL,
} name_mode;

static name_mode determine_auto_mode(const char *auto_name) {
    if( !auto_name )
        return SNAPNAME;
    if( !strcmp("hex", auto_name) )
        return HEXADECIMAL;
    if( !strcmp("tai", auto_name) )
        return TAI64N;
    if( !strcmp("utc", auto_name) )

```

Sep 10, 15 14:15

trig.c

Page 2/8

```

    return ISODATE;
    if( !strcmp("iso", auto_name) )
        return ISODATE;
    if( !strcmp("seq", auto_name) )
        return SEQUENTIAL;
    return SPECIAL;
}

/*
 * Global parameters for the snapshot program
 */

extern const char *snapshot_addr;
extern const char *auto_name;
extern uint32_t window_pre;
extern uint32_t window_pst;

param_t globals[] = {
    { "snapshot", "ipc://snapshot-CMD", &snapshot_addr, PARAM_TYPE(string), PARAM_SRC_ENV|PARAM_SRC_ARG,
      "address of snapshot command socket"
    },
    { "pre", "1000", &window_pre, PARAM_TYPE(int32), PARAM_SRC_ARG,
      "pre-trigger duration [ms]"
    },
    { "pst", "500", &window_pst, PARAM_TYPE(int32), PARAM_SRC_ARG,
      "post-trigger duration [ms]"
    },
    { "auto", AUTO_NAME_FORMAT_DEFAULT, &auto_name, PARAM_TYPE(string), PARAM_SRC_ARG,
      "format of auto-generated name"
    },
};

const int n_global_params = (sizeof(globals)/sizeof(param_t));

/*
 * Debugging print out control
 */

int verbose = 0;
char *program = NULL;

/* Command line syntax options */

struct arg_lit *h1, *vn1, *v1, *q1;
struct arg_end *e1;

BEGIN_CMD_SYNTAX(help) {
    v1 = arg_litn("v", "verbose", 0, 3, "Increase verbosity"),
    q1 = arg_lit0("q", "quiet", "Decrease verbosity"),
    h1 = arg_lit0("h", "help", "Print usage help message"),
    vn1 = arg_lit0(NULL, "version", "Print program version string"),
    e1 = arg_end(20)
} APPLY_CMD_DEFAULTS(help) {
    /* No defaults to apply here */
} END_CMD_SYNTAX(help)

struct arg_lit *v2, *q2, *w2;
struct arg_end *e2;
struct arg_str *u2;
struct arg_int *pb2, *pe2;
struct arg_str *n2;

BEGIN_CMD_SYNTAX(single) {
    v2 = arg_litn("v", "verbose", 0, 3, "Increase verbosity"),
    q2 = arg_lit0("q", "quiet", "Decrease verbosity"),
    u2 = arg_str0("s", "snapshot", "<url>", "URL of snapshotter command socket"),
    pb2 = arg_int0(NULL, "pre", "<int>", "Pre-trigger interval [ms]"),
    pe2 = arg_int0(NULL, "pst,post", "<int>", "Post-trigger interval [ms]"),
    w2 = arg_lit0("w", "wait-for-it", "Wait for keypress to trigger"),
    n2 = arg_str1(NULL, NULL, "<snapshot name>", "Name of the snapshot file"),
    e2 = arg_end(20)
} APPLY_CMD_DEFAULTS(single) {
    INCLUDE_PARAM_DEFAULTS(globals, n_global_params);
} END_CMD_SYNTAX(single);

```

Sep 10, 15 14:15

trig.c

Page 3/8

```

struct arg_lit *v3, *q3, *w3;
struct arg_end *e3;
struct arg_str *u3;
struct arg_rex *a3;
struct arg_int *pb3, *pe3;

BEGIN_CMD_SYNTAX(autoname) {
    v3 = arg_litn("v", "verbose", 0, 3, "Increase verbosity" ),
    q3 = arg_lit0("q", "quiet", "Decrease verbosity" ),
    u3 = arg_str0("s", "snapshot", "<url>", "URL of snapshotter command socket" ),
    pb3 = arg_int0(NULL, "pre", "<int>", "Pre-trigger interval [ms]" ),
    pe3 = arg_int0(NULL, "pst.post", "<int>", "Post-trigger interval [ms]" ),
    w3 = arg_lit0("w", "wait-for-it", "Wait for keypress to trigger" ),
    a3 = arg_rex1("a", "auto", AUTO_NAME_FORMAT_REX, "<format>", REG_EXTENDED, "Automatic snapshot name" ),
    e3 = arg_end(20)
} APPLY_CMD_DEFAULTS(autoname) {
    a3->hdr.flag |= ARG_HASOPTVALUE;
    INCLUDE_PARAM_DEFAULTS(globals, n_global_params);
} END_CMD_SYNTAX(autoname);

struct arg_lit *v4, *q4, *r4, *w4;
struct arg_end *e4;
struct arg_str *u4;
struct arg_rex *a4;
struct arg_int *pb4, *pe4;

BEGIN_CMD_SYNTAX(repeat) {
    v4 = arg_litn("v", "verbose", 0, 3, "Increase verbosity" ),
    q4 = arg_lit0("q", "quiet", "Decrease verbosity" ),
    u4 = arg_str0("s", "snapshot", "<url>", "URL of snapshotter command socket" ),
    pb4 = arg_int0(NULL, "pre", "<int>", "Pre-trigger interval [ms]" ),
    pe4 = arg_int0(NULL, "pst.post", "<int>", "Post-trigger interval [ms]" ),
    a4 = arg_rex0("a", "auto", AUTO_NAME_FORMAT_REX, "<format>", REG_EXTENDED, "Automatic snapshot name" ),
        arg_rem(NULL, "<format> is " AUTO_NAME_FORMAT_REX),
    w4 = arg_lit0("w", "wait-for-it", "Wait for keypress to trigger" ),
    r4 = arg_lit1("r", "repeat", "Loop, generating multiple triggers (implies -wa)" ),
    e4 = arg_end(20)
} APPLY_CMD_DEFAULTS(repeat) {
    a4->hdr.flag |= ARG_HASOPTVALUE;
    INCLUDE_PARAM_DEFAULTS(globals, n_global_params);
} END_CMD_SYNTAX(repeat);

/* Standard help routines: display the version banner */
void print_version(FILE *fp, int verbosity) {
    fprintf(fp, "%sVn.%s\n", program, PROGRAM_VERSION);
    if(verbosity > 0) { /* Verbose requested... */
        fprintf(fp, VERSION_VERBOSE_BANNER);
    }
}

/* Standard help routines: display the usage summary for a syntax */
void print_usage(FILE *fp, void **argtable, int verbosity, char *program) {
    if( !verbosity ) {
        fprintf(fp, "Usage: %s ", program);
        arg_print_syntax(fp, argtable, "\n");
        return;
    }
    if( verbosity ) {
        char *suffix = verbosity>1? "\n\n" : "\n";
        fprintf(fp, "Usage: %s ", program);
        arg_print_syntaxv(fp, argtable, suffix);
        if( verbosity > 1 )
            arg_print_glossary(fp, argtable, "%-25s %s\n");
    }
}

/*
 * Globals
 */

void *zmq_main_ctx; /* ZMQ context for messaging */

const char *auto_name; /* Auto-generate snapshot path value */
name_mode auto_mode; /* The basis for snapshot name generation */
const char *snap_name; /* The base name if not auto */

```


Sep 10, 15 14:15

trig.c

Page 4/8

```

const char *snapshot_addr;    /* URL of the snapshotter program */
int wait_for_it;              /* Wait for keypress before making message */
int repeat;                   /* Don't just do one, do many triggers */
uint32_t window_pre;          /* Window pre-trigger interval [ms] */
uint32_t window_pst;          /* Window post-trigger interval [ms] */

/*
 * Process a (possibly multipart) log message.
 * Collect the various pieces and write to stderr
 * Use a 1024 byte logging buffer
 */

#define LOGBUF_SIZE 1024

void print_message(void *socket) {
    char log_buffer[LOGBUF_SIZE];
    int used;

    used = zh_collect_multi(socket, &log_buffer[0], LOGBUF_SIZE-1, "");
    if( log_buffer[used-1] != '\n' ) {
        log_buffer[used] = '\n';
        fwrite(log_buffer, used+1, 1, stdout);
    }
    else {
        fwrite(log_buffer, used, 1, stdout);
    }
    fflush(stdout);
}

/*
 * Wait for a keypress to generate a trigger time.
 */
uint64_t wait_for_keypress(uint64_t *now_as_ns) {
    struct timespec now;

    if(now_as_ns == NULL)
        return -1;

    fputc('>', stdout);
    switch( fgetc(stdin) ) {
    case EOF:
    case 'q':
        return -1;

    case 's':
        repeat = 0;
        break;

    default:
        break;
    }

    /* Discover the current time, as trigger point */
    clock_gettime(CLOCK_MONOTONIC, &now);
    *now_as_ns = now.tv_sec;
    *now_as_ns = *now_as_ns * 1000000000 + now.tv_nsec;
    return 0;
}

/*
 * Construct the name of a snapshot file. Use the supplied name unless in auto mode.
 */
char *make_path_value(char buf[], int size, const char *snapname, uint64_t trigger, name_mode mode) {
    static int counter = 0;
    time_t trig;
    uint64_t secs;
    int ns;
    int used;
    struct tm *t;

    switch(mode) {

```

Sep 10, 15 14:15

trig.c

Page 5/8

```

case SNAPNAME:                /* Use the supplied snapshot name */
    snprintf(&buf[0], size, "%s", snapname);
    break;

case TAI64N:                   /* Use a TAI64N format timestamp */
    assertv(size >= 25, "Buffer too small (%d) for TAI path\n", size);
    secs = trigger / 1000000000;
    ns = trigger - secs * 1000000000;
    snprintf(&buf[0], size, "@%016llx%08lx", secs|0x4000000000000000, ns);
    break;

case ISODATE:
case ISOUTC:                   /* Use an ISO standard date with fractional seconds */
    assertv(size >= 26, "Buffer too small (%d) for ISO path\n", size);
    trig = trigger / 1000000000;
    ns = trigger - trig * 1000000000;
    t = (mode==ISOUTC? gmtime(&trig) : localtime(&trig));
    used = strftime(&buf[0], size, "%FT%T", t); /* 2015-07-14T16:55:32.nnnnnn */
    snprintf(&buf[used], size-used, "%06d", ns/1000);
    break;

case HEXADECIMAL:              /* Use a hexadecimal print of the trigger time */
    assertv(size >= 16, "Buffer too small (%d) for HEX path\n", size);
    snprintf(&buf[0], size, "%016llx", trigger);
    break;

case SEQUENTIAL:               /* Generate a sequentially incrementing snapshot name */
    assertv(size >= 10, "Buffer too small (%d) for SEQ path\n", size);
    snprintf(&buf[0], size, "snap%06d", counter++);
    break;

case SPECIAL:                  /* User-supplied format, not yet implemented */
default:
    snprintf(&buf[0], size, "%s%d", "unimplemented", counter++);
    break;
}
return &buf[0];
}

/*
 * Main entry point
 */

#define PATHBUF_SIZE 128

int main(int argc, char *argv[], char *envp[]) {
    char    buf[LOGBUF_SIZE];
    void    *snapshot;
    param_t *p;
    char    *v;
    int     ret, n;
    int     used, left;
    uint64_t trigger;
    struct timespec now;
    uint64_t now_as_ns;

    /* Discover the current time, as trigger point */
    clock_gettime(CLOCK_MONOTONIC, &now);
    now_as_ns = now.tv_sec;
    now_as_ns = now_as_ns * 1000000000 + now.tv_nsec;

    program = argv[0];

    /* Set up the standard parameters */
    /* 1. Process parameters: internal default, then environment. */
    set_param_from_env(envp, globals, n_global_params);

    /* 2. Process parameters: push values out to program globals */
    ret = assign_all_params(globals, n_global_params);
    assertv(ret == 0, "Push parameters failed on param %d out of %d\n", -ret, n_global_params);

    // fprintf(stderr, "Before command line processing, after environment\n");
    // debug_params(stderr, globals, n_global_params);

    /* 3. Create and parse the command lines -- installs defaults from parameter table */

```

Sep 10, 15 14:15

trig.c

Page 6/8

```

void **cmd_help      = arg_make_help();
void **cmd_single    = arg_make_single();
void **cmd_autoname  = arg_make_autoname();
void **cmd_repeat    = arg_make_repeat();

/* Try first syntax */
int err_help = arg_parse(argc, argv, cmd_help);
if( !err_help ) {
    int verbose = v1->count - q1->count;
    if(vn1->count)
        print_version(stdout, verbose);
    if(hl->count || !vn1->count) {
        print_usage(stdout, cmd_help, verbose>0, program);
        print_usage(stdout, cmd_single, verbose>0, program);
        print_usage(stdout, cmd_autoname, verbose>0, program);
        print_usage(stdout, cmd_repeat, verbose, program);
    }
    exit(0);
}

struct arg_end *found = NULL;
void **table = NULL;
int errs = 0, min_errs = 100;

/* Try remaining syntaxes */
errs = arg_parse(argc, argv, cmd_single);
if( !errs || errs < min_errs ) { /* Choose single trigger manual-named mode */
    found = e2;
    table = cmd_single;
    verbose = v2->count - q2->count;
    min_errs = errs;
    if( !errs ) {
        auto_name = NULL;
        repeat = 0;
        wait_for_it = w2->count;
        snap_name = n2->sval[0];
    }
}

if( errs ) {
    errs = arg_parse(argc, argv, cmd_autoname);
    if( !errs || errs < min_errs ) { /* Choose single trigger auto-named mode */
        found = e3;
        table = cmd_autoname;
        verbose = v3->count - q3->count;
        min_errs = errs;
        if( !errs ) {
            repeat = 0;
            wait_for_it = w2->count;
            snap_name = NULL;
        }
    }
}

if( errs ) {
    errs = arg_parse(argc, argv, cmd_repeat);
    if( !errs || errs < min_errs ) { /* Choose multi-trigger mode */
        found = e4;
        table = cmd_repeat;
        verbose = v4->count - q4->count;
        min_errs = errs;
        if( !errs ) {
            repeat = 1;
            wait_for_it = 1;
            if( !a4->count || !w4->count ) {
                if(verbose >= 0)
                    fprintf(stderr, "%s: Warning -- repeat (-r) implies -a and -w, using --auto=%s\n", program, auto_name);
            }
            snap_name = NULL;
        }
    }
}

/* Now found indicates the command line with minimum errors in parse */

```

Sep 10, 15 14:15

trig.c

Page 7/8

```

if( min_errs ) {
    arg_print_errors(stderr, found, program);
    print_usage(stderr, cmd_help, verbose>0, program);
    print_usage(stderr, cmd_single, verbose>0, program);
    print_usage(stderr, cmd_autoname, verbose>0, program);
    print_usage(stderr, cmd_repeat, verbose, program);
    exit(1);
}

// fprintf(stderr, "After commandline choice, before reverse push\n");
// debug_params(stderr, globals, n_global_params);

/* 4. Process parameters: copy argument values back through the parameter table */
ret = arg_results_to_params(table, globals, n_global_params);

// fprintf(stderr, "After reverse push\n");
// debug_params(stderr, globals, n_global_params);

/* Check the auto argument and compute the path generation mode */
auto_mode = determine_auto_mode(auto_name);

/* 5. All syntax tables are finished with now: clean up the mess :-)) */
arg_free(cmd_help);
arg_free(cmd_single);
arg_free(cmd_autoname);
arg_free(cmd_repeat);

if(verbose > 2)
    debug_params(stderr, globals, n_global_params);

/* Create the ZMQ contexts */
zmq_main_ctx = zmq_ctx_new();
if( !zmq_main_ctx ) {
    fprintf(stderr, "%s: Error -- ZeroMQ context creation failed: %s\n", program, strerror(errno));
    exit(2);
}

/* Create the socket to talk to the snapshot program */
snapshot = zh_connect_new_socket(zmq_main_ctx, ZMQ_REQ, snapshot_addr);
if( snapshot == NULL ) {
    fprintf(stderr, "%s: Error -- unable to create socket to snapshot at %s: %s\n",
        program, snapshot_addr, strerror(errno));
    zmq_ctx_term(zmq_main_ctx);
    exit(2);
}

/* Look at the parameters to construct the snap command */
if(window_pre + window_pst > 8000) {
    fprintf(stderr, "%s: Error -- maximum allowed capture window is 8000 [ms]\n");
    exit(3);
}

trigger = now_as_ns;
do {
    uint64_t time_start, time_stop;
    char path_buf[PATHBUF_SIZE];
    const char *path;
    int ret = 0;

    if( wait_for_it )
        ret = wait_for_keypress(&trigger);

    if( ret < 0 ) {
        break;
    }

    path = make_path_value(&path_buf[0], PATHBUF_SIZE-1, snap_name, trigger, auto_mode);

    // fprintf(stderr, "Window parameters pre %d post %d\n", window_pre, window_pst);

    time_start = trigger - 1000000 * (uint64_t) window_pre;
    time_stop = trigger + 1000000 * (uint64_t) window_pst;

    // fprintf(stderr, "Window parameters start %lld stop %lld\n", time_start, time_stop);

```

Sep 10, 15 14:15

trig.c

Page 8/8

```
/* Send the message, wait for the reply */
left = LOGBUF_SIZE-1;
used = snprintf(&buf[0], left, "snap begin=%lld,end=%lld,path=%s", time_start, time_stop, path);
buf[used] = '\0';
if(verbose > 1)
    fprintf(stderr, "Sending:%s\n", &buf[0]);
ret = zh_put_msg(snapshot, 0, used, buf);
if( ret < 0 ) {
    fprintf(stderr, "%s: Error -- sending message to %s failed\n", program, snapshot_addr);
    break;
}

/* Wait for reply */
if(verbose > 1)
    fprintf(stderr, "Awaiting reply...\n");
if(verbose > 0)
    print_message(snapshot);
} while(repeat);

/* Clean up ZeroMQ sockets and context */
zmq_close(snapshot);
zmq_ctx_term(zmq_main_ctx);
exit(0);
}
```

Sep 12, 15 17:15

util.c

Page 1/3

```

#
#include "general.h"

#include <stdio.h>
#include <stdlib.h>
#include <stdarg.h>
#include <string.h>
#include <errno.h>
#include "assert.h"

#include <zmq.h>

#include "util.h"

/*
 * Create, open and bind a ZMQ socket.
 */
public void *zh_bind_new_socket(void *ctx, int type, const char *url) {
    void *skt;

    skt = zmq_socket(ctx, type);
    if(skt != NULL) {
        int ret = zmq_bind(skt, url);
        if(ret < 0) {
            int safe_errno = errno;
            (void) zmq_close(skt);
            errno = safe_errno;
            skt = NULL;
        }
    }
    return skt;
}

/*
 * Create, open and connect a ZMQ socket.
 */
public void *zh_connect_new_socket(void *ctx, int type, const char *url) {
    void *skt;

    skt = zmq_socket(ctx, type);
    if(skt != NULL) {
        int ret = zmq_connect(skt, url);
        if(ret < 0) {
            int safe_errno = errno;
            (void) zmq_close(skt);
            errno = safe_errno;
            skt = NULL;
        }
    }
    return skt;
}

/*
 * Retrieve a ZMG message from a socket. Put it in the buffer buf and
 * transfer at most size bytes. If size is zero, we care only about
 * the arrival of the message, not its content.
 */
public int zh_get_msg(void *socket, int flags, size_t size, void *buf) {
    zmq_msg_t msg;
    int ret;
    size_t msg_size;

    ret = zmq_msg_init(&msg);
    assertv(ret == 0, "Message init failed\n");
    ret = zmq_msg_recv(&msg, socket, flags);
    if( ret < 0 )
        return ret;
    if( !size )
        return 0;
    msg_size = zmq_msg_size(&msg);
    if( !msg_size )

```

Sep 12, 15 17:15

util.c

Page 2/3

```

    return 0;
    if( msg_size < size )
        size = msg_size;
    assertv(buf != NULL, "Called with null buf argument\n");
    bcopy(zmq_msg_data(&msg), buf, size);
    ret = zmq_msg_close(&msg);
    assertv(ret == 0, "Message close failed\n");
    return size;
}

/*
 * Returns true if there is more of this message, otherwise false
 */

public int zh_any_more(void *socket) {
    int ret, more;
    size_t sz;

    sz = sizeof(more);
    ret = zmq_getsockopt(socket, ZMQ_RCVMORE, &more, &sz);
    assertv(ret == 0, "Attempt to get 'more' flag failed\n");
    return more != 0;
}

/*
 * Get a multipart message in a single buffer. Concatenate the
 * pieces, with 'spc' in between. End with \0. Return the size.
 */

public int zh_collect_multi(void *socket, char *buf, int bufsz, char *spc) {
    int used = 0,
        left = bufsz-1,
        nspc = strlen(spc);

    do {
        int ret, sz;

        sz = zh_get_msg(socket, 0, left-nspc, &buf[used]);
        assertv(sz >= 0, "Get message error\n");
        used += sz;
        left -= sz;
        if( !zh_any_more(socket) )
            break;
        bcopy(spc, &buf[used], nspc);
        used += nspc;
        left -= nspc;
    } while( left >= 0 );
    buf[used] = '\0';
    return used;
}

/*
 * Send a ZMQ message via a socket. If size is zero, send an empty
 * frame, and buf can be NULL. If ZMQ_SNDMORE is given as flag, this
 * is part of a multipart message.
 */

public int zh_put_msg(void *socket, int flags, size_t size, void *buf) {
    zmq_msg_t msg;
    int ret;

    assertv(size >= 0, "Put message with -ve size %d\n", size);
    ret = zmq_msg_init_size(&msg, size);
    assertv(ret == 0, "Message init failed\n");
    if( size ) {
        assertv(buf != NULL, "Non-zero size and NULL buf\n");
        bcopy(buf, zmq_msg_data(&msg), size);
    }
    return zmq_msg_send(&msg, socket, flags);
}

/*
 * Send an n-frame message via a socket given an argument list of strings.
 */

```

Sep 12, 15 17:15

util.c

Page 3/3

```
public int zh_put_multi(void *socket, int n, ...) {
    va_list ap;
    int ret;

    va_start(ap,n);
    while( n-- > 0 ) {
        char *next = va_arg(ap, char *);
        int sz = strlen(next);
        ret = zh_put_msg(socket, (n==0? 0 : ZMQ_SNDMORE), sz, next);
        if( ret < 0 )
            return ret;
    }
    va_end(ap);
    return 0;
}
```


Sep 12, 15 16:52

util.h

Page 1/1

```
#
#ifndef _UTIL_H
#define _UTIL_H

#include "general.h"

#include <stdarg.h>
#include <unistd.h>
#include "assert.h"

#define true 1
#define false 0

#define WAIT_FOR_CONDITION(cond,limit)          \
do { double l = (limit); int n = 0, max = 100*l; \
while( n<max && !(cond) ) usleep(10000), n++; \
assertv((cond), "Waited too long (%g [s]) for condition\n", l); \
} while(0)

/* Messaging utilities */

#include <zmq.h>

export int zh_get_msg(void *, int, size_t, void *);
export int zh_any_more(void *);

export int zh_put_msg(void *, int, size_t, void *);
export int zh_put_multi(void *, int, ...);

export void *zh_bind_new_socket(void *, int, const char *);
export void *zh_connect_new_socket(void *, int, const char *);

#endif /* _UTIL_H */
```

Sep 22, 15 16:05

writer.c

Page 1/21

```

#
#include "general.h"
#define _GNU_SOURCE

#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/mman.h>
#include <sys/capability.h>
#include <time.h>
#include <pwd.h>
#include <grp.h>

#include "assert.h"

#include <zmq.h>
#include <pthread.h>

#include <comedi.h>
#include <fcntl.h>
#include <unistd.h>
#include <errno.h>
#include <string.h>

#include "util.h"
#include "param.h"
#include "queue.h"
#include "mman.h"
#include "strbuf.h"
#include "chunk.h"
#include "adc.h"
#include "snapshot.h"
#include "reader.h"
#include "writer.h"

/* We import the READER's ADC object for its time conversion and activity check routines */
import adc reader_adc;

/*
 * -----
 *
 * TYPES INTERNAL TO THE WRITER THREAD
 *
 * -- snapshot descriptor
 * -- snapshot file descriptor
 * -- forward function declarations
 * -- local queue headers
 */

/* Snapshot Descriptor Structure */
typedef struct {
    queue    s_xQ[2];          /* Private snapshot descriptor structure used by writer */
#define s_Q    s_xQ[0]        /* Queue headers -- must be first member */
#define s_fileQhdr    s_xQ[1] /* Active snapshot queue header */
    uint16_t s_name;           /* Header for the queue of file descriptor structures */
    int      s_dirfd;          /* 'Name' for snapshot */
    uint64_t s_first;           /* Dirfd of the samples directory */
    uint64_t s_last;            /* First sample to collect for the next repetition */
    uint32_t s_samples;         /* Collect up to but not including this sample in the next repetition */
    int      s_bytes;           /* Number of samples to save */
    uint32_t s_count;           /* Total size of one sample file */
    int      s_pending;         /* Repetition count for this snapshot */
    int      s_done;            /* Count of pending repetitions */
    int      s_status;          /* Count of completed repetitions */
    const char *s_path;         /* Status of this snapshot */
    strbuf    s_error;          /* Directory path for this snapshot */
} snap_t;

/* Forward declarations of snapshot descriptor routines */

```

Sep 22, 15 16:05

writer.c

Page 2/21

```

private uint16_t snapshot_name(snap_t *);

#define qp2snap(qp) ((snap_t *)&(qp)[0])
#define snap2qp(s) (&((s)->s_xQ[0]))

#define fq2snap(fq) ((snap_t *)&(fq)[-1])
#define snap2fq(s) (&((s)->s_xQ[1]))

#define qp2sname(p) snapshot_name(qp2snap(p))

/* Snapshot File Descriptor Structure */

typedef struct _sfile {
    queue    f_Q;                /* Queue header for file descriptor structures */
    snap_t   *f_parent;          /* The snap_t structure that generated this file capture */
    int       f_fd;              /* System file descriptor -- only needed while pages left to map */
    int       f_indexnr;         /* Index number of this file in the full set for the snapshot */
    int       f_nchunks;         /* Number of chunks allocated for this file transfer */
    int       f_written;         /* Number of chunks actually written so far */
    int       f_pending;         /* Number of chunks controlled by the READER thread */
    chunk_t   *f_chunkQ;         /* Pointer to this file's writer chunk queue */
    int       f_status;          /* Status flags for this file */
    strbuf    f_error;           /* The strbuf to write error text into */
    uint16_t  f_name;            /* Unique number for debugging */
    char       f_file[FILE_NAME_SIZE]; /* Name of this file: the hexadecimal first sample number .s16 */
}
    snapfile_t;

/* Forward declarations of snapshot file descriptor routines needed by snapshot */

private snapfile_t *alloc_snapfile();
private int setup_snapfile(snapfile_t *, snap_t *);
private void abort_snapfile(snapfile_t *);
private void debug_snapfile(snapfile_t *);
import uint16_t snapfile_name(snapfile_t *);

#define qp2file(p) ((snapfile_t *) (p))
#define file2qp(f) (&(f)->f_Q)

#define qp2fname(p) snapfile_name(qp2file(p))

/* Local queue headers etc. used by the WRITER thread */

private QUEUE_HEADER(snapQ); /* The list of active snapshots */
private QUEUE_HEADER(WriterChunkQ); /* The list of chunks awaiting mapping, in order of first sample */

/*
 * -----
 *
 * INITIALISATION ROUTINES FOR WRITER THREAD:
 *
 * - Establish the communication endpoints needed
 * - Set up the required effective capabilities
 * - Set up RT priority scheduling (if requested)
 *
 * -----
 */

/*
 * Writer parameter structure.
 */

public wparams writer_parameters;

/* The values below are computed from writer_parameters by the verify() function */

private int wp_nframes; /* Number of transfer frames prepared */
private int wp_chunksamples; /* Number of samples in a chunk */
private int wp_snap_dirfd; /* Snapdir path fd */
private int wp_snap_curfd; /* Path fd of the 'working' directory */
private int wp_totxfrsamples; /* Total scheduled transfer samples remaining */
private int wp_nfiles; /* Number of files in progress */

/* Read-only access to chunk size, needed by READER */

```

Sep 22, 15 16:05

writer.c

Page 3/21

```

public int writer_chunksize_samples() {
    return wp_chunksamples;
}

/*
 * Reader thread comms initialisation (failure is fatal).
 * Called after the process-wide ZMQ context is created (elsewhere).
 */

private void *log;
private void *reader;
private void *command;

private void create_writer_comms() {
    import void *snapshot_zmq_ctx;
    /* Create necessary sockets */
    command = zh_bind_new_socket(snapshot_zmq_ctx, ZMQ_REP, WRITER_CMD_ADDR); /* Receive commands */
    assertv(command != NULL, "Failed to instantiate reader command socket\n");
    log = zh_connect_new_socket(snapshot_zmq_ctx, ZMQ_PUSH, LOG_SOCKET); /* Socket for log messages */
    assertv(log != NULL, "Failed to instantiate reader log socket\n");
    reader = zh_connect_new_socket(snapshot_zmq_ctx, ZMQ_PAIR, READER_QUEUE_ADDR);
    assertv(reader != NULL, "Failed to instantiate reader queue socket\n");
}

/* Close everything created above */

private void close_writer_comms() {
    zmq_close(log);
    zmq_close(reader);
    zmq_close(command);
}

/*
 * Copy the necessary capabilities from permitted to effective set (failure is fatal).
 *
 * The writer needs:
 *
 * CAP_IPC_LOCK -- ability to mmap and mlock pages.
 * CAP_SYS_NICE -- ability to set RT scheduling priorities
 * CAP_SYS_ADMIN (Writer) -- ability to set RT IO scheduling priorities (unused at present)
 *
 * These capabilities should be in the CAP_PERMITTED set, but not in CAP_EFFECTIVE which was cleared
 * when the main thread dropped privileges by changing to the desired non-root uid/gid.
 */

private int set_up_writer_capability() {
    cap_t c = cap_get_proc();
    const cap_value_t vs[] = { CAP_IPC_LOCK, CAP_SYS_NICE, CAP_SYS_ADMIN, };

    cap_set_flag(c, CAP_EFFECTIVE, sizeof(vs)/sizeof(cap_value_t), &vs[0], CAP_SET);
    return cap_set_proc(c);
}

/*
 * Set the WRITER thread to real-time priority, if RTPRIO is set...
 */

private int set_writer_rt_scheduling() {

    if( writer_parameters.w_schedprio > 0 ) { /* Then there is RT priority scheduling to set up */
        if( set_rt_scheduling(writer_parameters.w_schedprio) < 0 )
            return -1;

        /* Successfully applied RT scheduling */
        return 1;
    }

    /* RT scheduling not applicable: no RTPRIO set */
    return 0;
}

/*
 * Debug writer parameters
 */

```

Sep 22, 15 16:05

writer.c

Page 4/21

```

private void debug_writer_params() {
    char buf[MSGBUFSIZE];
    wparams *wp = &writer_parameters;

    if(verbose<1)
        return;

    snprintf(buf, MSGBUFSIZE, "WRITER: TMPDIR=%s, SNAPDIR=%s, RTprio=%d; WOF=%g; FrameRAM = %d[MiB], ChunkSize = %d[kiB], nFrames = %d xfrSampleQ = %d[ki]\n",
        tmpdir_path, wp->w_snapdir, wp->w_schedprio, wp->w_writeahead,
        wp->w_lockedram, wp->w_chunksize, wp_nframes, wp_totxfrrsamples/1024);
    zh_put_multi(log, 1, buf);
}

/*
 * -----
 * UTILITY FUNCTIONS USED ONLY BY THE WRITER THREAD
 * -----
 */

/*
 * Test for the presence of a directory by getting a path fd for it.
 */

private int test_directory(int dirfd, const char *name) {
    int ret;

    ret = openat(dirfd, name, O_PATH|O_DIRECTORY); /* Try to open the directory */
    return ret;
}

/*
 * Get a path handle to a directory, creating it if necessary.
 */

private int new_directory(int dirfd, const char *name) {
    int ret;

    ret = test_directory(dirfd, name); /* Try to open the directory */
    if(ret < 0 ) {
        if( errno != ENOENT ) /* OK if it doesn't exist, otherwise fail */
            return -1;
        ret = mkdirat(dirfd, name, 0750); /* Didn't exist, try to create it */
        if( ret < 0 )
            return -1;
        ret = openat(dirfd, name, O_PATH|O_DIRECTORY); /* Try again */
        if(ret < 0) /* Give up on failure */
            return -1;
    }
    return ret;
}

/* ===== Handle the Dir Command ===== */

/*
 * Snapshot working directory parameter(s), used by the D command line.
 */

private param_t snapwd_params[] ={
#define SNAP_SETWD 0
    { "path", NULL, NULL,
      PARAM_TYPE(string), PARAM_SRC_CMD,
      "working (sub-)directory for snapshots"
    },
};

private const int n_snapwd_params = (sizeof(snapwd_params)/sizeof(param_t));

/*
 * Manage the writer's 'working directory': clear the old, resetting to snapdir;
 * find/create and set a new one, clearing an old if necessary.
 */

```

Sep 22, 15 16:05

writer.c

Page 5/21

```

private void clear_writer_wd() {
    int fd = wp_snap_curfd;

    if( fd != wp_snap_dirfd ) {
        wp_snap_curfd = wp_snap_dirfd;
        close(fd);
    }
}

private int set_writer_new_wd(const char *dir) {
    int fd;

    fd = new_directory(wp_snap_dirfd, dir);
    if(fd < 0)
        return -1;
    wp_snap_curfd = fd;
    return 0;
}

/*
 * Process a D command to change the working directory. The command
 * comprises an introductory Dir verb followed by a path=... parameter.
 */

private int process_dir_command(strbuf c) {
    strbuf e = strbuf_next(c);
    param_t *ps = &snapwd_params[0];
    int nps = n_snapwd_params;
    char *path = NULL;
    int err;

    /* Initialise the parameter value pointer */
    setval_param(&ps[SNAP_SETWD], (void **)&path);
    err = set_opt_params_from_string(strbuf_string(c), ps, nps);
    if(err < 0) {
        strbuf_appendf(e, "parameter parsing error at position %d", -err);
        reset_param(&ps[SNAP_SETWD]);
        return -1;
    }
    err = assign_param(&ps[SNAP_SETWD]);
    /* If this string copy fails, it's a programming error! */
    assertv(err==0, "Dir PATH parameter assignment failed: %m");
    reset_param(&ps[SNAP_SETWD]);

    if( !path ) {
        clear_writer_wd();
        return 0;
    }

    /* Path is now instantiated to the given parameter string */
    if(set_writer_new_wd(path) < 0) {
        strbuf_appendf(e, "cannot create path=%s:%m", path);
        return -1;
    }
    return 0;
}

/* ===== Handle the Snap command ===== */

/*
 * Snapshot parameters, used by the S command line.
 * Local to this thread.
 *
 * Note the #defines, which are used to extract the parameter values
 * when building snapshot descriptors -- there is no need to search
 * for the parameter when we know exactly where it is.
 */

private param_t snapshot_params[] = {
#define SNAP_BEGIN 0
    { "begin", NULL, NULL,
      PARAM_TYPE(int64), PARAM_SRC_CMD,
      "start time of snapshot [ns from epoch]"

```

Sep 22, 15 16:05

writer.c

Page 6/21

```

    },
    #define SNAP_END 1
    { "end", NULL, NULL,
      PARAM_TYPE(int64), PARAM_SRC_CMD,
      "finish time of snapshot [ns from epoch]"
    },
    #define SNAP_START 2
    { "start", NULL, NULL,
      PARAM_TYPE(int64), PARAM_SRC_CMD,
      "start sample of snapshot"
    },
    #define SNAP_FINISH 3
    { "finish", NULL, NULL,
      PARAM_TYPE(int64), PARAM_SRC_CMD,
      "end sample of snapshot"
    },
    #define SNAP_LENGTH 4
    { "length", NULL, NULL,
      PARAM_TYPE(int32), PARAM_SRC_CMD,
      "length of snapshot [samples]"
    },
    #define SNAP_COUNT 5
    { "count", NULL, NULL,
      PARAM_TYPE(int32), PARAM_SRC_CMD,
      "repeat count of snapshot"
    },
    #define SNAP_PATH 6
    { "path", NULL, NULL,
      PARAM_TYPE(string), PARAM_SRC_CMD,
      "storage path of snapshot data"
    },
};

private const int n_snapshot_params = (sizeof(snapshot_params)/sizeof(param_t));

/*
 * -----
 * FUNCTIONS ETC. TO MANAGE SNAPSHOT DESCRIPTORS
 *
 * The writer maintains a list of "active" snapshot descriptors. A descriptor
 * is created in response to an S command and is "active" until it has been both
 * (a) completely processed and also (b) reported back in response to a Z
 * command. These data structures are entirely private to the writer.
 * -----
 */

/*
 * Allocate and free snap_t structures
 */

private uint16_t snap_counter = 0;

private snap_t *alloc_snapshot() {
    snap_t *ret = calloc(1, sizeof(snap_t));

    if( !snap_counter ) snap_counter++; /* Avoid snapshots called 0000 */

    if(ret) {
        init_queue( snap2qp(ret) );
        ret->s_dirfd = -1;
        ret->s_name = snap_counter++;
        init_queue( snap2fq(ret) );
    }
    return ret;
}

private void free_snapshot(snap_t *s) {
    if( !queue_singleton(snap2qp(s)) )
        de_queue(snap2qp(s));
    assertv(queue_singleton(snap2fq(s)),
            "Freeing snapshot %p with non-empty file queue %p", s, queue_next(snap2fq(s)));
    if(s->s_dirfd >= 0)
        close(s->s_dirfd);
    if(s->s_path)

```

Sep 22, 15 16:05

writer.c

Page 7/21

```

    free((void *)s->s_path);
    free( (void *)s );
}

/* Debugging routine to return unique name */
uint16_t snapshot_name(snap_t *s) {
    return s->s_name;
}

/*
 * Display snapshot status codes
 */

const char *snapshot_status(int st) {
    private const char *stab[] = {
        "INI", "ERR", "PRP", "RDY", "...", ">>>", "+++", "DON", "FIN",
    };
    if(st>=0 && st<sizeof(stab)/sizeof(char *))
        return stab[st];
    return "???" ;
}

/*
 * Manage the writer snapshot queue:
 *
 * - Check the parameters in an S command
 */

private int check_snapshot_params(param_t ps[], strbuf e) {
    int ret;

    /* path= is MANDATORY */
    if( !ps[SNAP_PATH].p_str ) {
        strbuf_appendf(e, "missing PATH parameter");
        return -1;
    }
    ret = assign_param(&ps[SNAP_PATH]);
    /* If this string copy fails, it's a programming error! */
    assertv(ret==0, "Snapshot PATH parameter assignment failed: %m");

    /* EITHER begin= OR start= is MANDATORY */
    if( !ps[SNAP_BEGIN].p_str && !ps[SNAP_START].p_str ) {
        strbuf_appendf(e, "neither BEGIN nor START present");
        return -1;
    }

    /* IF begin= THEN end= XOR length= AND NOT finish= is REQUIRED */
    if( ps[SNAP_BEGIN].p_str ) {
        if( ps[SNAP_FINISH].p_str ) {
            strbuf_appendf(e, "BEGIN with FINISH present");
            return -1;
        }
        if( !ps[SNAP_END].p_str && !ps[SNAP_LENGTH].p_str ) {
            strbuf_appendf(e, "BEGIN but neither END nor LENGTH present");
            return -1;
        }
        if( ps[SNAP_END].p_str && ps[SNAP_LENGTH].p_str ) {
            strbuf_appendf(e, "BEGIN with both END and LENGTH present");
            return -1;
        }
    }
    ret = assign_param(&ps[SNAP_BEGIN]); /* Error implies bad number */
    if(ret < 0) {
        strbuf_appendf(e, "cannot assign BEGIN value %s: %m", ps[SNAP_BEGIN].p_str);
        return -1;
    }
    if(ps[SNAP_END].p_str) {
        ret = assign_param(&ps[SNAP_END]); /* Error implies bad number */
        if(ret < 0) {
            strbuf_appendf(e, "cannot assign END value %s: %m", ps[SNAP_END].p_str);
            return -1;
        }
    }
    if(ps[SNAP_LENGTH].p_str) {
        ret = assign_param(&ps[SNAP_LENGTH]); /* Error implies bad number */
        if(ret < 0) {

```


Sep 22, 15 16:05

writer.c

Page 8/21

```

        strbuf_appendf(e, "cannot assign LENGTH value %s: %m", ps[SNAP_LENGTH].p_str);
        return -1;
    }
}

/* IF start= THEN finish= XOR length= AND NOT end= is REQUIRED */
if( ps[SNAP_START].p_str ) {
    if( ps[SNAP_END].p_str ) {
        strbuf_appendf(e, "START with END present");
        return -1;
    }
    if( !ps[SNAP_FINISH].p_str && !ps[SNAP_LENGTH].p_str ) {
        strbuf_appendf(e, "START but neither FINISH nor LENGTH present");
        return -1;
    }
    if( ps[SNAP_FINISH].p_str && ps[SNAP_LENGTH].p_str ) {
        strbuf_appendf(e, "START with both FINISH and LENGTH present");
        return -1;
    }
    ret = assign_param(&ps[SNAP_START]); /* Error implies bad number */
    if(ret < 0) {
        strbuf_appendf(e, "cannot assign START value %s: %m", ps[SNAP_START].p_str);
        return -1;
    }
    if(ps[SNAP_FINISH].p_str) {
        ret = assign_param(&ps[SNAP_FINISH]); /* Error implies bad number */
        if(ret < 0) {
            strbuf_appendf(e, "cannot assign FINISH value %s: %m", ps[SNAP_FINISH].p_str);
            return -1;
        }
    }
    if(ps[SNAP_LENGTH].p_str) {
        ret = assign_param(&ps[SNAP_LENGTH]); /* Error implies bad number */
        if(ret < 0) {
            strbuf_appendf(e, "cannot assign LENGTH value %s: %m", ps[SNAP_LENGTH].p_str);
            return -1;
        }
    }
}

/* count= is OPTIONAL */
if(ps[SNAP_COUNT].p_str) {
    ret = assign_param(&ps[SNAP_COUNT]); /* Error implies bad number */
    if(ret < 0) {
        strbuf_appendf(e, "cannot assign COUNT value %s: %m", ps[SNAP_COUNT].p_str);
        return -1;
    }
}

/* All required parameters present in legal combination and values parse */
return 0;
}

/*
 * Complete the snap_t structure sample-range contents -- we know the parameter
 * subset is correct We can also assume that the various members of the snap_t
 * structure have been instantiated by parameter assignment handled by the
 * caller. We need the param[] array to determine which case we are handling.
 * No errors can occur here because they are dealt with by the caller(s) of this
 * routine.
 */

private void setup_snapshot_samples(snap_t *s, param_t p[]) {

    /* Start with length= -- if present, no finish= or end= spec. needed */
    if( p[SNAP_LENGTH].p_str ) { /* Length was stored in s_samples, round up to integral number of pages */
        s->s_bytes = s->s_samples * sizeof(sampl_t);
        s->s_bytes += (sysconf(_SC_PAGE_SIZE) - (s->s_bytes % sysconf(_SC_PAGE_SIZE))) % sysconf(_SC_PAGE_SIZE);
        s->s_samples = s->s_bytes / sizeof(sampl_t);
    }

    /* Mandatory EITHER begin= OR start= -- it was begin= */
    if( p[SNAP_BEGIN].p_str ) { /* Begin time was stored in s_first */
        s->s_first = adc_time_to_sample(reader_adc, s->s_first);
    }
}

```

Sep 22, 15 16:05

writer.c

Page 9/21

```

s->s_first = s->s_first - (s->s_first % NCHANNELS); /* Fix to NCHANNELS boundary */
if( !s->s_samples ) { /* No length given, need end from s_last */
    s->s_last = adc_time_to_sample(reader_adc, s->s_last);
    s->s_last = s->s_last + ((NCHANNELS - (s->s_last % NCHANNELS)) % NCHANNELS); /* Round up to integral number of channel sweeps */
    s->s_samples = s->s_last - s->s_first;
    s->s_bytes = s->s_samples * sizeof(sampl_t);
    s->s_bytes += (sysconf(_SC_PAGE_SIZE) - (s->s_bytes % sysconf(_SC_PAGE_SIZE))) % sysconf(_SC_PAGE_SIZE);
    s->s_samples = s->s_bytes / sizeof(sampl_t); /* Round up to integral number of system pages */
}
s->s_last = s->s_first + s->s_samples; /* Calculate end point using rounded-up sample count */
}

/* Mandatory EITHER begin= OR start= -- it was start= */
if( p[SNAP_START].p_str ) { /* Start sample was stored in s_first */
    if( !s->s_samples ) { /* No length given, need end from s_last */
        s->s_last += ((NCHANNELS - (s->s_last % NCHANNELS)) % NCHANNELS); /* Round up to integral number of channel sweeps */
        s->s_samples = s->s_last - s->s_first; /* Compute requested length */
        s->s_bytes = s->s_samples * sizeof(sampl_t);
        s->s_bytes += (sysconf(_SC_PAGE_SIZE) - (s->s_bytes % sysconf(_SC_PAGE_SIZE))) % sysconf(_SC_PAGE_SIZE);
        s->s_samples = s->s_bytes / sizeof(sampl_t); /* Round up to integral number of system pages */
    }
    s->s_last = s->s_first + s->s_samples; /* Calculate end point using rounded-up sample count */
}

/* Optional count=, default is 1 */
if( !p[SNAP_COUNT].p_str ) { /* The count parameter was written to s_count */
    s->s_count = 1;
}

s->s_pending = 0;
s->s_status = 0;
}

/*
 * Build snapshot from S command line: the main thread passes a ring
 * of strbufs comprising the command buffer and the error buffer.
 *
 * The sequence of operations is:
 * - allocate a snap_t structure and bind the parameter val pointers to it
 * - populate the parameter structures from the string in the command buffer
 * - check the parameter set for correctness (check_snapshot_params)
 * - check the snapshot path and create the dirfd
 * - populate the sample value elements (setup_snapshot_samples)
 * - return the complete structure
 *
 * Errors arising during the above process cause an error status mark and are
 * reported in the error buffer.
 */

private snap_t *build_snapshot_descriptor(strbuf c) {
    strbuf e = strbuf_next(c);
    param_t *ps = &snapshot_params[0];
    int nps = n_snapshot_params;
    const char *path = NULL;
    snap_t *ret;
    int err;
    int i;

    if( !(ret = alloc_snapshot()) ) { /* Allocation failed */
        strbuf_appendf(e, "unable to allocate snapshot descriptor. %m");
        return ret;
    }

    /* Initialise the targets for the parameters */
    setval_param(&ps[SNAP_BEGIN], (void **) &ret->s_first);
    setval_param(&ps[SNAP_END], (void **) &ret->s_last);
    setval_param(&ps[SNAP_START], (void **) &ret->s_first);
    setval_param(&ps[SNAP_FINISH], (void **) &ret->s_first);
    setval_param(&ps[SNAP_LENGTH], (void **) &ret->s_samples);
    setval_param(&ps[SNAP_COUNT], (void **) &ret->s_count);
    setval_param(&ps[SNAP_PATH], (void **) &path);

    /* Process the S command parameters */
    err = set_params_from_string(strbuf_string(c), ps, nps);
    if(err < 0) { /* Error parsing command string */

```

Sep 22, 15 16:05

writer.c

Page 10/21

```

    strbuf_appendf(e, "parameter parsing error at position %d", -err);
    goto FAIL;
}
/* Check the populated parameters and assign to values */
err = check_snapshot_params(ps, e);
if(err < 0) {
    /* Problems put into strbuf by check function */
    goto FAIL;
}

if(ret->s_last <= ret->s_first) { /* Parameter error: end before start */
    strbuf_appendf(e, "end %016llx before start %016llx", ret->s_last, ret->s_first);
    goto FAIL;
}

/* Path may not already exist */
ret->s_dirfd = test_directory(wp_snap_curfd, path);
if(ret->s_dirfd >= 0) {
    /* Then directory already exists */
    strbuf_appendf(e, "requested dir path=%s already exists", path);
    goto FAIL;
}

if( !adc_is_running(reader_adc) ) {
    strbuf_appendf(e, "data acquisition is currently stopped", path);
    goto FAIL;
}

/* Now try to create required directory */
ret->s_dirfd = new_directory(wp_snap_curfd, path);
if(ret->s_dirfd < 0) {
    strbuf_appendf(e, "unable to create dir path=%s: %m", path);
    goto FAIL;
}
ret->s_path = strdup(path);

/* Set up the sample-dependent values -- cannot fail */
setup_snapshot_samples(ret, ps);

/* Finished with the parameters, their values etc. now */
for(i=0; i<nps; i++) reset_param(&ps[i]);

/* All done, no errors */
ret->s_status = SNAPSHOT_PREPARE; /* Structure complete but no files/chunks yet... */
return ret;

FAIL:
for(i=0; i<nps; i++) reset_param(&ps[i]);
free_snapshot(ret);
return NULL;
}

/*
 * Set up snapshot -- create the necessary file descriptor structures etc.
 */

private void setup_snapshot(snap_t *s) {
    snapfile_t *f = alloc_snapfile();

    if(f == NULL) {
        strbuf_appendf(s->s_error, "Failed to allocate file %d/%d", s->s_pending+s->s_done+1, s->s_count);
        s->s_status = SNAPSHOT_ERROR;
        return;
    }
    if( setup_snapfile(f, s) < 0 ) {
        s->s_status = SNAPSHOT_ERROR;
        return;
    }
    s->s_first += s->s_samples; /* Move current sample indices to next file */
    s->s_last += s->s_samples;
    debug_snapfile(f);
}

/*
 * Called when a snapshot file has just been written.
 */

```

Sep 22, 15 16:05

writer.c

Page 11/21

```

private void refresh_snapshot(snap_t *s) {

    if(s->s_status == SNAPSHOT_ERROR) { /* Tidy up after an error */
        while(s->s_pending) { /* There are files that have not got the message */
            assertv(!queue_singleton(snap2fq(s)),
                "Pending file count %d and file header Q mismatch in snapshot %p\n", s->s_pending, s);
            abort_snapfile(qp2file(queue_next(snap2fq(s))));
        }
        return;
    }
    else if(s->s_done == s->s_count) { /* No files left to request */
        s->s_status = SNAPSHOT_COMPLETE;
        return;
    }
    else if(s->s_done + s->s_pending == s->s_count) { /* All required files are in progress */
        return;
    }
    else { /* See if this snapshot should have another file */
        if(wp_nfiles < 2 || s->s_pending == 0) {
            setup_snapshot(s);
        }
    }
}

/*
 * Debugging function for snapshot descriptors...
 */

private void debug_snapshot_descriptor(snap_t *s) {
    char buf[MSGBUFSIZE];

    snprintf(buf, MSGBUFSIZE,
        "Snap %04hx at %p: path '%s' fd %d status %s "
        "sQ[s:%04hx.s:%04hx] "
        "fQ[f:%04hx.f:%04hx] "
        "files %d/%d/%d "
        "S:%08lx B:%08lx F:%016llx L:%016llx\n",
        s->s_name, s, s->s_path, s->s_dirfd, snapshot_status(s->s_status),
        qp2sname(queue_prev(&s->s_Q)), qp2sname(queue_next(&s->s_Q)),
        qp2fname(queue_prev(&s->s_fileQhdr)), qp2fname(queue_next(&s->s_fileQhdr)),
        s->s_done, s->s_pending, s->s_count,
        s->s_samples, s->s_bytes, s->s_first, s->s_last);
    zh_put_multi(log, 1, &buf[0]);
}

/* ===== Handle a Z(Status) Command ===== */

/*
 * Snapshot status request parameter(s), used by the Z command line.
 */

private param_t status_params[] = {
#define SNAP_NAME 0
    { "name", NULL, NULL,
      PARAM_TYPE(int16), PARAM_SRC_CMD,
      "snapshot name"
    },
};

private const int n_status_params = (sizeof(status_params)/sizeof(param_t));

/*
 * The snapshot s should report its status as follows. If it is a
 * pending snapshot, it should append a status line to the given
 * strbuf x. If it is completed (with or without error) it should
 * transfer its own error strbuf to the chain by inserting it
 * immediately following x. The idea is that on success the caller
 * will ignore the c strbuf and the chain following will give status
 * reports for completed snapshots..
 */

private void snapshot_report_status(strbuf x, snap_t *s) {

    if(s->s_status == SNAPSHOT_DONE) { /* If completed, attach its error strbuf */
        queue_ins_after(strbuf2qp(x), strbuf2qp(s->s_error));
    }
}

```

Sep 22, 15 16:05

writer.c

Page 12/21

```

    s->s_error = (strbuf)NULL;
    return;
}
/* Snapshot is in progress: append a status line to x */
strbuf_appendf(x, "Snap %04hx: %s %d/%d/%d\n",
               s->s_name, snapshot_status(s->s_status),
               s->s_done, s->s_pending, s->s_count);
}

/*
 * Process a Z command to collect and return snapshot status. The command
 * comprises an introductory Z verb followed by an optional name=... parameter.
 *
 * The caller has written an initial NO: prefix into the e strbuf, for
 * the error case. For success, it will rewrite an OK line. The c
 * strbuf is not cleared here or in the caller, since it is used by
 * snapshot_report_status for snapshots in progress.
 */

private int process_status_command(strbuf c) {
    strbuf e = strbuf_next(c);
    param_t *ps = &status_params[0];
    int nps = n_status_params;
    uint16_t name = 0;
    int err;
    snap_t *s = NULL;

    /* Initialise the parameter value pointer */
    setval_param(&ps[SNAP_NAME], (void **)&name);
    err = set_opt_params_from_string(strbuf_string(c), ps, nps);
    if(err < 0) {
        strbuf_appendf(e, "parameter parsing error at position %d", -err);
        reset_param(&ps[SNAP_NAME]);
        return -1;
    }
    err = assign_param(&ps[SNAP_NAME]);
    /* If this string copy fails, it's a programming error! */
    assertv(err==0, "Status NAME parameter assignment failed: %m");
    reset_param(&ps[SNAP_NAME]);

    if(queue_singleton(&snapQ)) { /* There are no snapshots in the queue */
        if(name) {
            strbuf_appendf(e, "Snapshot %hd not found: queue empty", name);
            return -1;
        }
        else {
            strbuf_printf(c, "Files: %d, Xfr samples %d [Mi]\n",
                          wp_nfiles, wp_totxfrsamples/(1024*1024));
            return 0;
        }
    }

    if(name) { /* A specific snapshot is requested */
        for_nxt_in_Q(queue *p, queue_next(&snapQ), &snapQ)
            if(name == qp2snap(p)->s_name) {
                s = qp2snap(p);
                break;
            }
        end_for_nxt;
        if(s == NULL) {
            strbuf_appendf(e, "Snapshot %hd not found", name);
            return -1;
        }
        /* ... we got one */
        strbuf_printf(c, "\n");
        snapshot_report_status(c, s);
        if(s->s_status == SNAPSHOT_DONE) /* If completed, free it */
            free_snapshot(s);
    }
    else { /* Otherwise, look at all the snapshots in the queue */
        /*
         * Note that, the loop below, we alter the queue being traversed
         * since free_snapshot unlinks the current snapshot. This is OK,
         * since the loop macros have already determined whether the node
         * being worked is the last one or not.
         */
    }
}

```

Sep 22, 15 16:05

writer.c

Page 13/21

```

    */
    strbuf_printf(c, " Files: %d, Xfr space %d [MiB]\n", wp_nfiles, wp_totxfrsamples*sizeof(sampl_t)/(1024*1024));
    for_nxt_in_Q(queue *p, queue_next(&snapQ), &snapQ)
    {
        s = qp2snap(p);
        snapshot_report_status(c, s);      /* Report the status of each one */
        if(s->s_status == SNAPSHOT_DONE) /* If completed, free it */
            free_snapshot(s);
        end_for_nxt;
    }
    return 0;
}

/* ===== Deal with the Snapshot File Queue ===== */

/*
 * -----
 * FUNCTIONS ETC. FOR SNAPSHOT FILE DESCRIPTOR STRUCTURES: ONE OF THESE PER FILE TO CAPTURE.
 * -----
 */

/*
 * Allocate and free snapfile_t structures
 */

private uint16_t snapfile_counter;

private snapfile_t *alloc_snapfile() {
    snapfile_t *ret = calloc(1, sizeof(snapfile_t));

    if(ret) {
        init_queue(&ret->f_Q);
        ret->f_fd = -1;
        ret->f_name = ++snapfile_counter;
    }
    return ret;
}

private void free_snapfile(snapfile_t *f) {
    if(f->f_fd >= 0)
        close(f->f_fd);
    assertv(f->f_chunkQ == NULL, "Freeing snapfile %p with remaining chunks %p\n", f, f->f_chunkQ);
    free((void *)f);
}

/* Debugging routine to return unique name */
public uint16_t snapfile_name(snapfile_t *f) {
    return f->f_name;
}

/*
 * Initialise a snapfile_t structure from a snap_t structure.
 */

private int setup_snapfile(snapfile_t *f, snap_t *s) {
    wparams *wp = &writer_parameters;
    int fd;
    int ret;

    f->f_indexnr = s->s_done+s->s_pending;

    snprintf(&f->f_file[0], FILE_NAME_SIZE, "%016llx.s16", s->s_first);
    fd = openat(s->s_dirfd, &f->f_file[0], O_RDWR|O_CREAT|O_EXCL, 0600);
    if(fd < 0) {
        strbuf_appendf(s->s_error, "Unable to open sample file %s in path %s: %m\n", &f->f_file[0], s->s_path);
        return -1;
    }

    ret = ftruncate(fd, s->s_bytes); /* Pre-size the file */
    if(ret < 0) {
        /* Try to tidy up... */
        strbuf_appendf(s->s_error, "Unable to truncate sample file %s to size %d [B]: %m\n", &f->f_file[0], s->s_bytes);
        unlinkat(s->s_dirfd, &f->f_file[0], 0);
        close(fd);
        return -1;
    }
}

```

Sep 22, 15 16:05

writer.c

Page 14/21

```

}

/* Allocate and initialise the chunks */
int nc = s->s_bytes / wp->w_chunksize; /* Number of milli-chunks to use (because chunksize is in [kiB] */
f->f_nchunks = (nc+1023) / 1024;
f->f_chunkQ = alloc_chunk(f->f_nchunks);
if( f->f_chunkQ == NULL ) {
    strbuf_appendf(s->s_error, "Cannot allocate %d chunks for file %s: %m\n", f->f_nchunks, &f->f_file[0]);
    unlinkat(s->s_dirfd, &f->f_file[0], 0);
    close(fd);
    return -1;
}

/* Basic book-keeping entries from here: no options for failure */
f->f_fd = fd;
f->f_parent = s;
f->f_error = s->s_error;
f->f_written = 0;

/*
 * This next variable accounts for the number of samples we have
 * committed to write. It is initialised by verify() from the
 * locked RAM and overbooking parameters.
 *
 * It is decremented here when we set up a file for capture. It is
 * later incremented in one of two places: for a successfully
 * written chunk it is incremented by the queue message handler; for
 * a failed chunk, it is incremented by abort_file when it processes
 * the chunks in the file's chunk list.
 */

wp_totxfirsamples -= s->s_samples;

/* Go through the chunk queue writing in data */
uint64_t first = s->s_first;
uint64_t rest = s->s_samples;
uint32_t chunk = wp_chunksamples;
uint32_t offset = 0;

for_nxt_in_Q(queue *p, chunk2qp(f->f_chunkQ), chunk2qp(f->f_chunkQ))
    chunk_t *c = qp2chunk(p);
    /* Determine chunk parameters */
    c->c_status = SNAPSHOT_READY;
    c->c_parent = f;
    c->c_error = f->f_error;
    c->c_fd = f->f_fd;
    c->c_ring = NULL; /* The ADC object computes this pointer */
    c->c_frame = NULL; /* The transfer frames are allocated elsewhere */
    c->c_first = first;
    if(rest > chunk && rest < 2*chunk) /* Deal with final partial chunk(s) */
        chunk = rest / 2;
    c->c_samples = chunk;
    c->c_last = first + chunk;
    c->c_offset = offset;
    offset += chunk*sizeof(sampl_t);
    first += chunk;

    /* Add the chunk to the WRITER chunk queue */
    queue *pos = &WriterChunkQ;
    if( !queue_singleton(&WriterChunkQ) ) {
        for_nxt_in_Q(queue *p, queue_next(&WriterChunkQ), &WriterChunkQ);
        chunk_t *h = rq2chunk(p);
        if(h->c_first > c->c_first) {
            pos = p;
            break;
        }
    }
    end_for_nxt;
    queue_ins_before(pos, chunk2rq(c));

end_for_nxt;

f->f_status = SNAPSHOT_READY;
s->s_pending++;
wp_nfiles++; /* One more file in progress */

```

Sep 22, 15 16:05

writer.c

Page 15/21

```

}
return 0;
}

/*
 * Completed file descriptor -- called when file acquisition ends,
 * both normally and exeptionally.
 *
 * We assume that the READER has cleared up any assigned frames when
 * deleting the file chunks in the READER queue. Therefore, at this
 * point, only the file on disk remains -- remove it if there was an
 * error. Adjust the book-keeping in the snap_t structure to show
 * this file as done. Release the chunk descriptors.
 *
 * The file is finally written/gone when the TIDY thread has unmapped
 * the frames released by the READER.
 */

private void completed_snapfile(snapfile_t *f) {
    snap_t *s = f->f_parent;

    if(f->f_fd >= 0)
        close(f->f_fd);

    s->s_pending--;
    wp_nfiles--;          /* One less file in progress */

    release_chunk(f->f_chunkQ); /* Finished with these now */
    f->f_chunkQ = NULL;

    if(f->f_status == SNAPSHOT_ERROR) {
        s->s_status = SNAPSHOT_ERROR;
        unlinkat(s->s_dirfd, &f->f_file[0], 0); /* If the file failed, remove it */
    }
    else {
        s->s_done++;          /* This file is done, it was pending before */
        if(s->s_done == s->s_count) {
            s->s_status = SNAPSHOT_COMPLETE;
            strbuf_printf(s->s_error, "OK Snap %04hx: FIN %d/%d files", snapshot_name(s), s->s_done, s->s_count);
        }
    }
    de_queue(file2qp(f));          /* Remove this one from the snapshot */
    free_snapfile(f);              /* And free the structure */
}

/*
 * Abort a file from the WRITER thread's viewpoint: remove all chunks
 * from the WRITER's chunk queue and mark the file in ERROR state.
 *
 * N.B. The READER AND WRITER both use the rq chunk linkage, and keep
 * track of who has it by means of exchanged messages.
 *
 * Adjust the w_totxfersamples parameter to match new situation.
 */

private void abort_snapfile(snapfile_t *f) {
    snap_t *s = f->f_parent;

    f->f_status = SNAPSHOT_ERROR;

    assertv(f->f_chunkQ != NULL, "Aborted file f:%04hx at %p has an empty chunk queue\n", snapshot_name(f), f);

    for_nxt_in_Q(queue *p, chunk2qp(f->f_chunkQ), chunk2qp(f->f_chunkQ));
    chunk_t *c = qp2chunk(p);
    if(queue_singleton(chunk2rq(c))) {          /* These were chunks in the READER queue */
        if(c->c_status == SNAPSHOT_WAITING) { /* These were pending chunks in transit from WRITER to READER */
            c->c_status = SNAPSHOT_ERROR;
            wp_totxfersamples += c->c_samples;
        }
        continue;
    }
    de_queue(chunk2rq(c));          /* Remove from WRITER chunk queue */
    if(c->c_status == SNAPSHOT_ERROR) /* Release the write commitment for this chunk */
        wp_totxfersamples += c->c_samples;
    end_for_nxt;
}

```


Sep 22, 15 16:05

writer.c

Page 16/21

```

/*
 * Emit debugging data for a given file descriptor.
 */

private void debug_snapfile(snapfile_t *f) {
    snap_t *s = f->f_parent;
    int left = MSGBUFSIZE-1,
        used = 0;
    int i;
    char buf[MSGBUFSIZE];

    used = snprintf(&buf[used], left,
        "File %s (f:%04hx) of snapshot %04hx, at %p: "
        "Q [f:%04hx,f:%04hx] "
        "fd %d ix %d nc %d/%d st %s\n",
        &f->f_file[0], f->f_name, s->s_name, f,
        qp2fname(queue_prev(&f->f_Q)), qp2fname(queue_next(&f->f_Q)),
        f->f_fd, f->f_indexnr, f->f_written, f->f_nchunks, snapshot_status(f->f_status)
    );

    if(used >= left) used = left;
    left -= used;
    i = 0;
    for_nxt_in_Q(queue *p, chunk2qp(f->f_chunkQ), chunk2qp(f->f_chunkQ))
        int u = snprintf(&buf[used], left,
            ">%03d: ", i++);

        if(u >= left) u = left;
        used += u;
        left -= u;
        u = debug_chunk(&buf[used], left, qp2chunk(p));
        used += u;
        left -= u;
    end_for_nxt;
    zh_put_multi(log, 1, &buf[0]);
}

/*
 * -----
 *
 * MAIN LOOP TASKS: deal with command and queue messages as they arrive and transfer
 * chunks to the READER when possible.
 *
 * -----
 */

/*
 * Service the WRITER queue, i.e try to find frames to attach to
 * chunks, and pass such chunks to the READER. Steps are:
 *
 * - check to see what is in the WRITER queue
 * - allocate at least one frame and pass chunk to READER
 * - loop while time remains...
 *
 * The READER receives messages for chunks that are now in Waiting
 * state, and it returns chunks in either Written state or Error
 * state.
 *
 * Note that when the READER returns a chunk in error state there may
 * be other chunks in transit as messages between the WRITER and
 * READER... The WRITER needs to keep track of these and make sure
 * they are released in an orderly fashion.
 */

private uint64_t writer_service_queue(uint64_t start) {
    uint64_t now = start;
    uint64_t stop = start + WRITER_MAX_CHUNK_DELAY;
    int max;

    for(max=WRITER_MAX_CHUNKS_TRANSFER; max > 0 && !queue_singleton(&WriterChunkQ) && now < stop; --max) { /* Only ever do max chunks at the most */
        chunk_t *c = rq2chunk(queue_next(&WriterChunkQ));

        if( map_chunk_to_frame(c) < 0 ) {
            if(c->c_status == SNAPSHOT_ERROR) { /* Something nasty went wrong! */
                abort_snapfile(c->c_parent);
            }
        }
    }
}

```

Sep 22, 15 16:05

writer.c

Page 17/21

```

    }
    max = 0;                /* Couldn't get a frame, so we are done */
}
else {                    /* We succeeded */
    de_queue(chunk2rq(c)); /* Hand the chunk over to the READER thread */
    c->c_status = SNAPSHOT_WAITING;
    c->c_parent->f_pending++;
    int ret = zh_put_msg(reader, 0, sizeof(chunk_t *), (void *)&c);
    assertv(ret==sizeof(chunk_t *), "Message to READER has wrong size %d not %d\n", ret, sizeof(chunk_t *));
}
now = monotonic_ns_clock();
}
return now;                /* Current end-of-loop time */
}

/*
 * Deal with a queue message from the READER thread. These messages
 * are chunk pointers and fall into two disjoint classes. In either
 * case any chunk received here has been detached from the READER's
 * chunk queue and its frame has been released.
 *
 * - a chunk in SNAPSHOT_WRITTEN state:
 *   release the write commitment.
 *
 * If this was the last chunk of a snapfile, then run completed_snapfile.
 *
 * - a chunk in SNAPSHOT_ERROR state:
 *   abort the snapfile.
 *
 * In this case the READER will have released all chunks in its queue
 * and marked them in SNAPSHOT_ERROR state so that the abort_snapfile
 * routine can clean them up. Chunks in transit between WRITER and
 * READER will still be in SNAPSHOT_WAITING state and are tidied by
 * abort_snapfile which runs for the first erroneous chunk. The
 * snapfile structure is tidied by completed_snapfile which runs when
 * the last pending chunk is returned.
 */

private int process_reader_message(void *s) {
    chunk_t *c;
    int ret;
    snapfile_t *f;

    /* We are expecting a chunk pointer message */
    ret = zh_get_msg(s, 0, sizeof(chunk_t *), (void *)&c);
    assertv(ret == sizeof(chunk_t *), "Queue message size wrong %d vs %d\n", ret, sizeof(chunk_t *));
    assertv(c != NULL, "Queue message from READER was NULL pointer\n");

    f = c->c_parent;
    f->f_pending--;

    if(c->c_status == SNAPSHOT_WRITTEN) {
        f->f_written++;
        wp_totxfrrsamples += c->c_samples;
        if(f->f_written == f->f_nchunks) {
            f->f_status = SNAPSHOT_WRITTEN;
            completed_snapfile(f); /* This file is finished -- all chunks were written */
        }
        return true;
    }
    if(c->c_status == SNAPSHOT_ERROR) {
        if(f->f_status != SNAPSHOT_ERROR)
            abort_snapfile(f); /* Tidy the chunk list, marking all into SNAPSHOT_ERROR state */
        if(f->f_pending == 0)
            completed_snapfile(f); /* This file is finished -- no pending chunks in transit */
        return true;
    }

    assertv(false, "Chunk c:%04hx received in unexpected state %s\n", c->c_name, snapshot_status(c->c_status));
}

/* ===== Process Command Messages ===== */

private int process_writer_command(void *s) {
    int used;
    int ret;

```

Sep 22, 15 16:05

writer.c

Page 18/21

```

char *p;
strbuf cmd;
char *cmd_buf;
strbuf err;

used = zh_get_msg(s, 0, sizeof(strbuf), &cmd);
if( !used ) {
    return false;
}

cmd_buf = strbuf_string(cmd);
err = strbuf_next(cmd);

switch(cmd_buf[0]) {
case 'd':
    /* Dir command */
case 'D':
    /* Call the command handler for Dir */
    strbuf_printf(err, "NO: Dir -- ");
    ret = process_dir_command(cmd);
    if(ret == 0) {
        strbuf_printf(err, "OK Dir");
        strbuf_clear(cmd);
    }
    break;

case 'z':
case 'Z':
    strbuf_printf(err, "NO: Zstatus -- ");
    ret = process_status_command(cmd);
    if(ret == 0) {
        strbuf_printf(err, "OK Zstatus.");
    }
    break;

case 's':
    /* Snap command */
case 'S':
    /* Try to build a snapshot descriptor */
    strbuf_printf(err, "NO: Snap -- ");
    snap_t *s = build_snapshot_descriptor(cmd);
    if(s != NULL) {
        /* Snapshot building succeeded */
        queue_ins_after(&snapQ, snap2qp(s));
        strbuf_printf(err, "OK Snap %04hx", s->s_name);
        s->s_error = (strbuf)de_queue((queue *)cmd);
        strbuf_clear(cmd);
        if(verbose > 0)
            debug_snapshot_descriptor(s);
        refresh_snapshot(s);
    }
    else {
        ret = -1;
    }
    break;

default:
    strbuf_printf(err, "NO: WRITER -- unexpected writer command");
    ret = -1;
    break;
}

if(ret < 0) {
    strbuf_revert(cmd);
    zh_put_multi(log, 4, strbuf_string(err), "\n>", &cmd_buf[0], ""); /* Error occurred, log the problem */
    strbuf_clear(cmd);
}
zh_put_msg(s, 0, sizeof(strbuf), (void *)&err);
return true;
}

/*
 * WRITER thread message loop
 */

private void writer_thread_msg_loop() {
    int borrowedtime;
    int ret;

```

Sep 22, 15 16:05

writer.c

Page 19/21

```

int running;
int n;

zmq_pollitem_t poll_list[] =
{
    { reader, 0, ZMQ_POLLIN, 0 },
    { command, 0, ZMQ_POLLIN, 0 },
};

#define N_POLL_ITEMS    (sizeof(poll_list)/sizeof(zmq_pollitem_t))
int (*poll_responders[N_POLL_ITEMS])(void *) =
{
    process_reader_message,
    process_writer_command,
};

/* WRITER initialisation is complete */
writer_parameters.w_running = !die_die_die_now;
zh_put_multi(log, 1, "WRITER thread is initialised");

running = writer_parameters.w_running;
borrowedtime = 0; /* Keeps track of the number of [ms] we owe */

while( running && !die_die_die_now ) {
    int delay = borrowedtime + WRITER_POLL_DELAY; /* This is how long we wait normally in [ms] */

    int ret = zmq_poll(&poll_list[0], N_POLL_ITEMS, (delay<=0? -1 : delay));

    if( ret < 0 && errno == EINTR ) { /* Interrupted */
        zh_put_multi(log, 1, "WRITER loop interrupted");
        break;
    }
    if(ret < 0)
        break;

    if(delay >= 0) /* We did some waiting, we owe no time */
        borrowedtime = 0;

    uint64_t tick = monotonic_ns_clock();

    for(n=0; n<N_POLL_ITEMS; n++) {
        if( poll_list[n].revents & ZMQ_POLLIN ) {
            if( (*poll_responders[n])(poll_list[n].socket) )
                running = true;
        }
    }

    uint64_t tock = writer_service_queue(tick);
    borrowedtime -= (tock-tick+5000000)/1000000; /* Rounded elapsed time in [ms] */
}

/* ===== Thread Startup ===== */

/*
 * WRITER thread main routine
 */

public void *writer_main(void *arg) {
    int ret;

    create_writer_comms();

    if( set_up_writer_capability < 0 ) {
        zh_put_multi(log, 1, "WRITER thread capabilities are deficient");
    }

    ret = set_writer_rt_scheduling();
    switch(ret) {
    case 1:
        zh_put_multi(log, 1, "WRITER RT scheduling succeeded");
        break;
    case 0:
        zh_put_multi(log, 1, "WRITER using normal scheduling: RTPRIO unset");
        break;
    default:
        zh_put_multi(log, 2, "WRITER RT scheduling setup failed: ", strerror(errno));
        break;
    }
}

```

Sep 22, 15 16:05

writer.c

Page 20/21

```

}

debug_writer_params();
writer_thread_msg_loop();
zh_put_multi(log, 1, "WRITER thread terminates by return");

/* Clean up our ZeroMQ sockets */
close_writer_comms();
writer_parameters.w_running = false;
return (void *)"normal exit";
}

/*
 * Verify the parameters for the WRITER and construct the WRITER state.
 * Called by the MAIN thread during start up initialisation.
 */

public int verify_writer_params(wparams *wp, strbuf e) {
    import int tmpdir_dirfd; /* Imported from snapshot.c */
    int ret;

    if( wp->w_schedprio != 0 ) { /* Check for illegal value */
        int max, min;

        min = sched_get_priority_min(SCHED_FIFO);
        max = sched_get_priority_max(SCHED_FIFO);
        if(wp->w_schedprio < min || wp->w_schedprio > max) {
            strbuf_appendf(e, "RT scheduling priority %d not in kernel's acceptable range [%d,%d]",
                           wp->w_schedprio, min, max);

            return -1;
        }
    }

    /*
     * Check that the requested mmap'd transfer RAM size and the
     * transfer chunk size are reasonable.
     */
    if(wp->w_lockedram < MIN_RAM_MB || wp->w_lockedram > MAX_RAM_MB) {
        strbuf_appendf(e, "Transfer Locked RAM parameter %d MiB outwith compiled-in range [%d,%d] MiB",
                       wp->w_lockedram, MIN_RAM_MB, MAX_RAM_MB);

        return -1;
    }
    if(wp->w_chunksize < MIN_CHUNK_SZ || wp->w_chunksize > MAX_CHUNK_SZ) {
        strbuf_appendf(e, "Transfer chunk size %d KiB outwith compiled-in range [%d,%d] KiB",
                       wp->w_chunksize, MIN_CHUNK_SZ, MAX_CHUNK_SZ);

        return -1;
    }

    /* Compute the number of frames available */
    const int pagesize = sysconf(_SC_PAGESIZE);
    int sz = wp->w_chunksize*1024;
    int nfr;
    sz = pagesize * ((sz + pagesize - 1) / pagesize); /* Round up to multiple of PAGE SIZE */
    wp->w_chunksize = sz / 1024;
    nfr = (wp->w_lockedram * 1024*1024) / sz; /* Number of frames that fit in locked RAM */
    if(nfr < MIN_NFRAMES) {
        strbuf_appendf(e, "Adjusted chunk size %d KiB and given RAM %d MiB yield too few (%d < %d) frames",
                       wp->w_chunksize, wp->w_lockedram, nfr, MIN_NFRAMES);

        return -1;
    }
    wp_nframes = nfr;
    wp_chunksamples = wp->w_chunksize * 1024 / sizeof(sampl_t);

    /*
     * Check the writeahead fraction -- this is the proportion by which
     * the locked transfer RAM may be "overbooked". Should be positive
     * and not too big :-)).
     */
    if(wp->w_writeahead < 0 || wp->w_writeahead > 1) {
        strbuf_appendf(e, "Transfer writeahead fraction %g out of compiled-in range [0,1]", wp->w_writeahead);
        return -1;
    }
    wp_totxfamples = nfr*wp->w_chunksize*1024*(1 + wp->w_writeahead) + pagesize-1;

```

Sep 22, 15 16:05

writer.c

Page 21/21

```
wp_totxfrrsamples = pagesize * (wp_totxfrrsamples / pagesize);
wp_totxfrrsamples = wp_totxfrrsamples / sizeof(sampl_t);

wp_nfiles = 0;           /* Currently no files in progress */

/*
 * Check the snapdir directory exists and get a path fd for it.
 * Assumes we are already running as the non-privileged user.
 */
wp_snap_dirfd = new_directory(tmpdir_dirfd, wp->w_snapdir);
if( wp_snap_dirfd < 0 ) { /* Give up on failure */
    strbuf_appendf(e, "Snapdir %s inaccessible: %m", wp->w_snapdir);
    return -1;
}
wp_snap_curfd = wp_snap_dirfd; /* Initial default */

/*
 * Now try to get the memory for the transfer RAM... This maps in a set of
 * anonymous pages so requires CAP_IPC_LOCK capability.
 */
ret = init_frame_system(e, nfr, wp->w_lockedram, wp->w_chunksize);
return ret;
}
```

Sep 16, 15 15:39

writer.h

Page 1/1

```

#
#include "general.h"

/*
 * The ZMQ address for the writer thread
 */

#define WRITER_CMD_ADDR "inproc://Writer-CMD"

#define MIN_RAM_MB      16
#define MAX_RAM_MB      256
#define MIN_CHUNK_SZ    128
#define MAX_CHUNK_SZ    4096
#define MIN_NFRAMES     4

#define WRITER_MAX_CHUNK_DELAY 100 /* [ms] */
#define WRITER_MAX_CHUNKS_TRANSFER 8
#define WRITER_POLL_DELAY 50 /* [ms] */

typedef struct {
    /* These values come from environment and/or argument parameters */
    const char *w_snapdir;
    int w_schedprio;
    int w_lockedram;
    int w_chunksize;
    double w_writeahead;
    /* Thread is running and ready -- set by main routine */
    int w_running;
}
wparams;

export int verify_writer_params(wparams *, strbuf);
export void *writer_main(void *);

#define FILE_NAME_SIZE 32

#define SNAPSHOT_INIT 0 /* Structure just created */
#define SNAPSHOT_ERROR 1 /* Error found during checking or execution */
#define SNAPSHOT_PREPARE 2 /* Structure filled in, but files/chunks not done yet */
#define SNAPSHOT_READY 3 /* Snapshot etc. is ready, but waiting for READER queue space */
#define SNAPSHOT_WAITING 4 /* Snapshot etc. is ready, but waiting for data */
#define SNAPSHOT_WRITING 5 /* Snapshot file's chunks are being written */
#define SNAPSHOT_WRITTEN 6 /* Snapshot's chunk has been successfully written */
#define SNAPSHOT_COMPLETE 7 /* Snapshot written correctly (off queue) */
#define SNAPSHOT_DONE 8 /* Structure is finished with */

export const char *snapshot_status(int);

```